

# Cloud Computing

SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# The Vision

- ▶ In 1969, Leonard Kleinrock, one of the chief scientists of the original Advanced Research Projects Agency Network (ARPANET) predicted Computer utilities like water, electricity
- ▶ I don't care where my servers are, who manages them, where my documents are stored, or where my applications are hosted. I just want them always available and access them from any device connected through Internet. And I am willing to pay for this service for as much as I need it.
- ▶ IT services are traded as utilities in an open market
- ▶ No infrastructure needed on consumer side

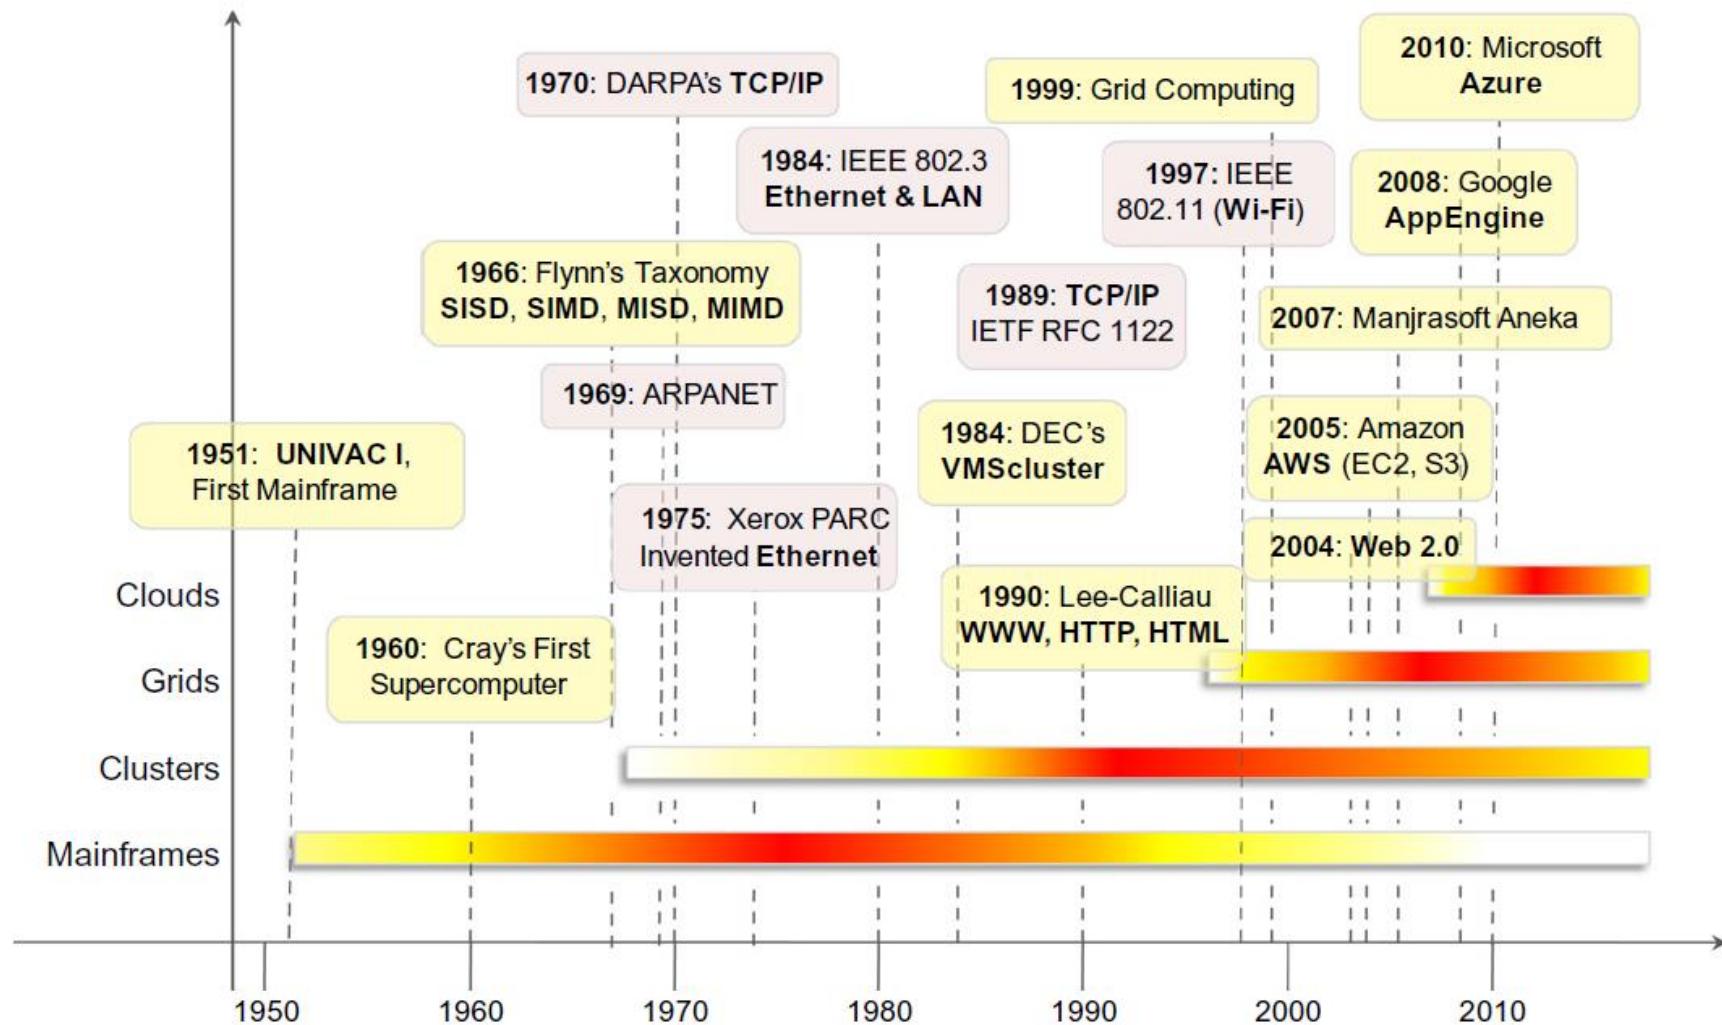
# Definition

- ▶ Internet plays important role
- ▶ Armbrust: Cloud computing refers to both the applications delivered as services over the Internet and the hardware and system software in the datacenters that provide those services.
- ▶ Gartner: a style of computing in which massively scalable IT-enabled capabilities are delivered as a service to external customers using Internet technologies
- ▶ U.S. National Institute of Standards and Technology (NIST): Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction.
- ▶ Zero capital expenditure is necessary to get started.

# Benefits

- ▶ No up-front commitments
- ▶ On-demand access
- ▶ Nice pricing
- ▶ Simplified application acceleration and scalability
- ▶ Efficient resource allocation
- ▶ Energy efficiency
- ▶ Seamless creation and use of third-party services

# Evolution



# Mainframe



# Web 2.0

- ▶ Web is the primary interface
- ▶ Interactivity and flexibility into web pages
- ▶ XML, AJAX, Web Services
- ▶ Variety of devices: mobile phones, car dashboards, TV sets, IoT
- ▶ Extremely dynamic applications, lightweight, loose coupling
- ▶ Wide availability via media accessibility or affordability
- ▶ Google maps, flickr, facebook, twitter, youtube

# Evolution (contd)

- ▶ Service oriented
- ▶ Utility oriented
- ▶ Grid computing
- ▶ HPC
- ▶ Parallel Programming
- ▶ Distributed Computing

# Challenges

- ▶ How many and how long for maximizing benefit
- ▶ Tradeoff between availability and price
- ▶ Longevity – how long your provider will be in business?
- ▶ Business Continuity: 9/11, multi region
- ▶ SLA
- ▶ Quality of Service (QoS): measure of user experience
- ▶ Security
- ▶ Regulatory, Legal

# Summary

- ▶ Vision
- ▶ Definition
- ▶ Benefits
- ▶ Evolution
- ▶ Challenges

# Assignment

- ▶ Signup for AWS
- ▶ Create Ubuntu VM
- ▶ Create private – public keypair
- ▶ Connect to VM using putty and .ppk
- ▶ Sudo apt install apache2
- ▶ Edit /var/www/html/index
- ▶ Edit inbound rules to allow port 80
- ▶ Connect to <your VM's public address> from browser

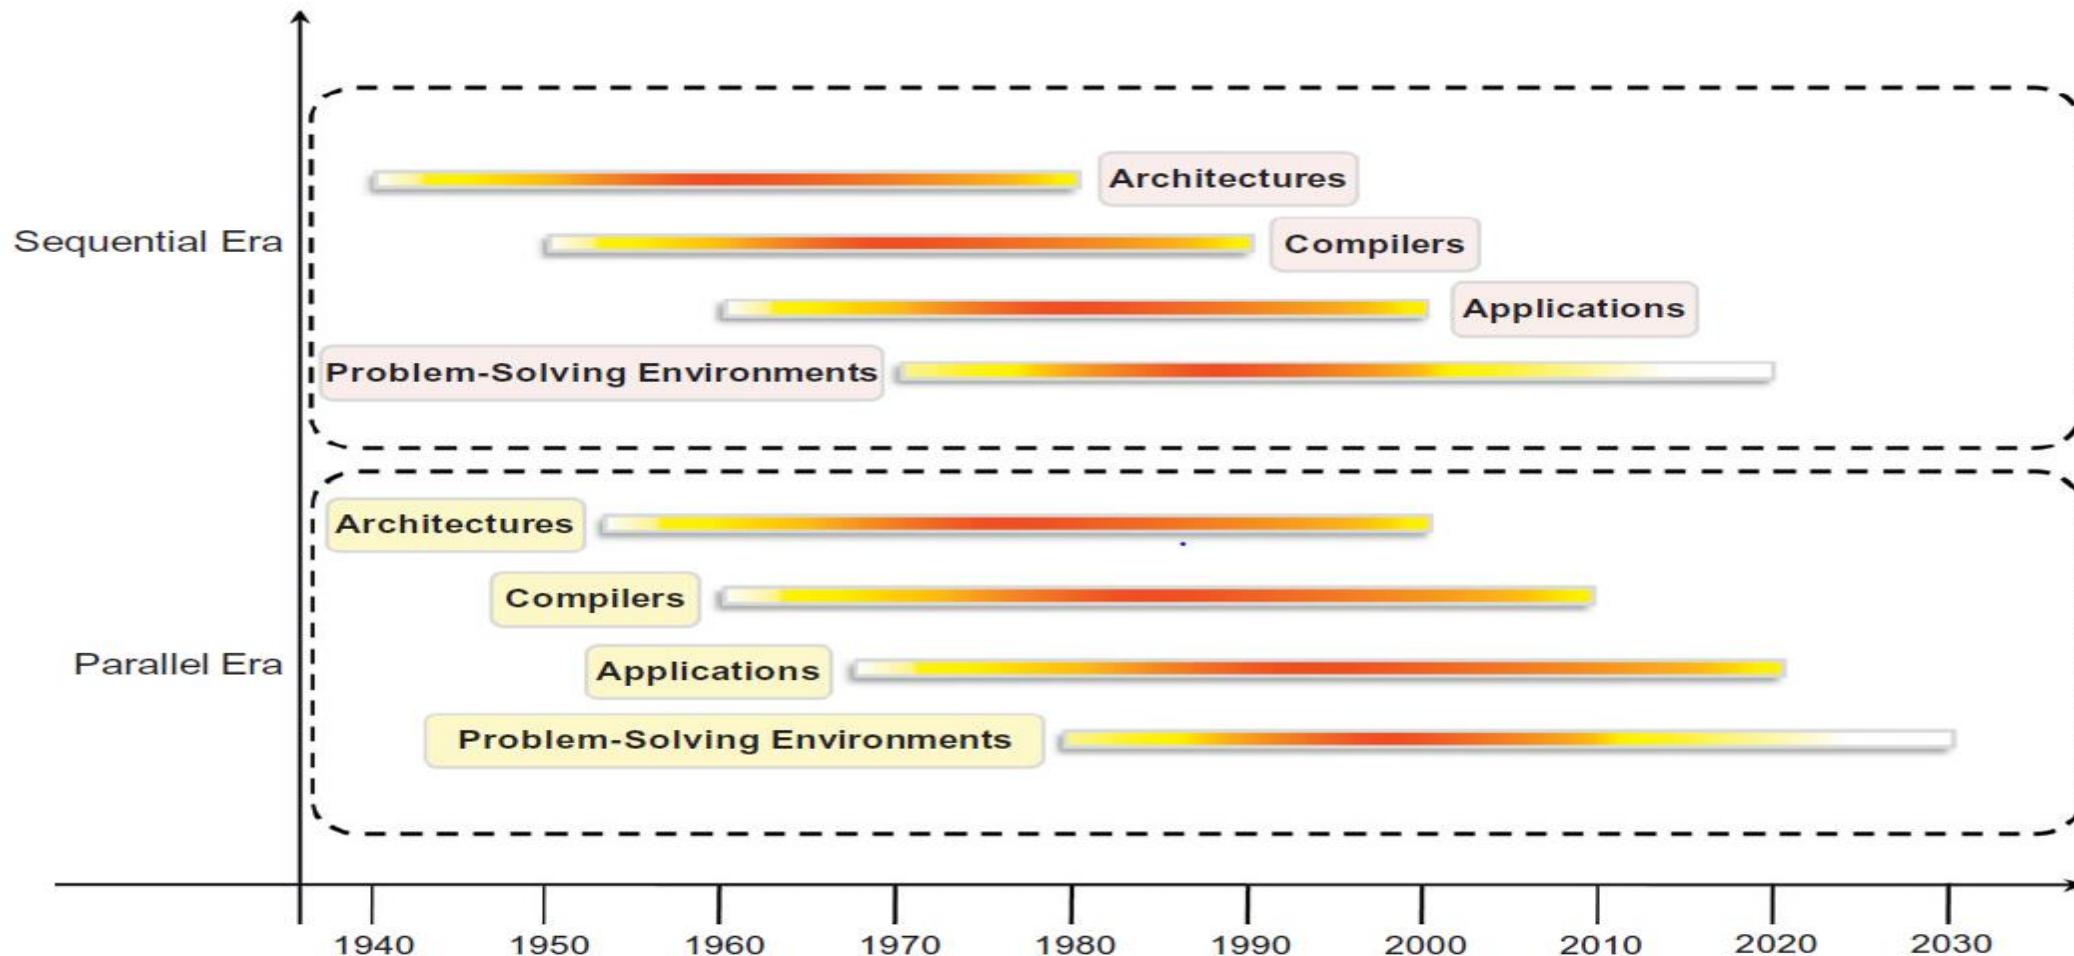
Complete video: <https://youtu.be/RLeQ9INROr0>

# Parallel and Distributed Computing

SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# Computing Era



# Factors for parallel and distributed computing

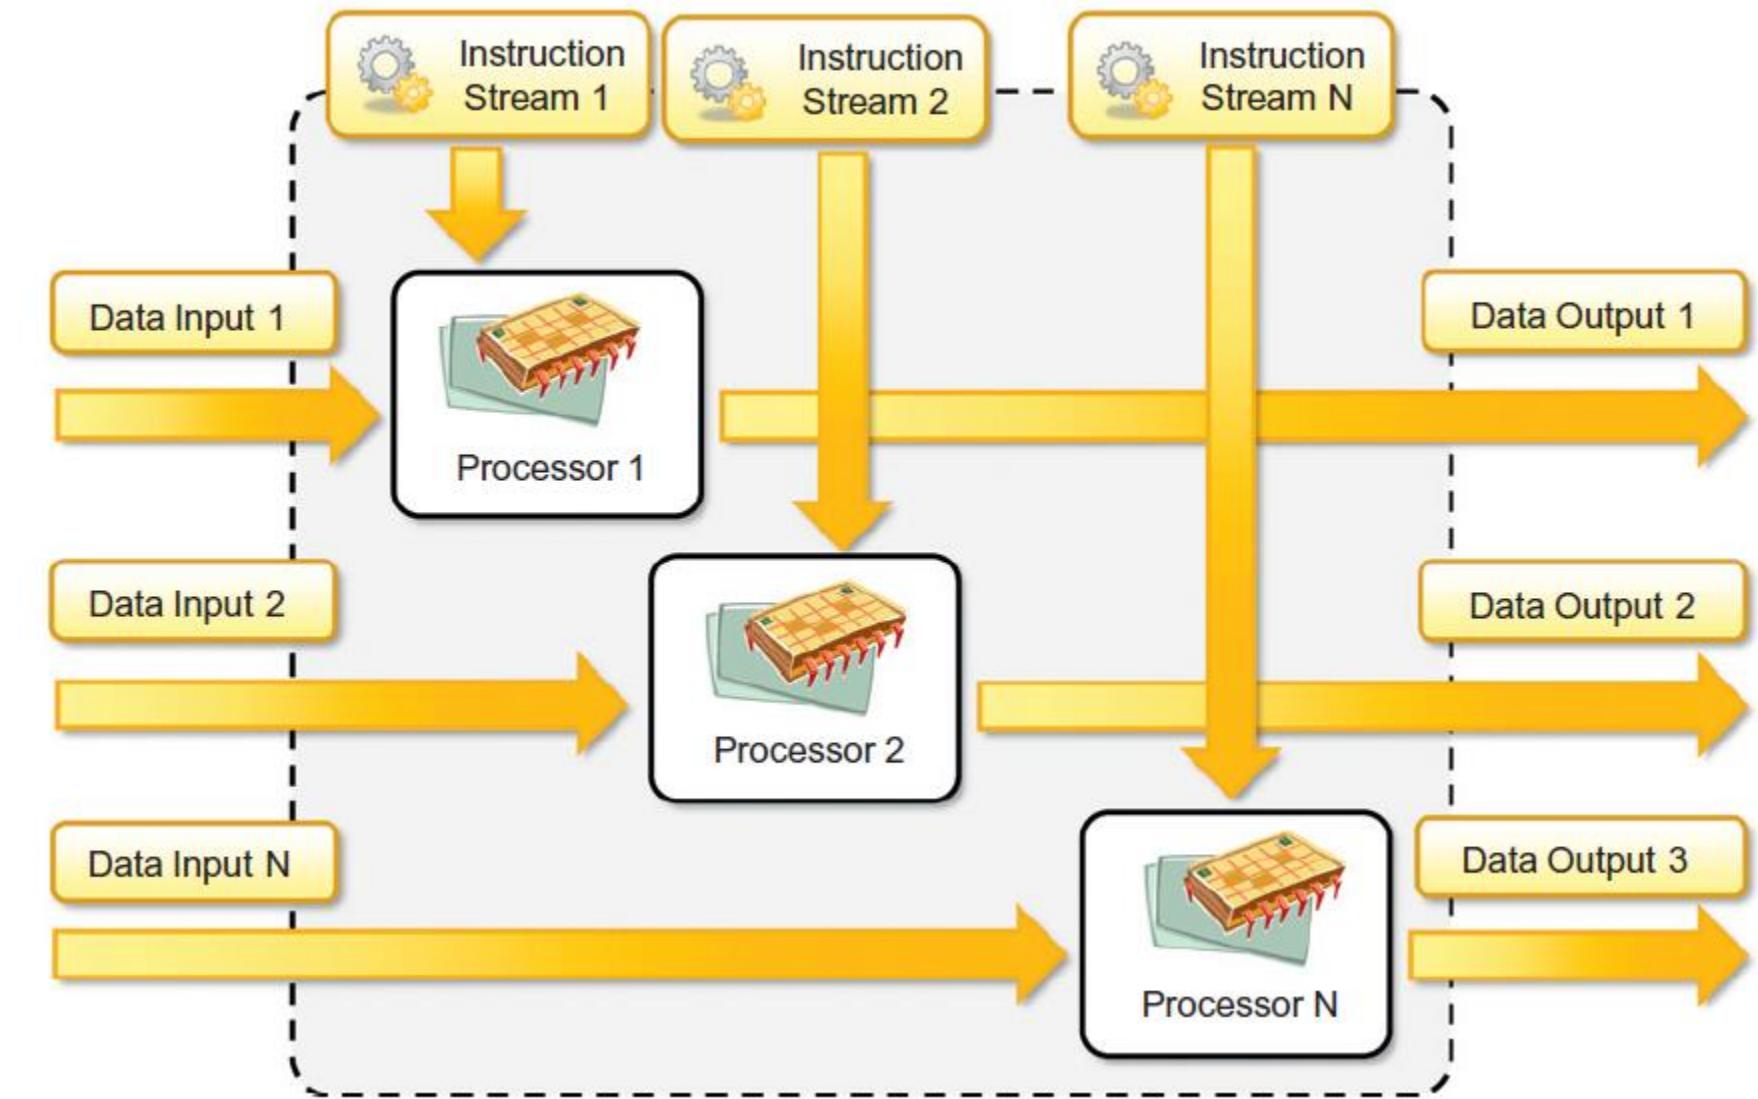
- ▶ Ever increasing computation requirements
- ▶ Saturation reached of sequential computation
- ▶ Hardware improvements in pipelining are non scalable, required compiler technologies are difficult
- ▶ Vector processing is useful only for scientific applications with matrices and not generic e.g. for databases
- ▶ Improvements in networking
- ▶ Parallel processing technology is mature, commercially viable, wide support with development tools and environments
- ▶ Scaling: scale up, scale out

# Parallel and Distributed Computing

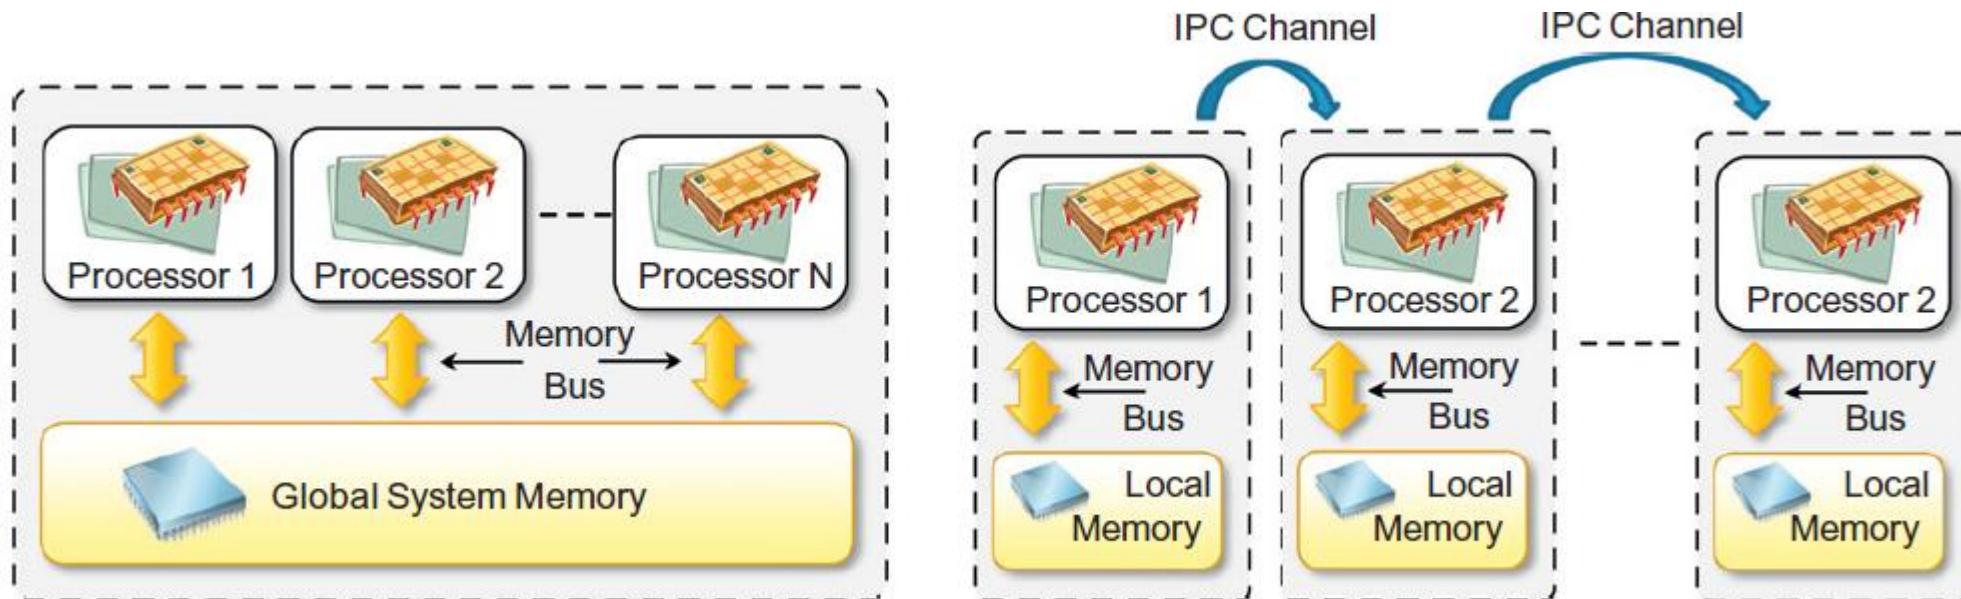
- ▶ Parallel
  - ▶ Computation is divided among several processors
  - ▶ Sharing the same memory
  - ▶ Homogeneity of components
  - ▶ Distributed shared memory connected by InfiniBand
- ▶ Distributed
  - ▶ Computation broken down into units and executed concurrently
  - ▶ Wider range of systems and applications
  - ▶ Heterogeneous

# Hardware architectures for parallel processing

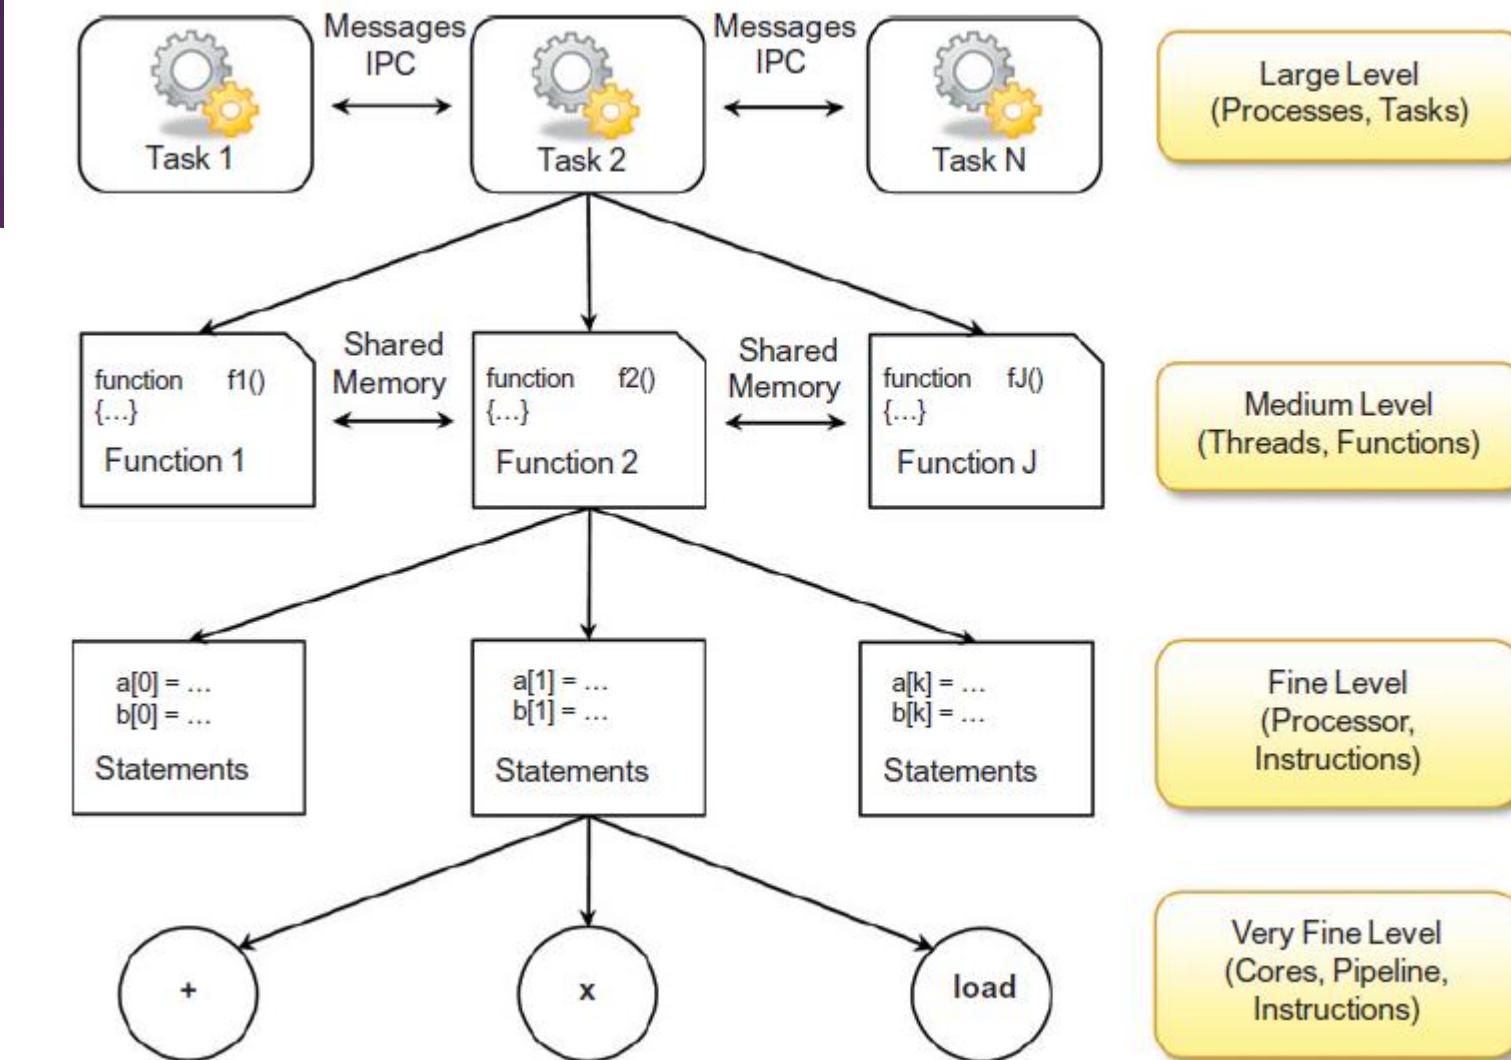
- ▶ SISD
- ▶ SIMD
- ▶ MISD
- ▶ MIMD
- ▶ Data / Process Parallelism



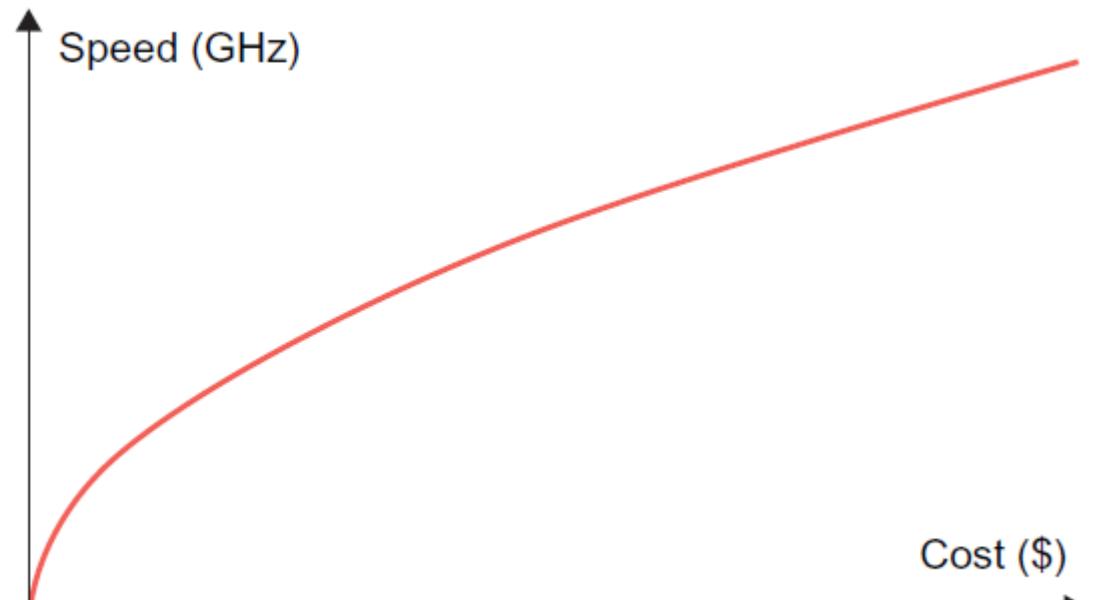
# Shared and Distributed Memory MIMD



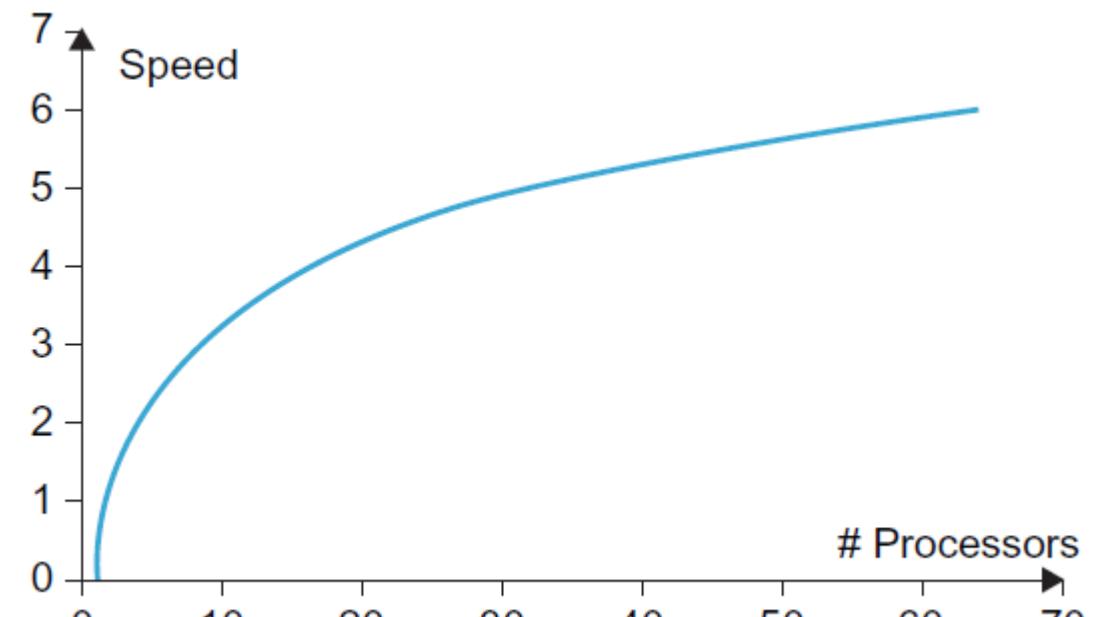
# Levels of Parallelism



# Speed up graphs



Square root

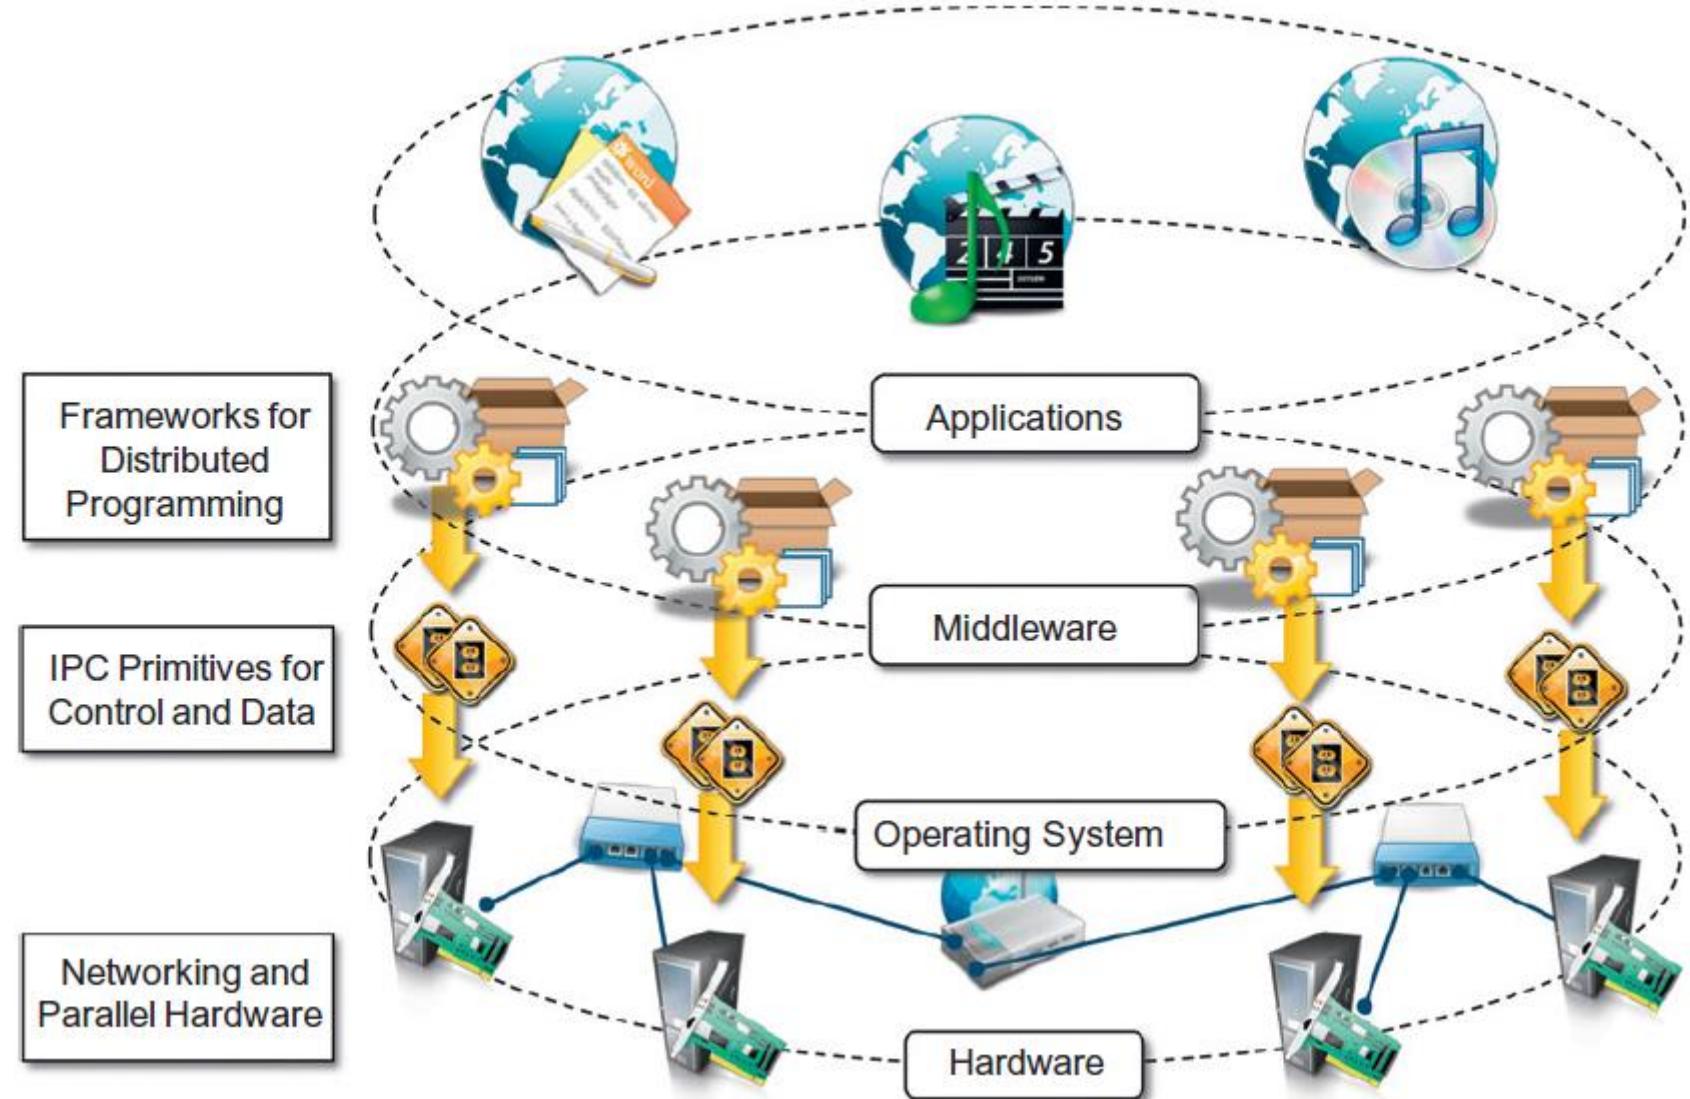


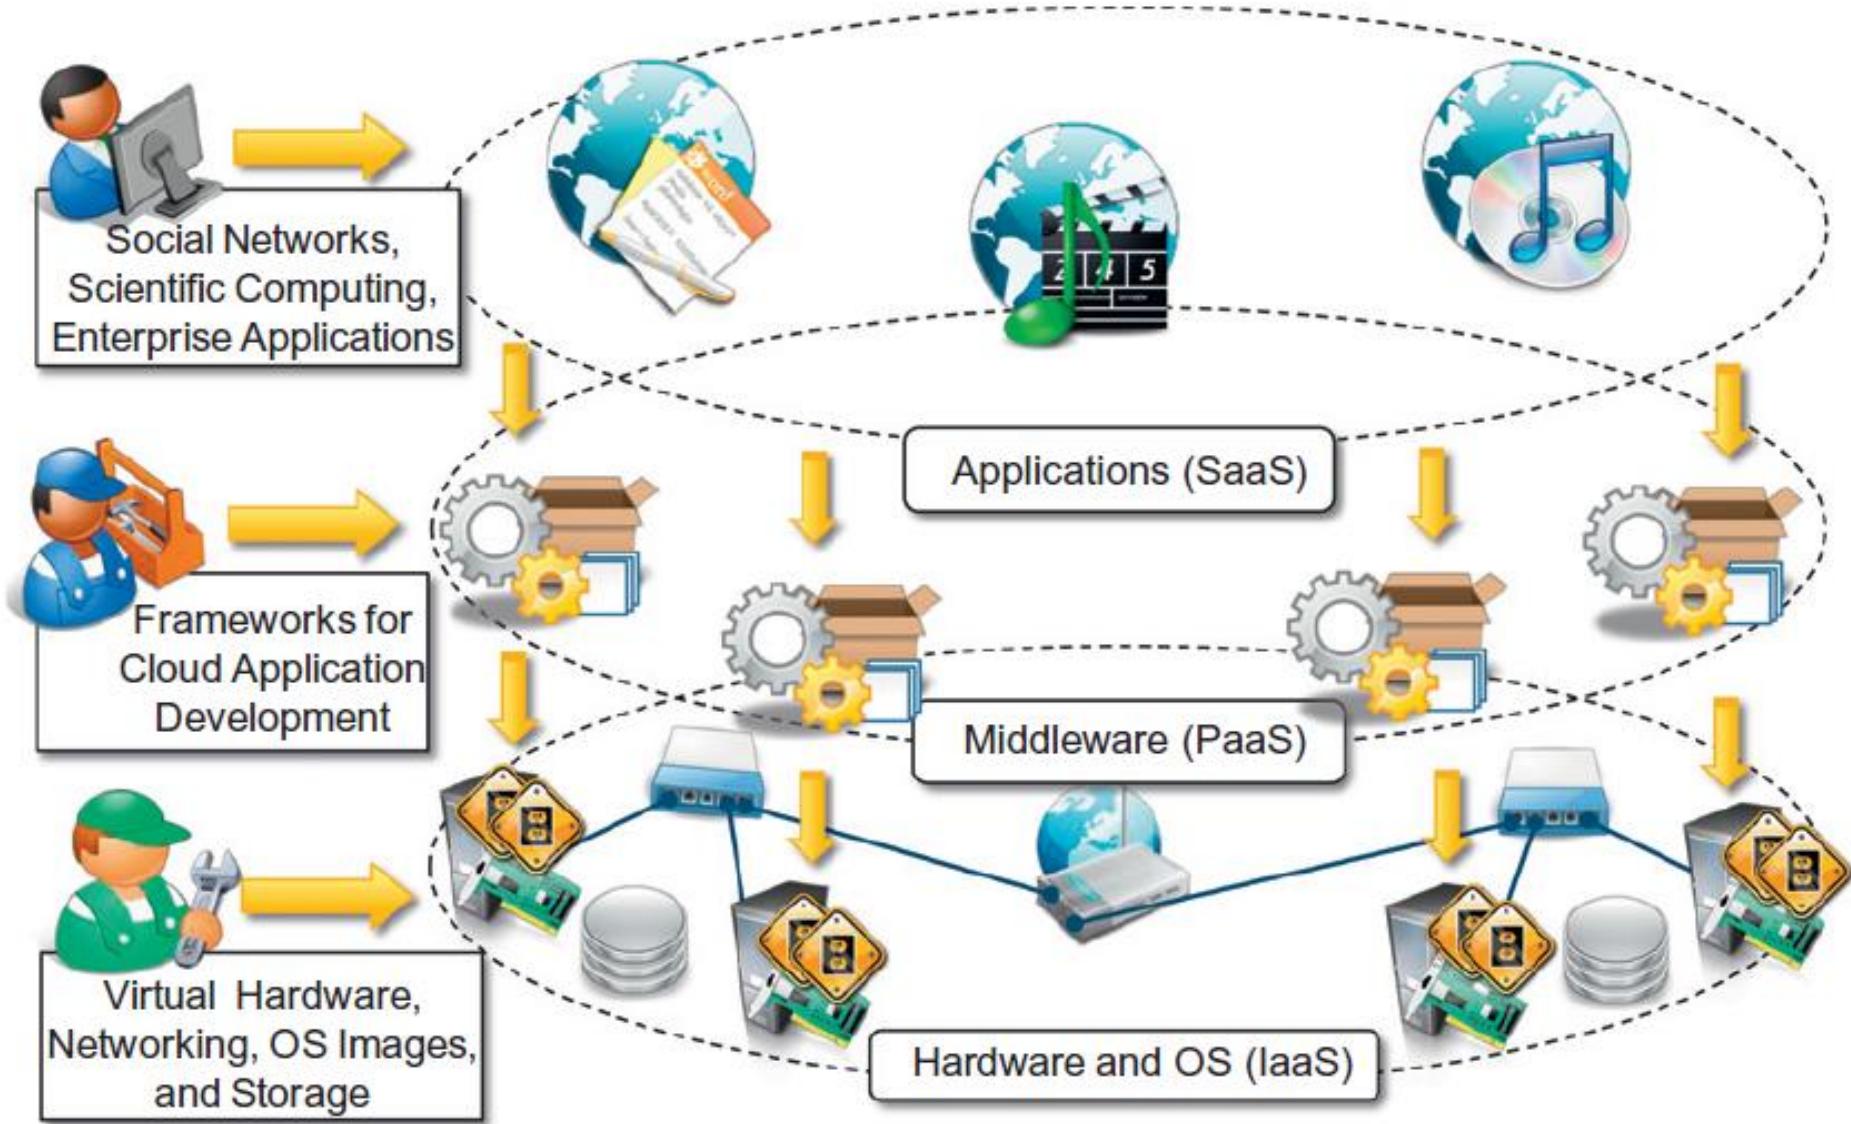
Logarithmic



## Distributed Computing:

Collection of independent computers that appears to its users as a single coherent system



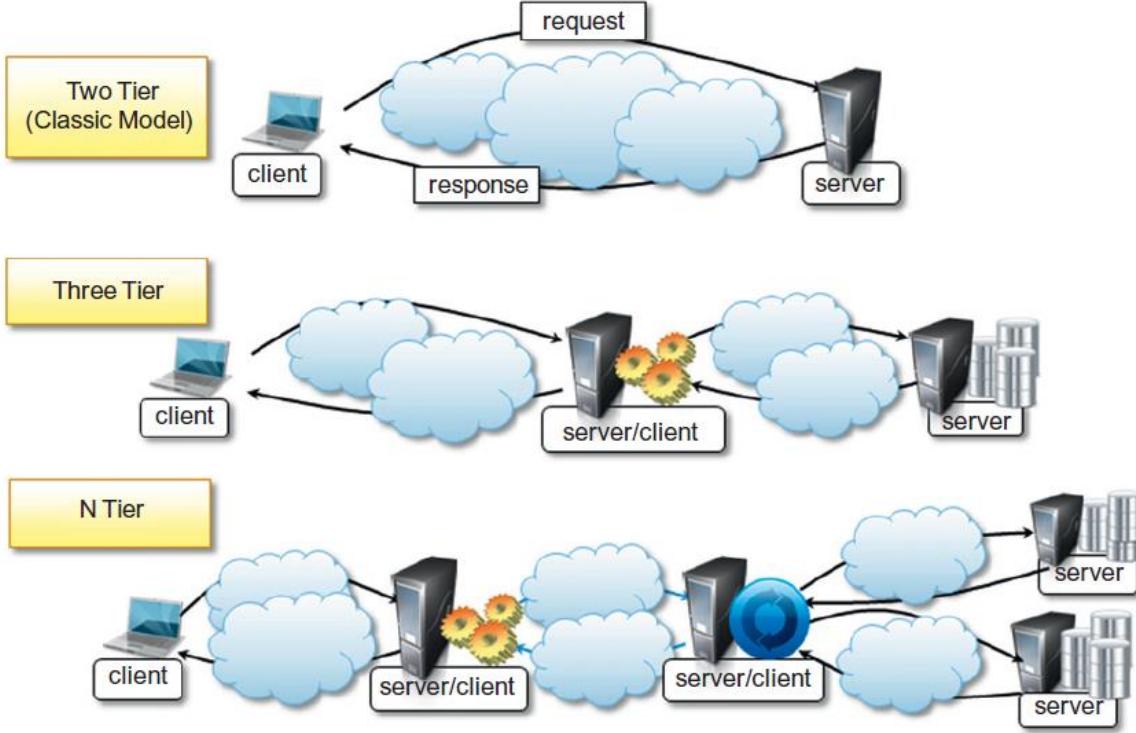


# Software Architectural Styles

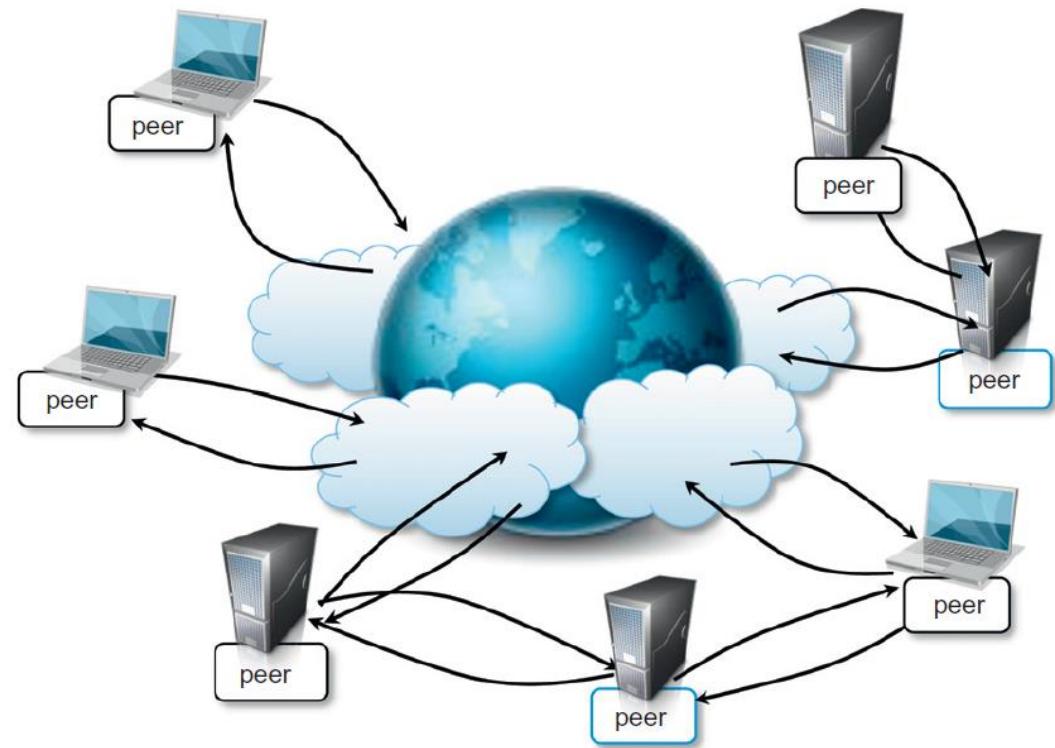
Category	Most Common Architectural Styles
Data-centered	Repository Blackboard
Data flow	Pipe and filter Batch sequential
Virtual machine	Rule-based system Interpreter
Call and return	Main program and subroutine call/top-down systems Object-oriented systems Layered systems
Independent components	Communicating processes Event systems

# System Architectural Styles

Client - Server



Peer to Peer



# Inter Process Communication (IPC)

Message: discrete amount of information that is passed from one entity to another

- ▶ Message passing: MPI, OpenMP
- ▶ Remote Procedure call (RPC): marshaling parameters and return values
- ▶ Distributed objects: state management and lifetime – CORBA, COM, .Net Remoting
- ▶ Distributed agents and active objects: objects with control thread
- ▶ Web Services: RPC over HTTP, Simple Object Access Protocol (SOAP) or Representational State Transfer (REST) protocols

# Messaging Models

- ▶ Point to point: direct addressing to identify the message receiver, no central infrastructure, initiated by sender
  - ▶ Direct
  - ▶ Queue based
- ▶ Publish and subscribe: topic or event, one to many, kafka
  - ▶ Push: publisher notifies all interested subscribers
  - ▶ Pull: subscriber poll for the messages of their interest
- ▶ Request-reply: each request has a reply, mostly point to point and not publish subscribe

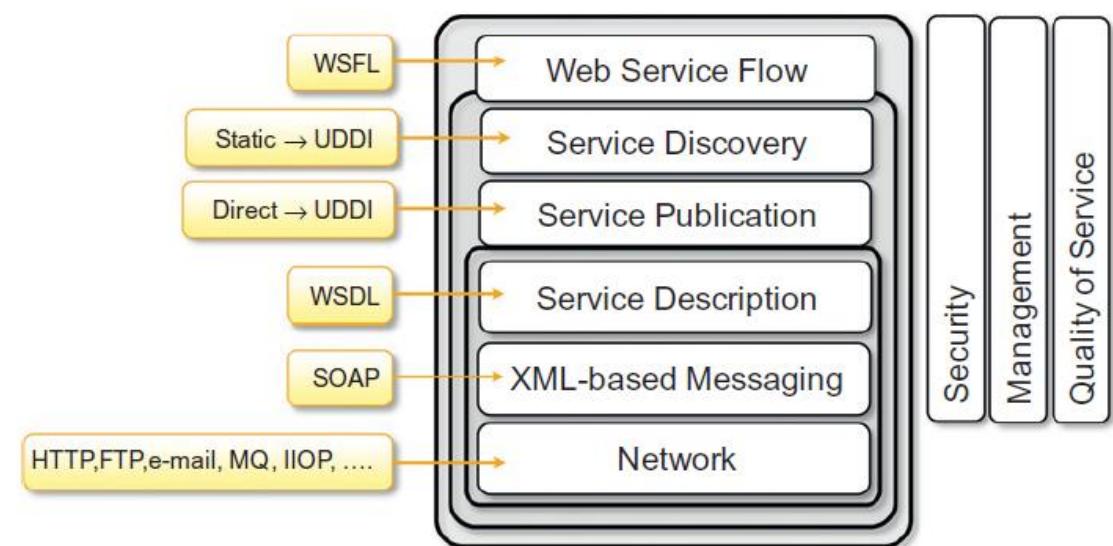
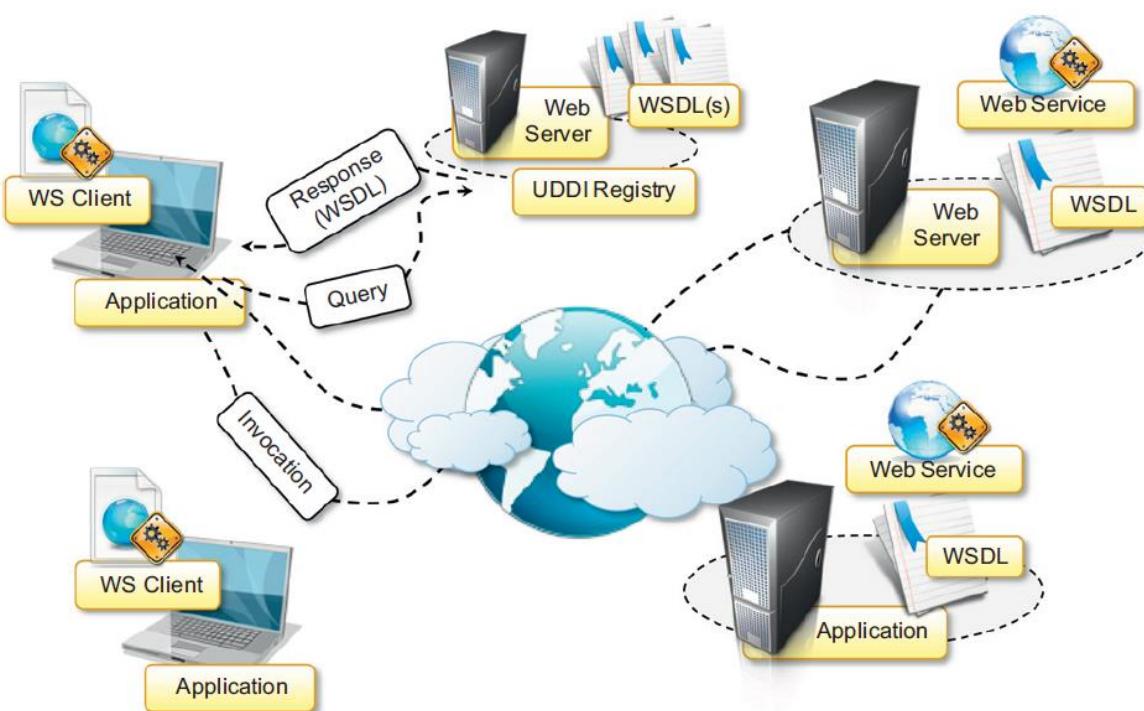
# Technologies

- ▶ RPC: gRPC
  - ▶ Server procedures
  - ▶ Register remote procedures with RPC Server
  - ▶ Client invoking remote procedures
- ▶ Distributed Object Frameworks: proxy skeleton
  - ▶ Objects distributed across heterogeneous network acting coherently
  - ▶ Server registry of active objects is published
  - ▶ Client accesses these objects using interfaces and invokes methods

# Service Oriented Architecture (SOA)

- ▶ Service is a unit of functionality (Don Box)
  - ▶ Explicit boundaries
  - ▶ Autonomous
  - ▶ Expressed as schema and contracts, not class or interface definitions
  - ▶ Semantic compatibility is determined based on policy
- ▶ Guiding Principles
  - ▶ Standardized service contracts
  - ▶ Loose coupling
  - ▶ Abstraction
  - ▶ Reusability
  - ▶ Autonomy
  - ▶ Stateless
  - ▶ Discoverable
  - ▶ Composable

# Web Services



# Abstraction and Virtualization

SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# Abstraction

- Connecting multiple systems together leveraging their interfaces, not worrying about internal details
- Uber / Ola as an example for transportation service
- OSI – 7 layers of networking
- ODBC – abstracting databases
- OpenGL – images
- Internet is an abstraction layer to cloud computing
- Translation from abstraction layer to underlying physical layer

# Open Systems Interconnection

Vendor  
Interconnection

Layer  
Interconnection

Independent  
development at  
each layer

## OSI Model

Layer 7 - Application

Layer 6 - Presentation

Layer 5 - Session

Layer 4 - Transport

Layer 3 - Network

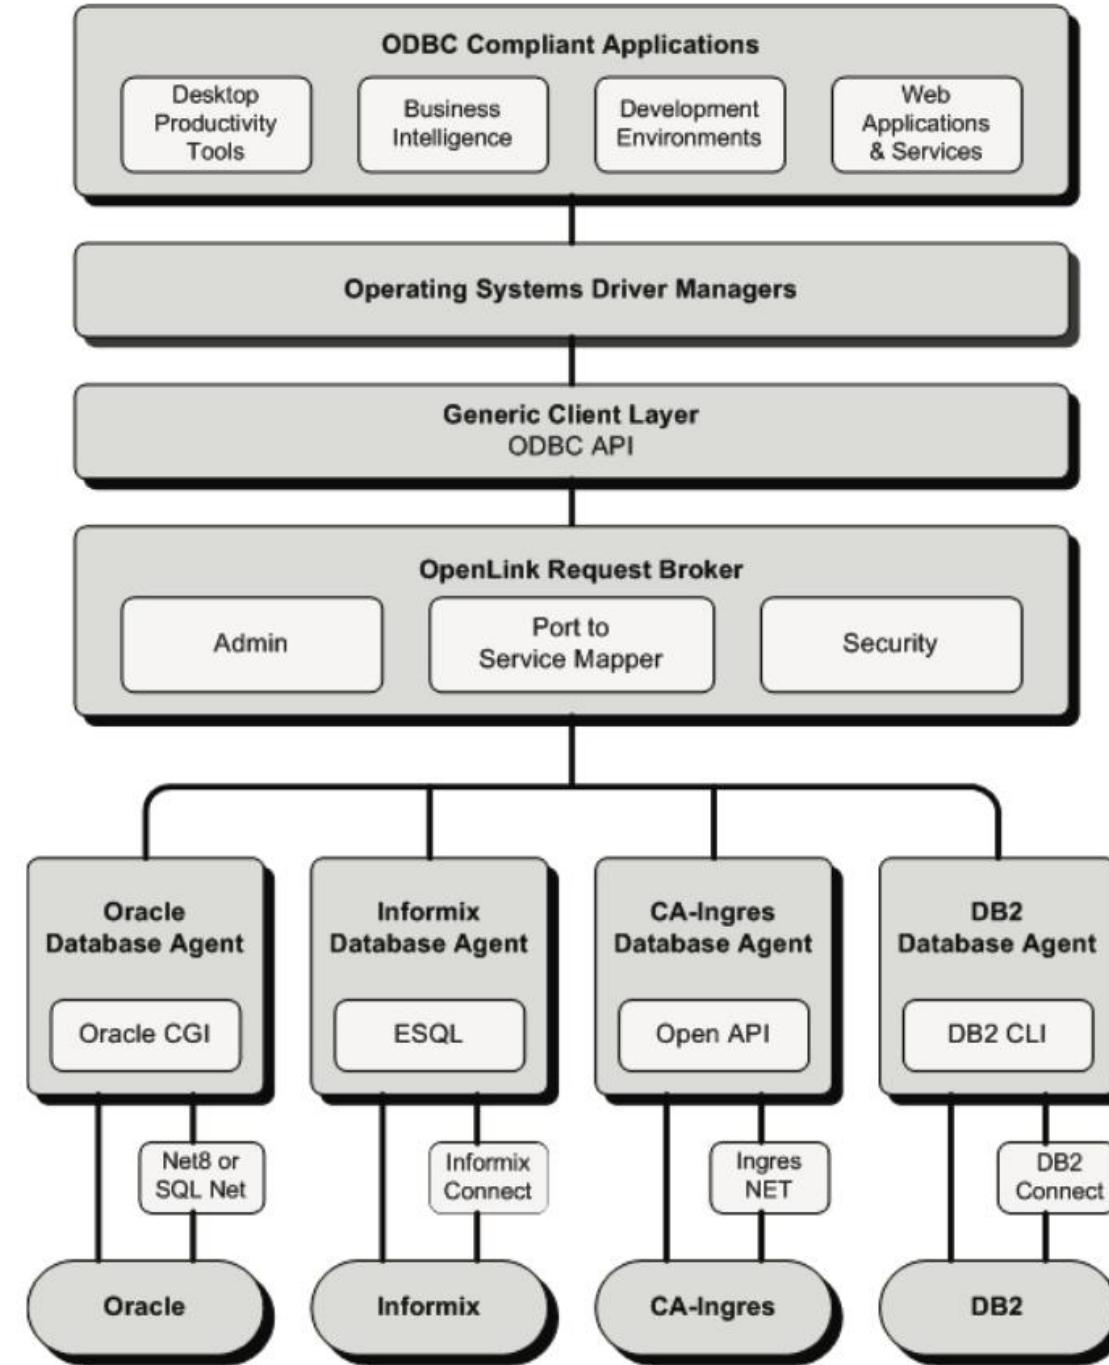
Layer 2 - Data Link

Layer 1 - Physical

# Open DataBase Connection

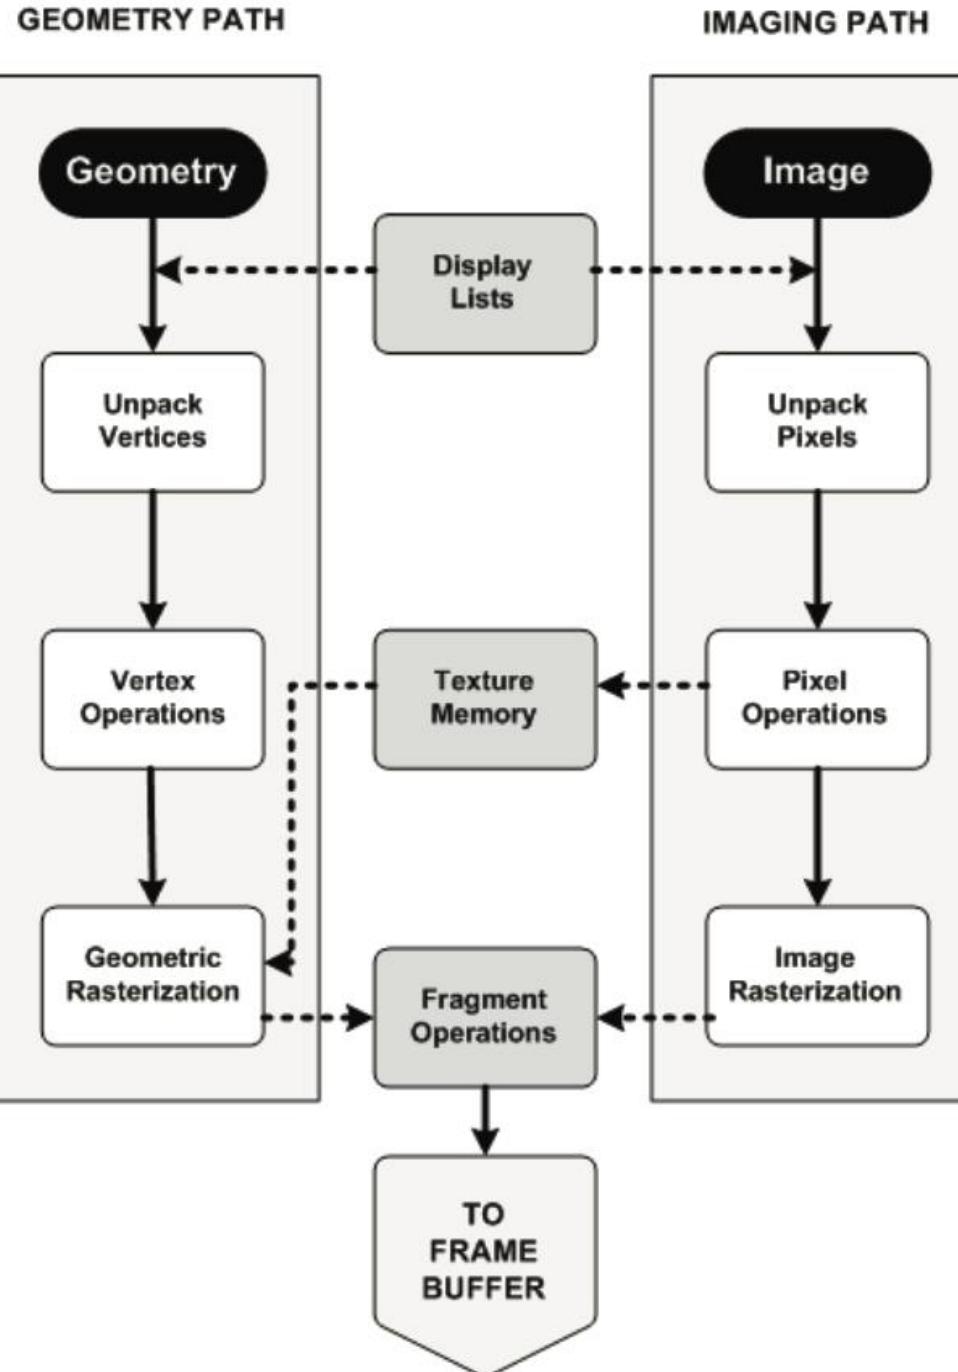
Config  
change

ODBC  
driver



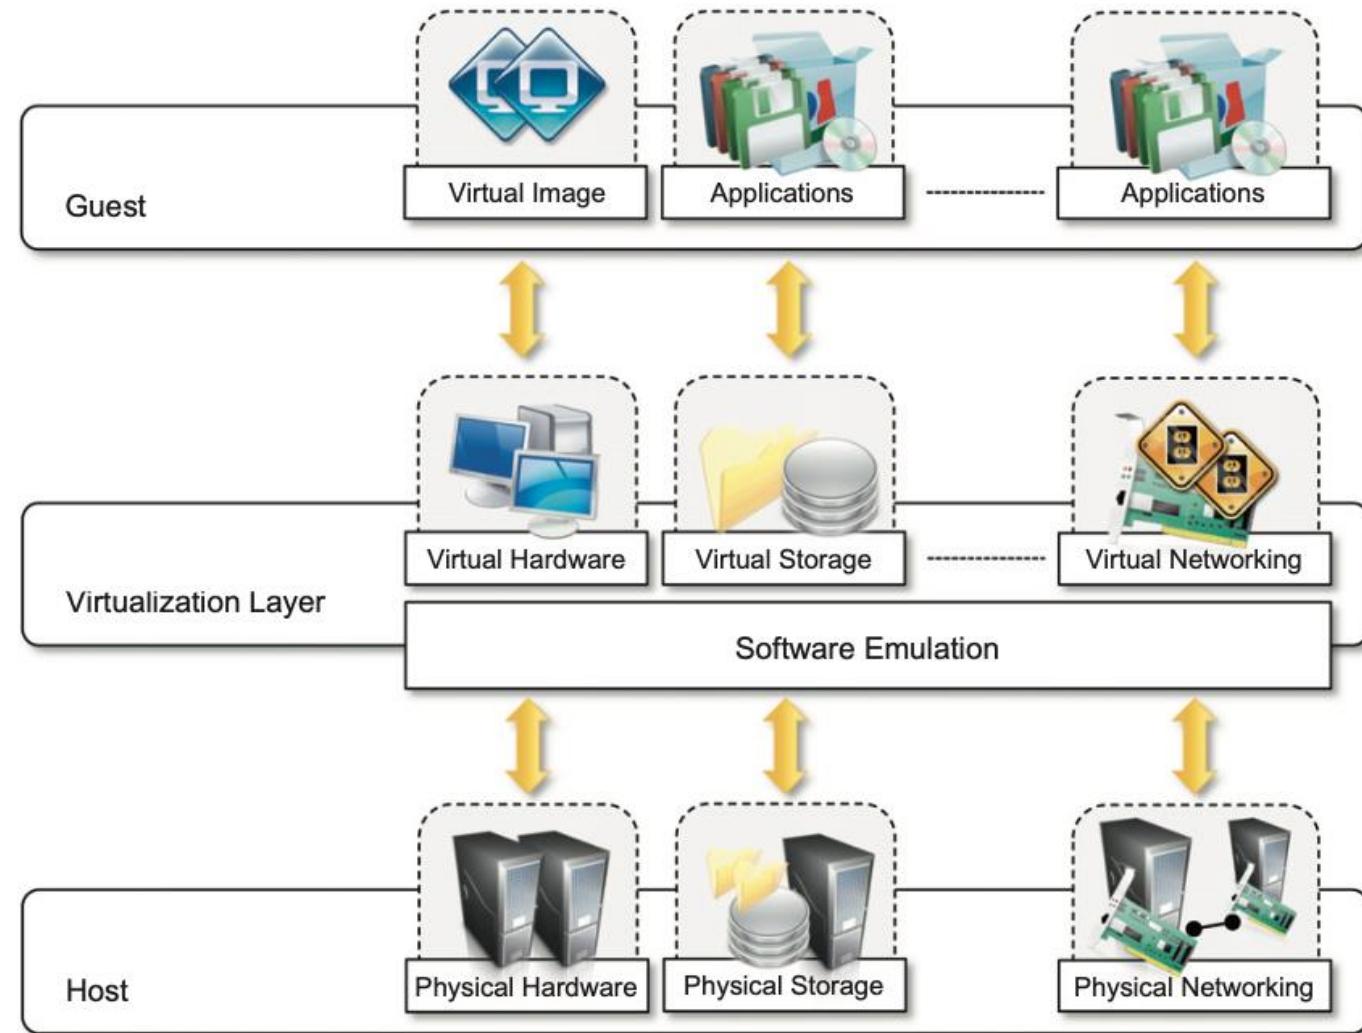
# Open Graphics Library

Application  
development  
and  
Display  
adapter



# Virtualization

- Abstraction of hardware, runtime environments, storage and networking
- Began in 1960s but widely adopted since 2000s
- Physical Host – virtual guest, 1 physical server hosting tens of virtual guests
- Virtual machine (VM) instances operate in complete isolation of each other
- Customized software stack on demand in each VM instance



# Factors

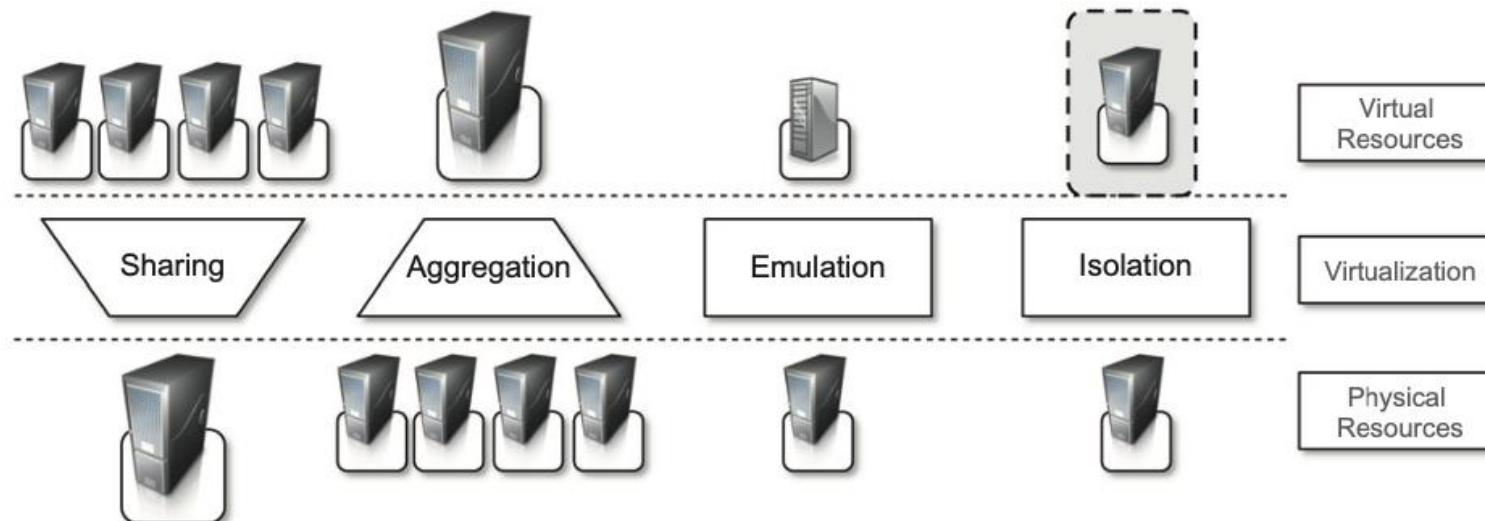
- Increased computing capacity: server can host tens of VMs
- Underutilized resources: using office desktops during night without affecting day work
- Lack of space – server consolidation
- Energy efficiency / Green initiatives
- Increased administrative costs
- 1995 – Sun Java, 2002 - .Net

# Characteristics

- hardware virtualization
- physical machine is the host
- virtualization layer / virtual machine manager / hypervisor – software
- virtual machines / guests interact with virtualization layer

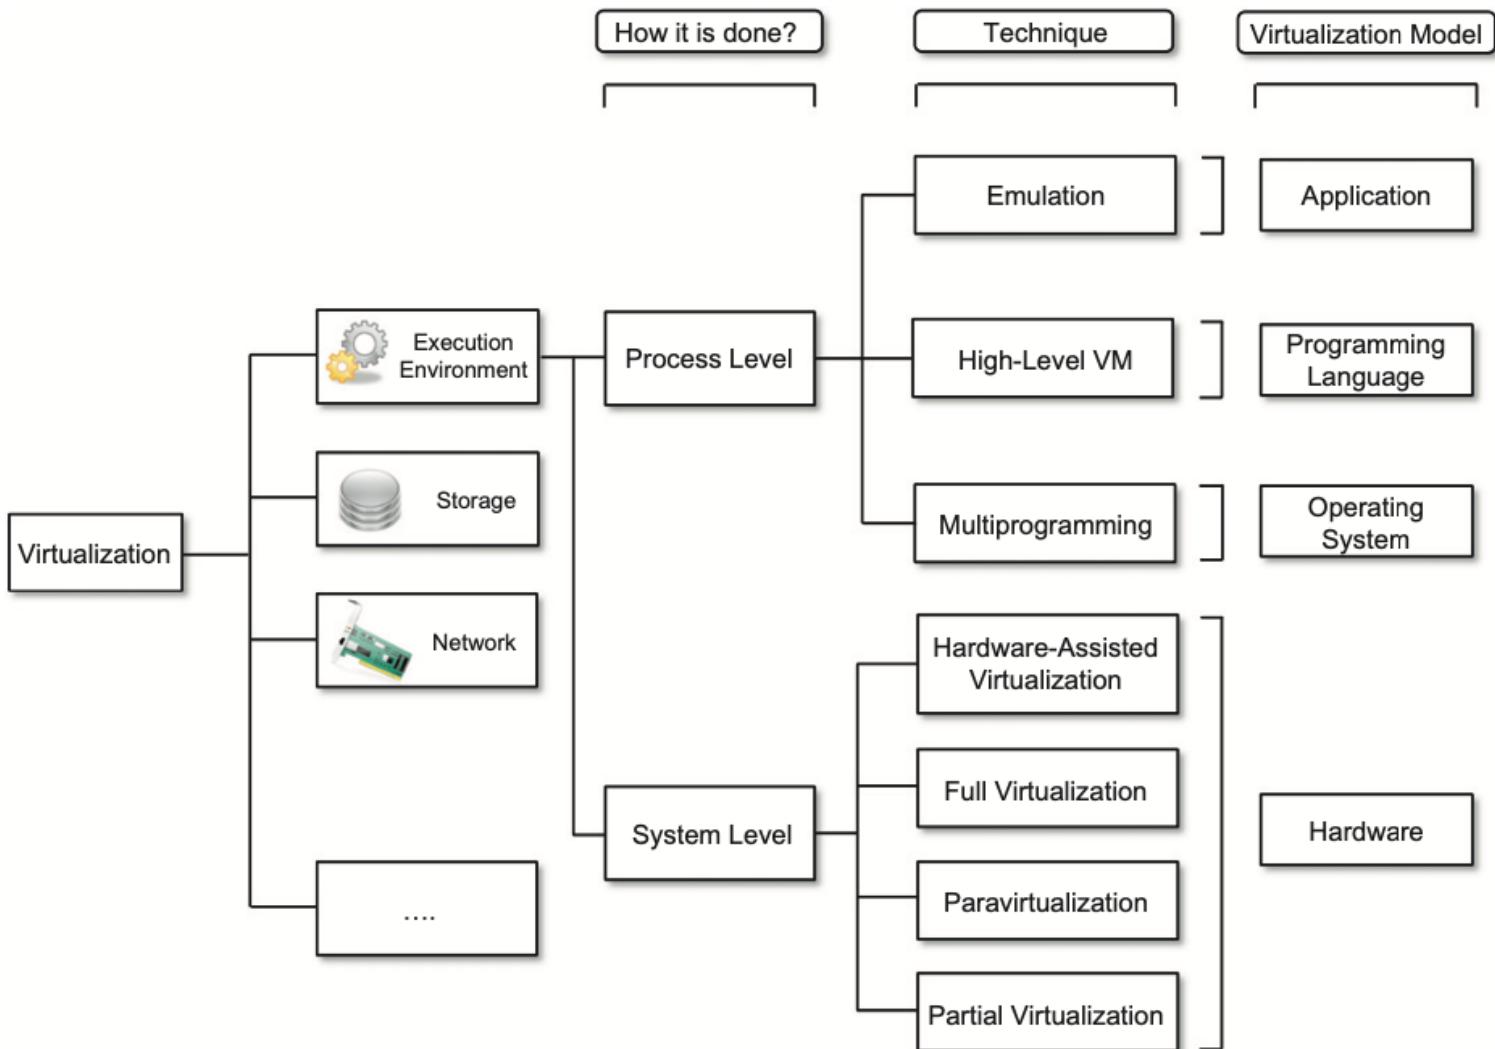
# Advantages

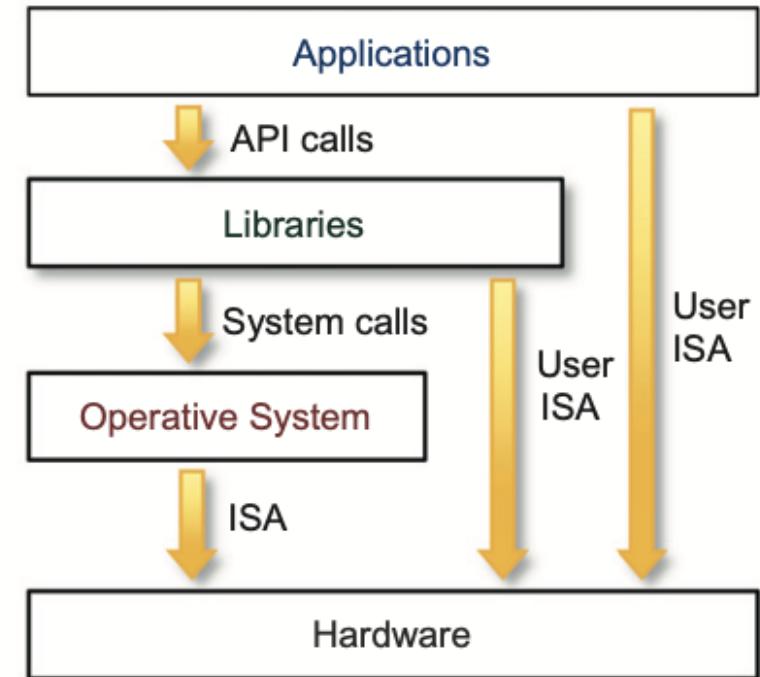
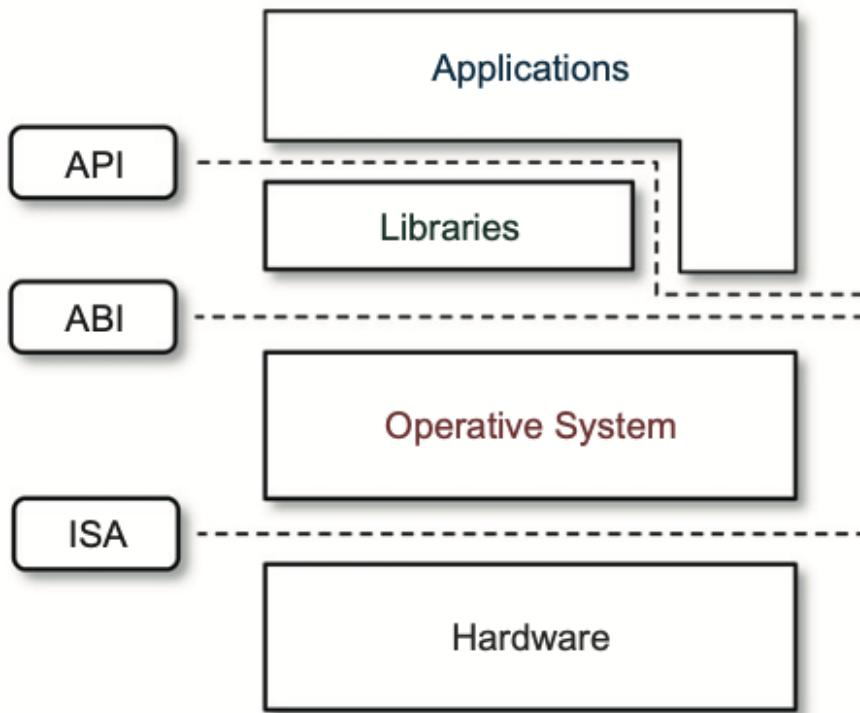
- Increased security – translation, control and filter guest activity, hide things on host from guest(s), limited access, separate FS
- Managed execution



# Advantages (contd)

- Performance tuning
- VM snapshot, migration
- Portability : JVM jar, .Net assembly, VMWare VMDK – images; don't need all applications and services but just VM manager





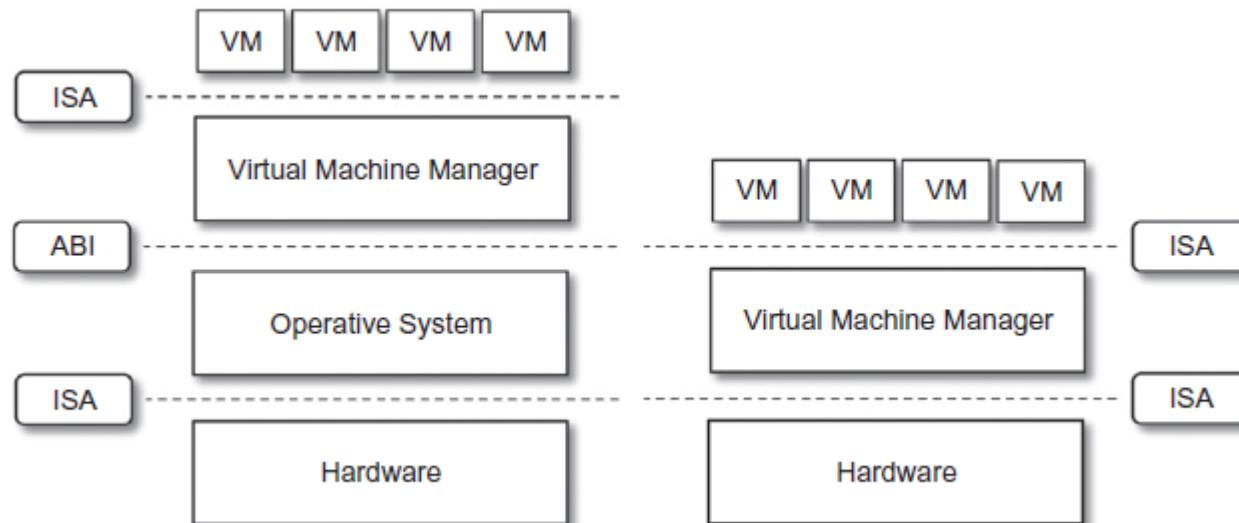
# Levels of virtualization

- ▶ Hardware level
- ▶ Operating system level
- ▶ Programming language level
- ▶ Application level
- ▶ Storage
- ▶ Network
- ▶ Desktop
- ▶ Application server

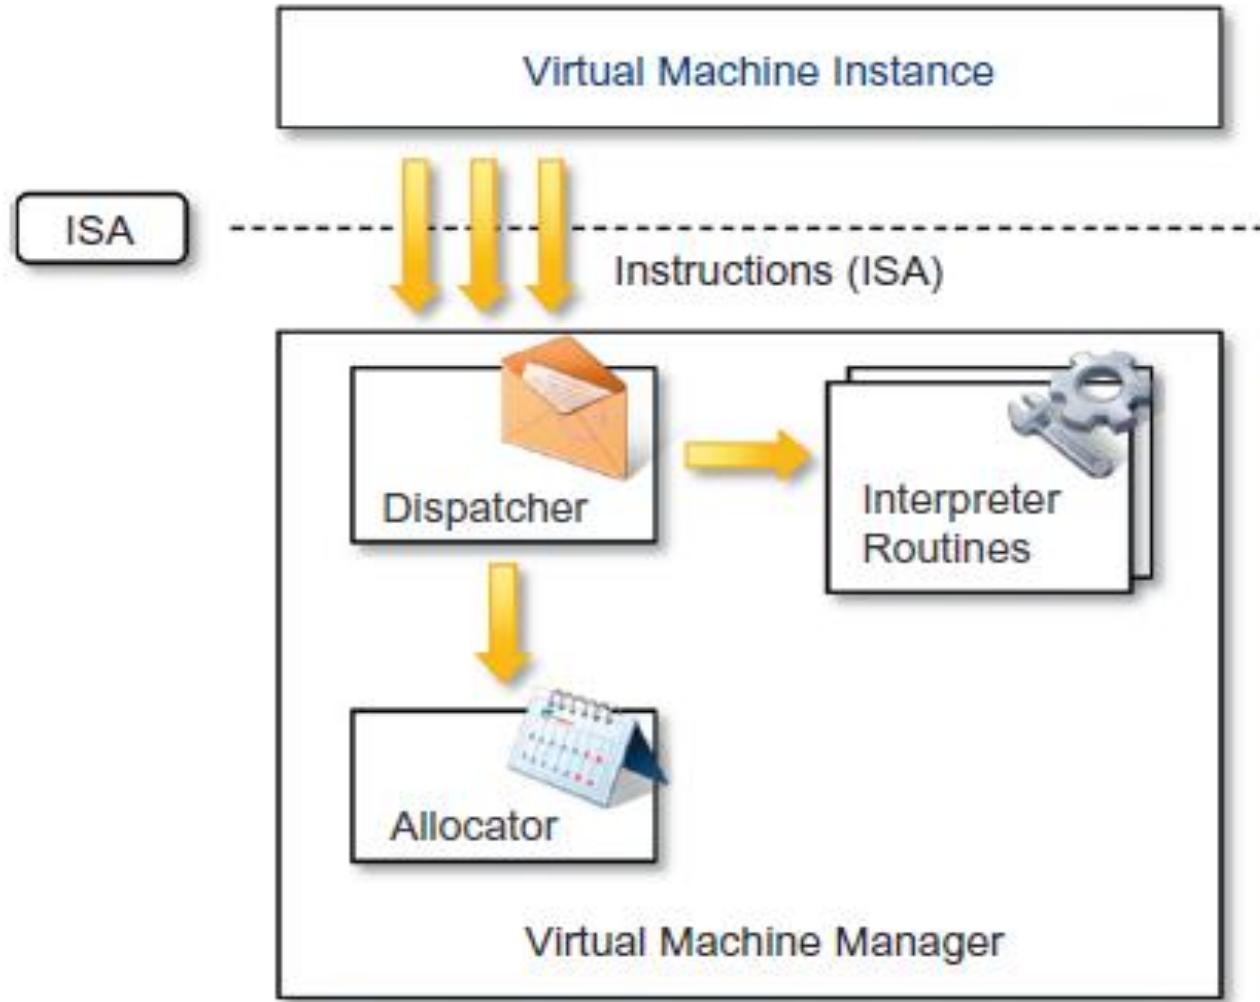
# Hardware level virtualization

## ► Hypervisor

- ▶ Type I : on bare hardware, like OS, interact directly with ISA, native
- ▶ Type II : require support of OS, interact with ABI, hosted



## Hypervisor Reference Architecture



# Virtual Machine Manager Criteria

- ▶ **Equivalence:** should be same as that of physical host
- ▶ **Resource control:** should be complete control of virtualized resource
- ▶ **Efficiency.** A statistically dominant fraction of the machine instructions should be executed without intervention from the VMM.

## Classification of instructions set

- ▶ Sensitive instructions are subset of privileged instructions
- ▶ Recursive virtualization: A VMM without any timing dependencies can be constructed for it
- ▶ In hybrid VMM user sensitive instructions are subset of privileged instructions

# Hardware Virtualization Techniques

- ▶ Hardware assisted : architectural support, Intel x86-64 VT-x, AMD V (2006), less overhead
- ▶ Full virtualization: running complete OS in VM like physical machine
- ▶ Paravirtualization: thin virtualization layer, guest needs to be modified to use VMM interface, need OS source, Linux on Xen
- ▶ Partial virtualization: partial execution of guest, Address space virtualization

# Levels of Virtualization

- ▶ OS Level: no VMM, within OS, evolution of Unix chroot, less overhead, solaris zones, containers (namespace)
- ▶ PL level: JVM, .Net assembly
- ▶ Application level: not installed on runtime hardware but run as if they were
  - ▶ Interpretation: every instruction is emulated, low startup but high overhead
  - ▶ Binary translation: block converted to native instructions, cached and reused
    - Wine: Unix like on Windows
    - WABI: Win16 on Solaris
    - CrossOver: Windows on Mac OS X
    - VMWare ThinApp

# Other than execution virtualization

- ▶ Storage Virtualization: decouple physical from logical, Storage Area Networks (SAN), NAS, enabled SDS
- ▶ Network Virtualization: VLAN, Network Address Translation (NAT) enabled SDN
- ▶ Desktop Virtualization: making the same desktop environment accessible from everywhere, VDI
- ▶ Application server virtualization: achieved using load balancers

# Enablers for Cloud Computing

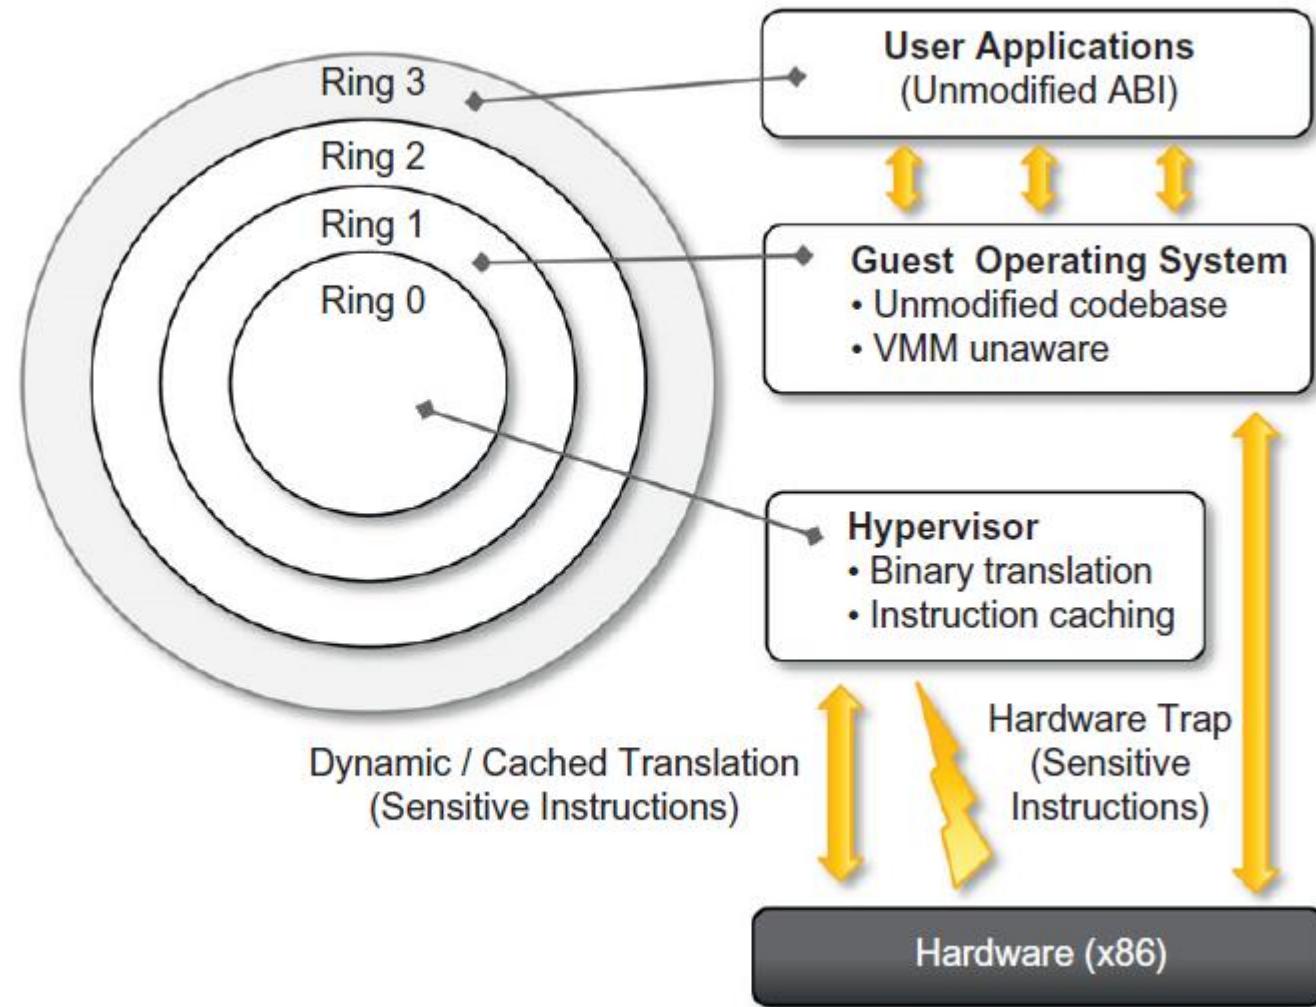
- ▶ Hardware level – IaaS
- ▶ PL level – PaaS
- ▶ Migrate VMs to free up servers, server consolidation
- ▶ Storage services
- ▶ Network services
- ▶ VDI: your entire desktop can come from cloud

# Advantages & Disadvantages

- ▶ Sandbox, Prevent any harmful operation cross the boundary
- ▶ Portable across physical systems, self contained without any dependencies, represented by one or two files
- ▶ Low maintenance
  
- ▶ Performance degradation
- ▶ Inefficiency and degraded user experience: inaccessible features
- ▶ Security holes, threats, phishing: extract sensitive information from guest, BluePil, SubVirt

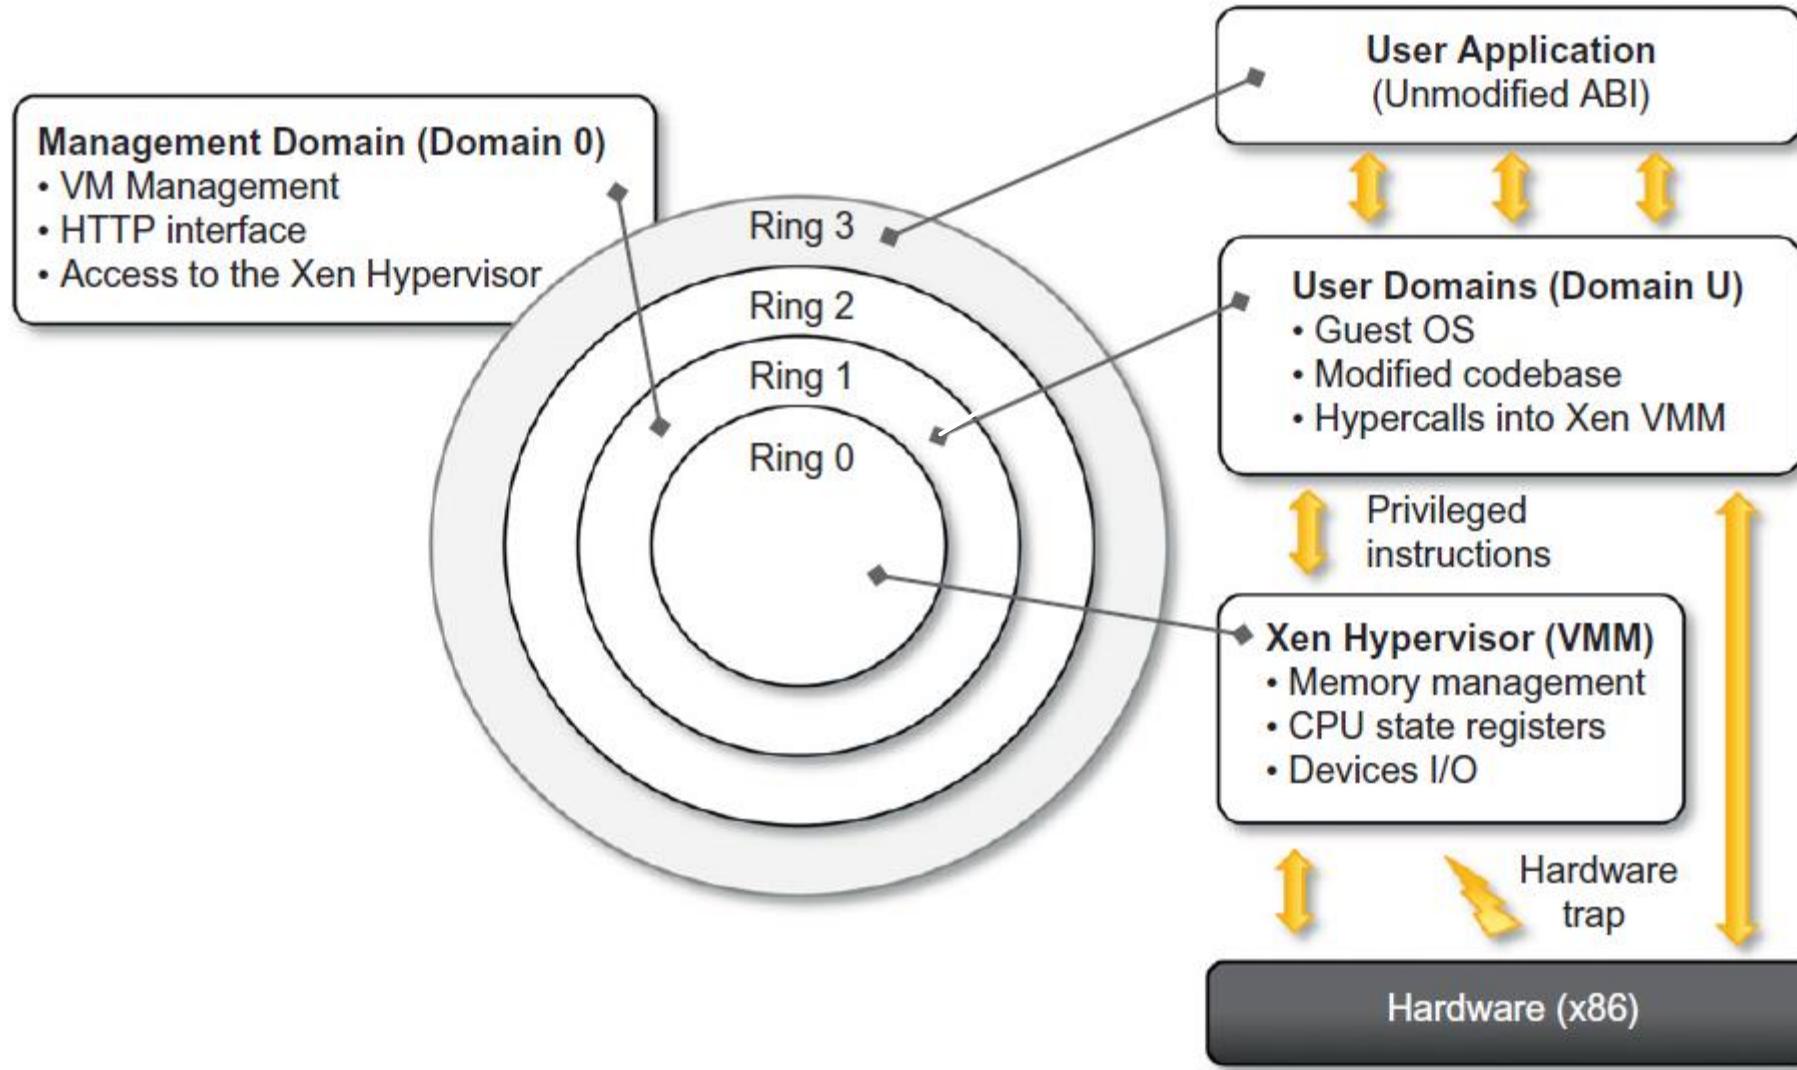
# VMWare: full virtualization

- ▶ Hardware replicated, guest not modified and not aware of abstraction layer in between, desktop type II, server type I, direct execution for non sensitive instructions, binary translation for sensitive instructions, ecosystem
- ▶ X86 architecture does not meet theorem I, dynamic binary translation, trap is generated for privileged instructions and translation performed, caching of translation to improve performance
- ▶ CPU, memory and I/O devices (keyboard, mouse, USB, network controllers) virtualization
- ▶ Translation Look-aside Buffer (TLB) to improve MMU performance

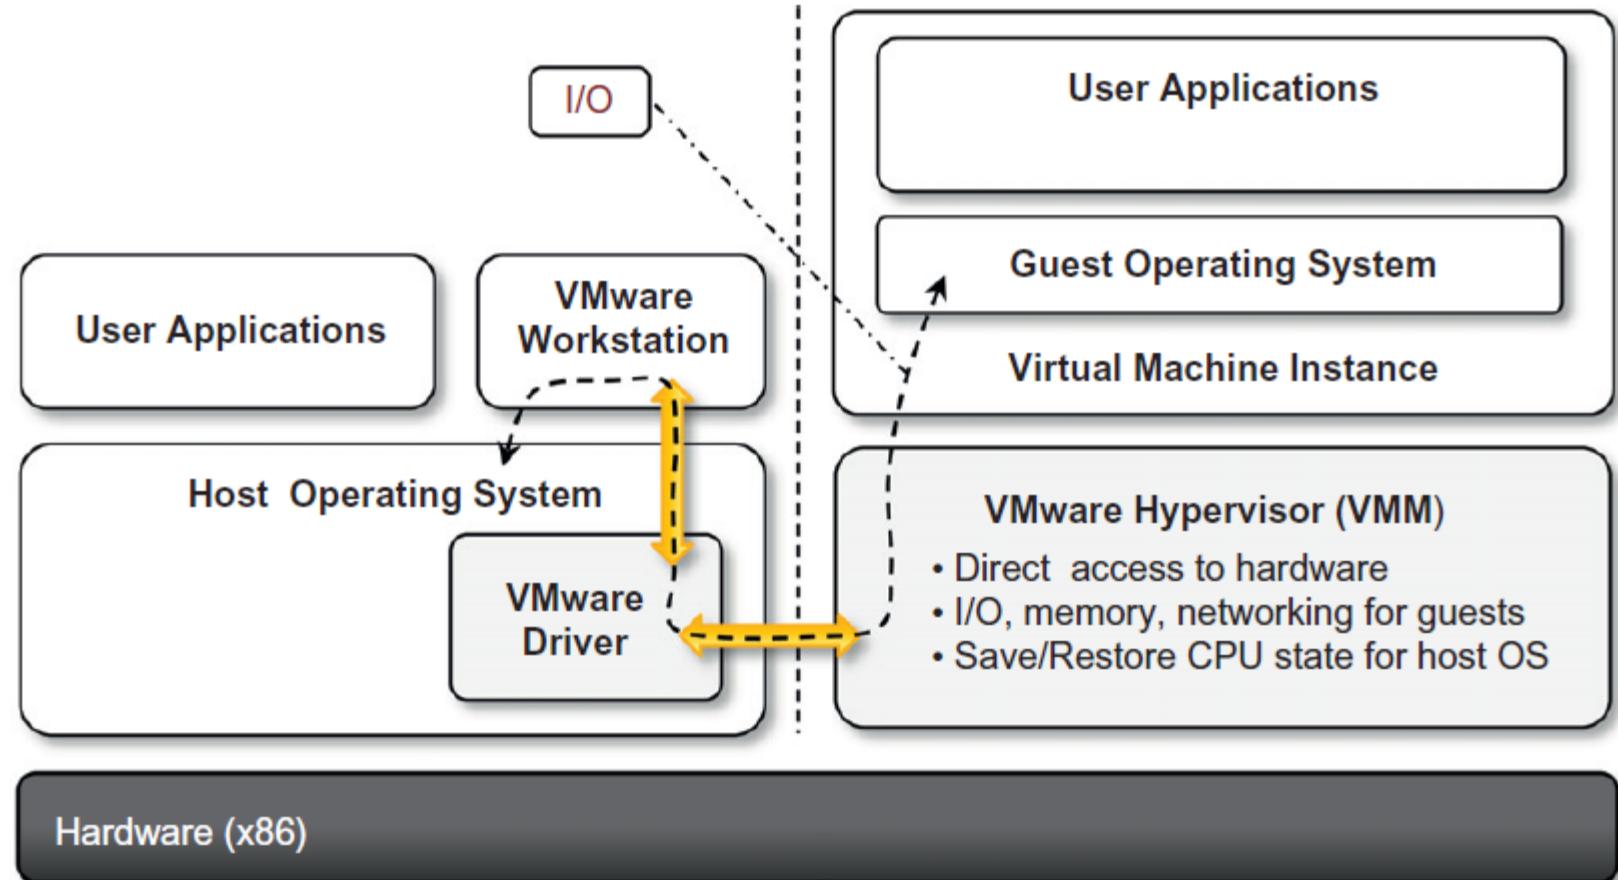


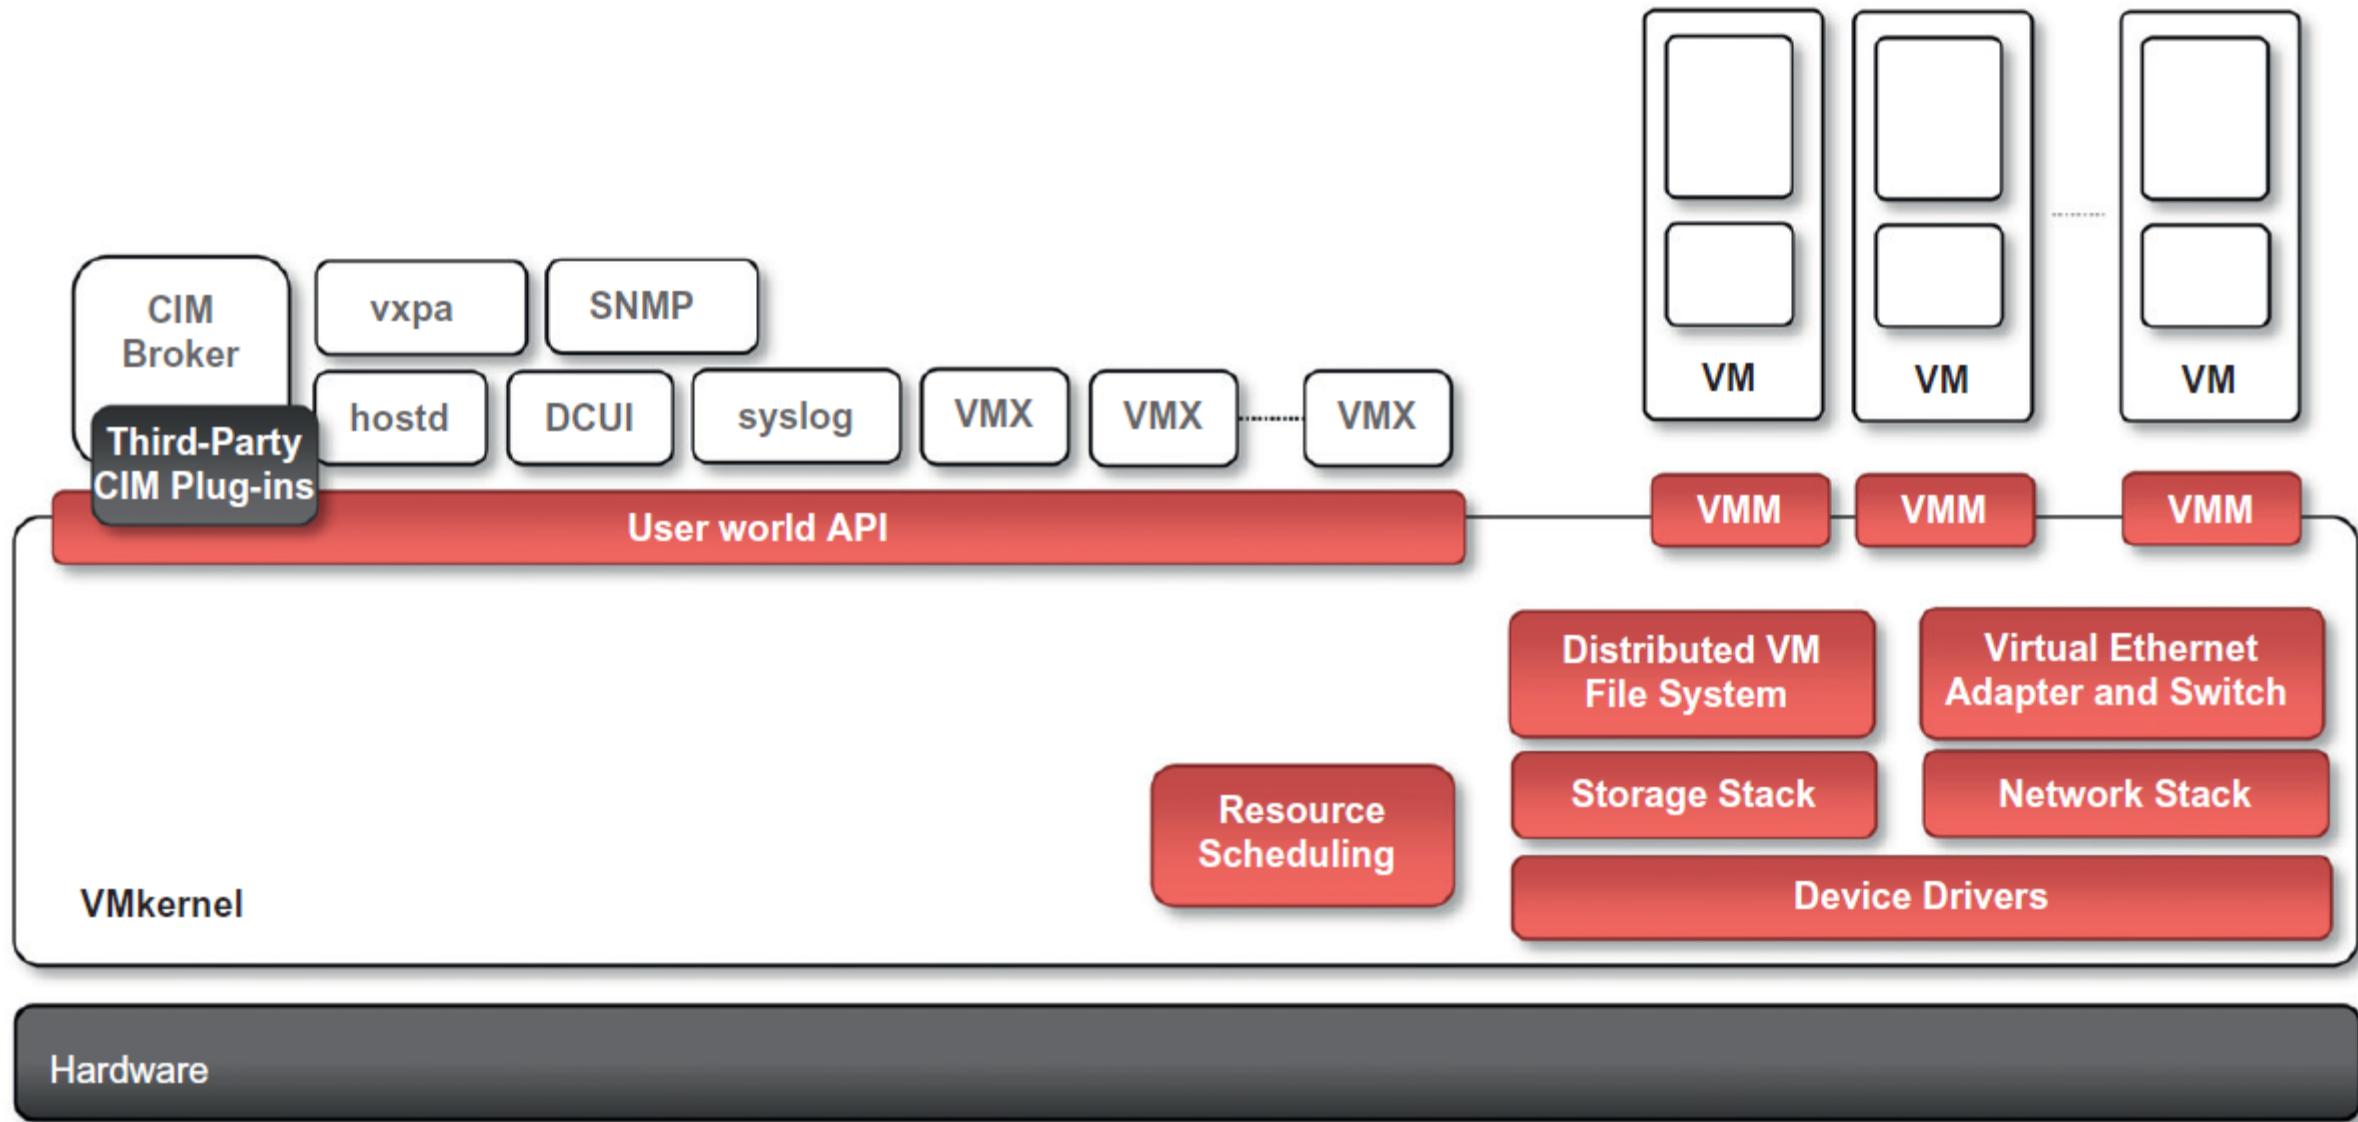
# Xen: paravirtualization

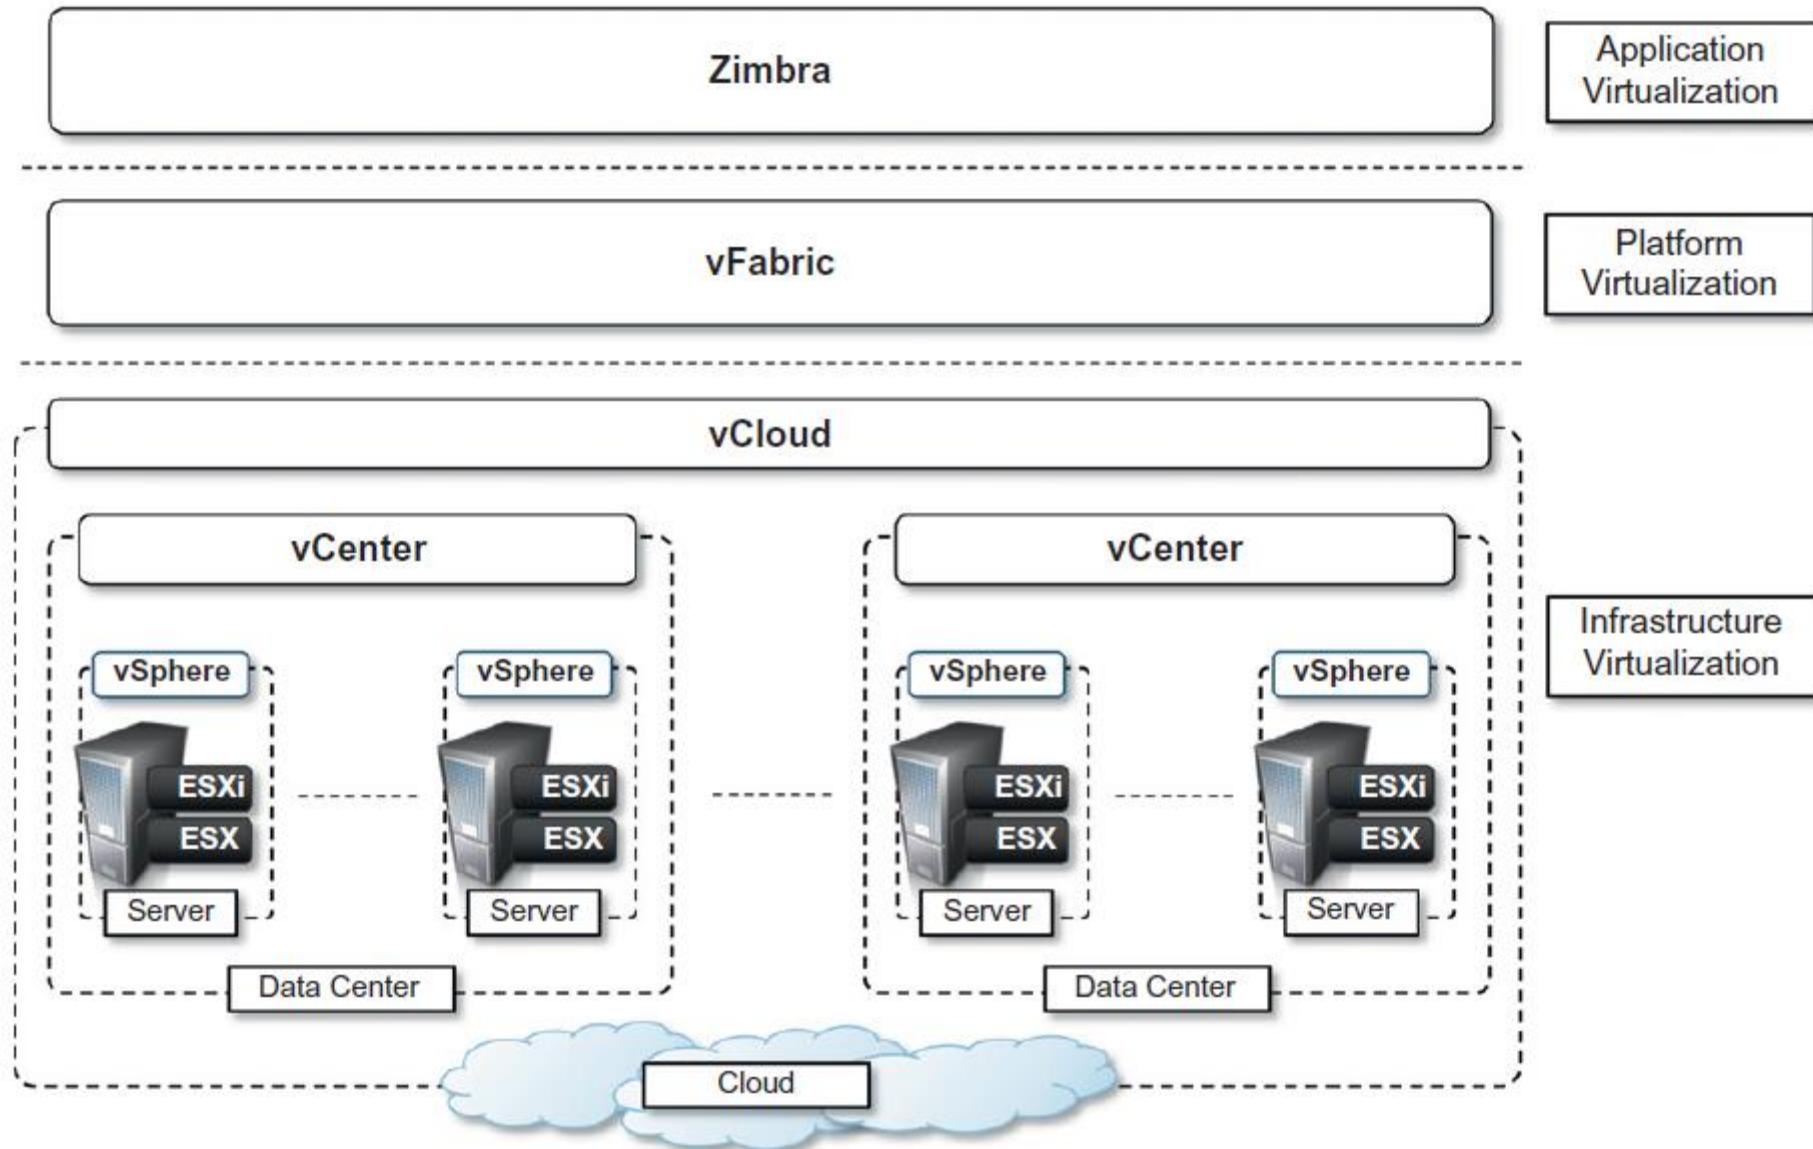
- ▶ Research by Cambridge University, UK
- ▶ Open source community
- ▶ Citix offers commercial offering of XenSource
- ▶ Desktop or server virtualization
- ▶ Xen Cloud Platform XCP
- ▶ Full virtualization with hardware assistance
- ▶ Xen Hypervisor, domains, hypercalls

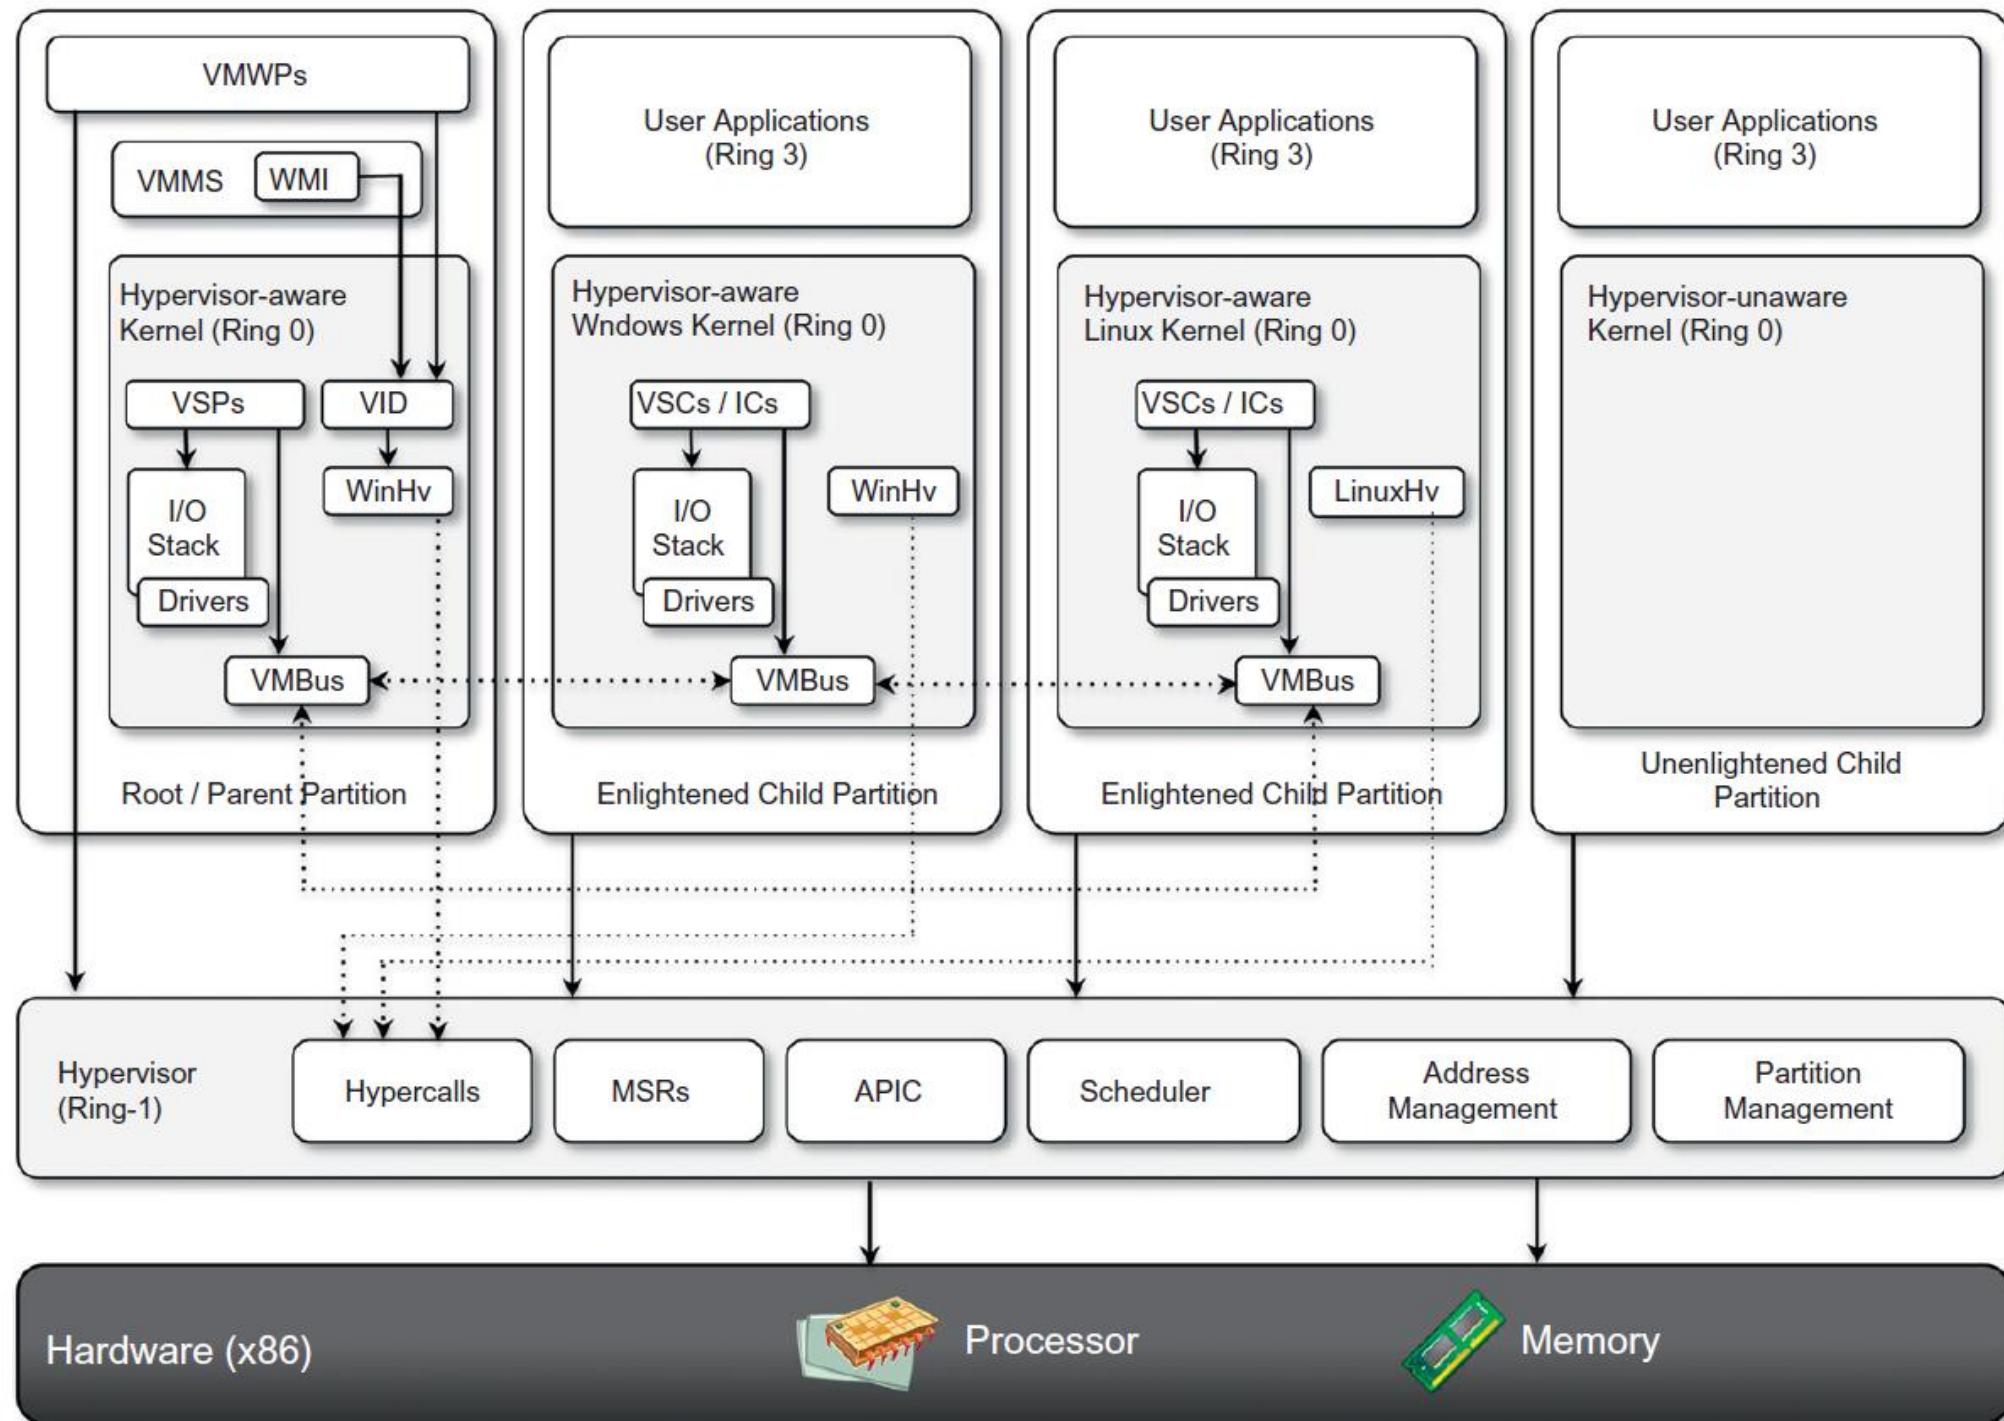


## End user VMWare Workstation Architecture









# Releases

- ▶ Hyper-V currently shipped as a component of Windows Server 2008 R2 that installs the hypervisor as a role within the server
- ▶ Windows Server Core to increase performance with smaller footprint, reduced maintenance, reduced management, use of “removed” rich features by remote management through WMI (e.g. powershell)
- ▶ SystemCenter Virtual MachineManager(SCVMM)2008
  - ▶ Management portal
  - ▶ V2V, P2V conversions
  - ▶ Deep powershell integration

# Assignment

- ▶ Install VMWare ESXi server
- ▶ Create Ubuntu VM
- ▶ Connect to Ubuntu VM

# Cloud Computing Architecture

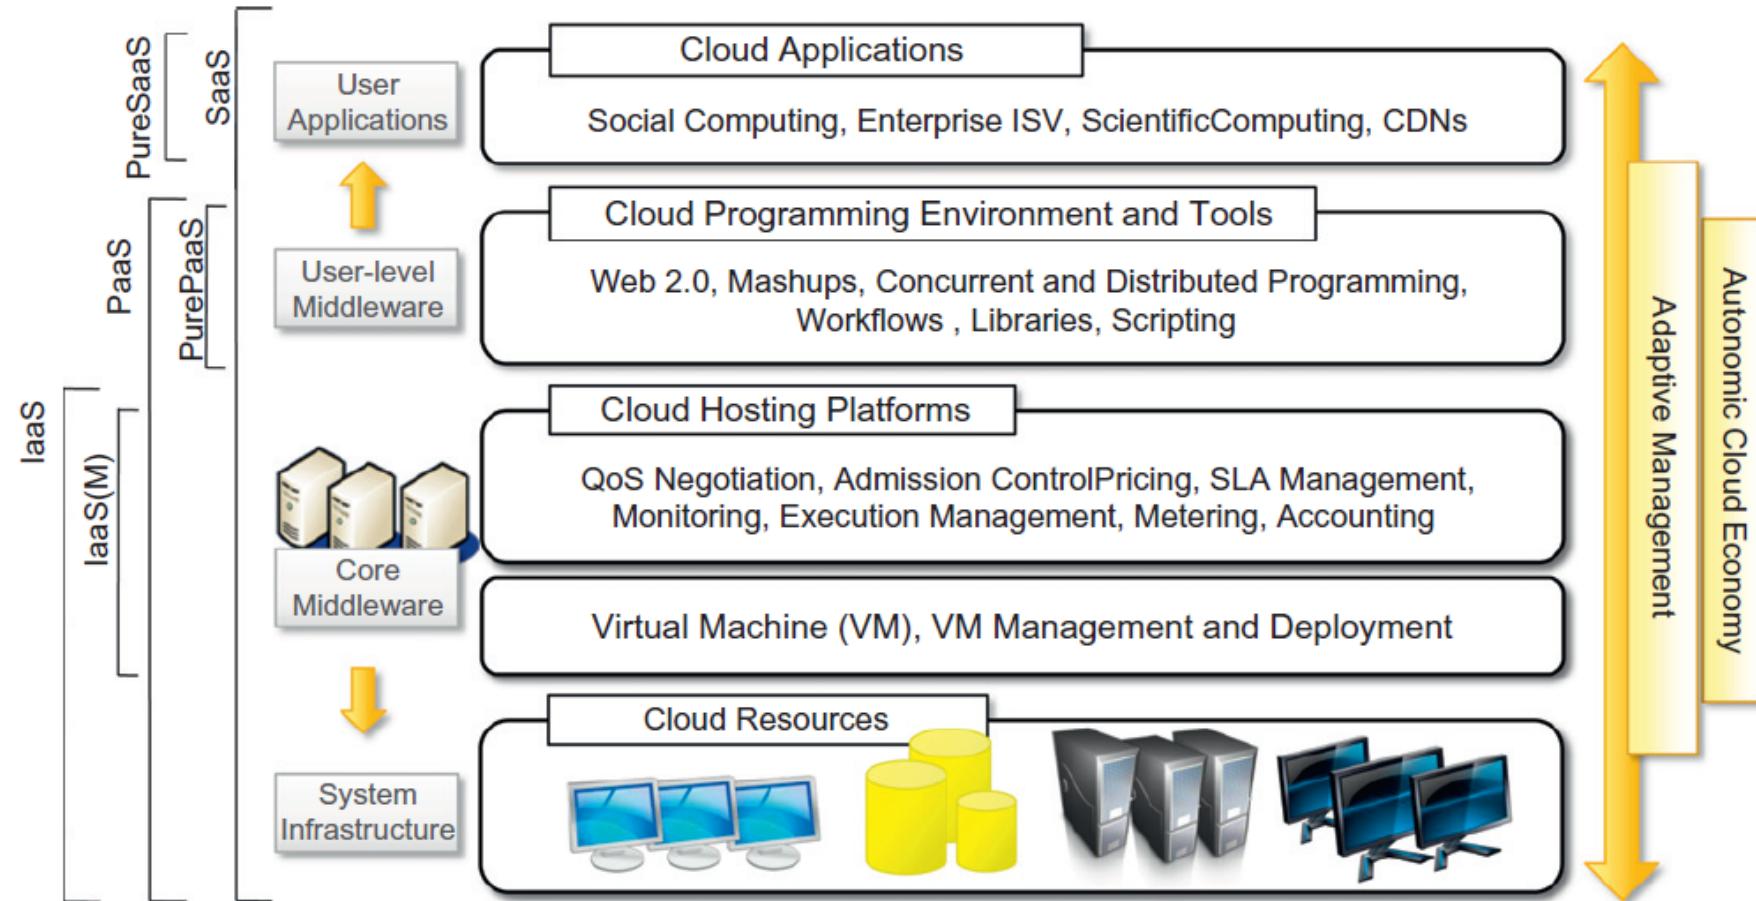
SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# Cloud Computing Architecture

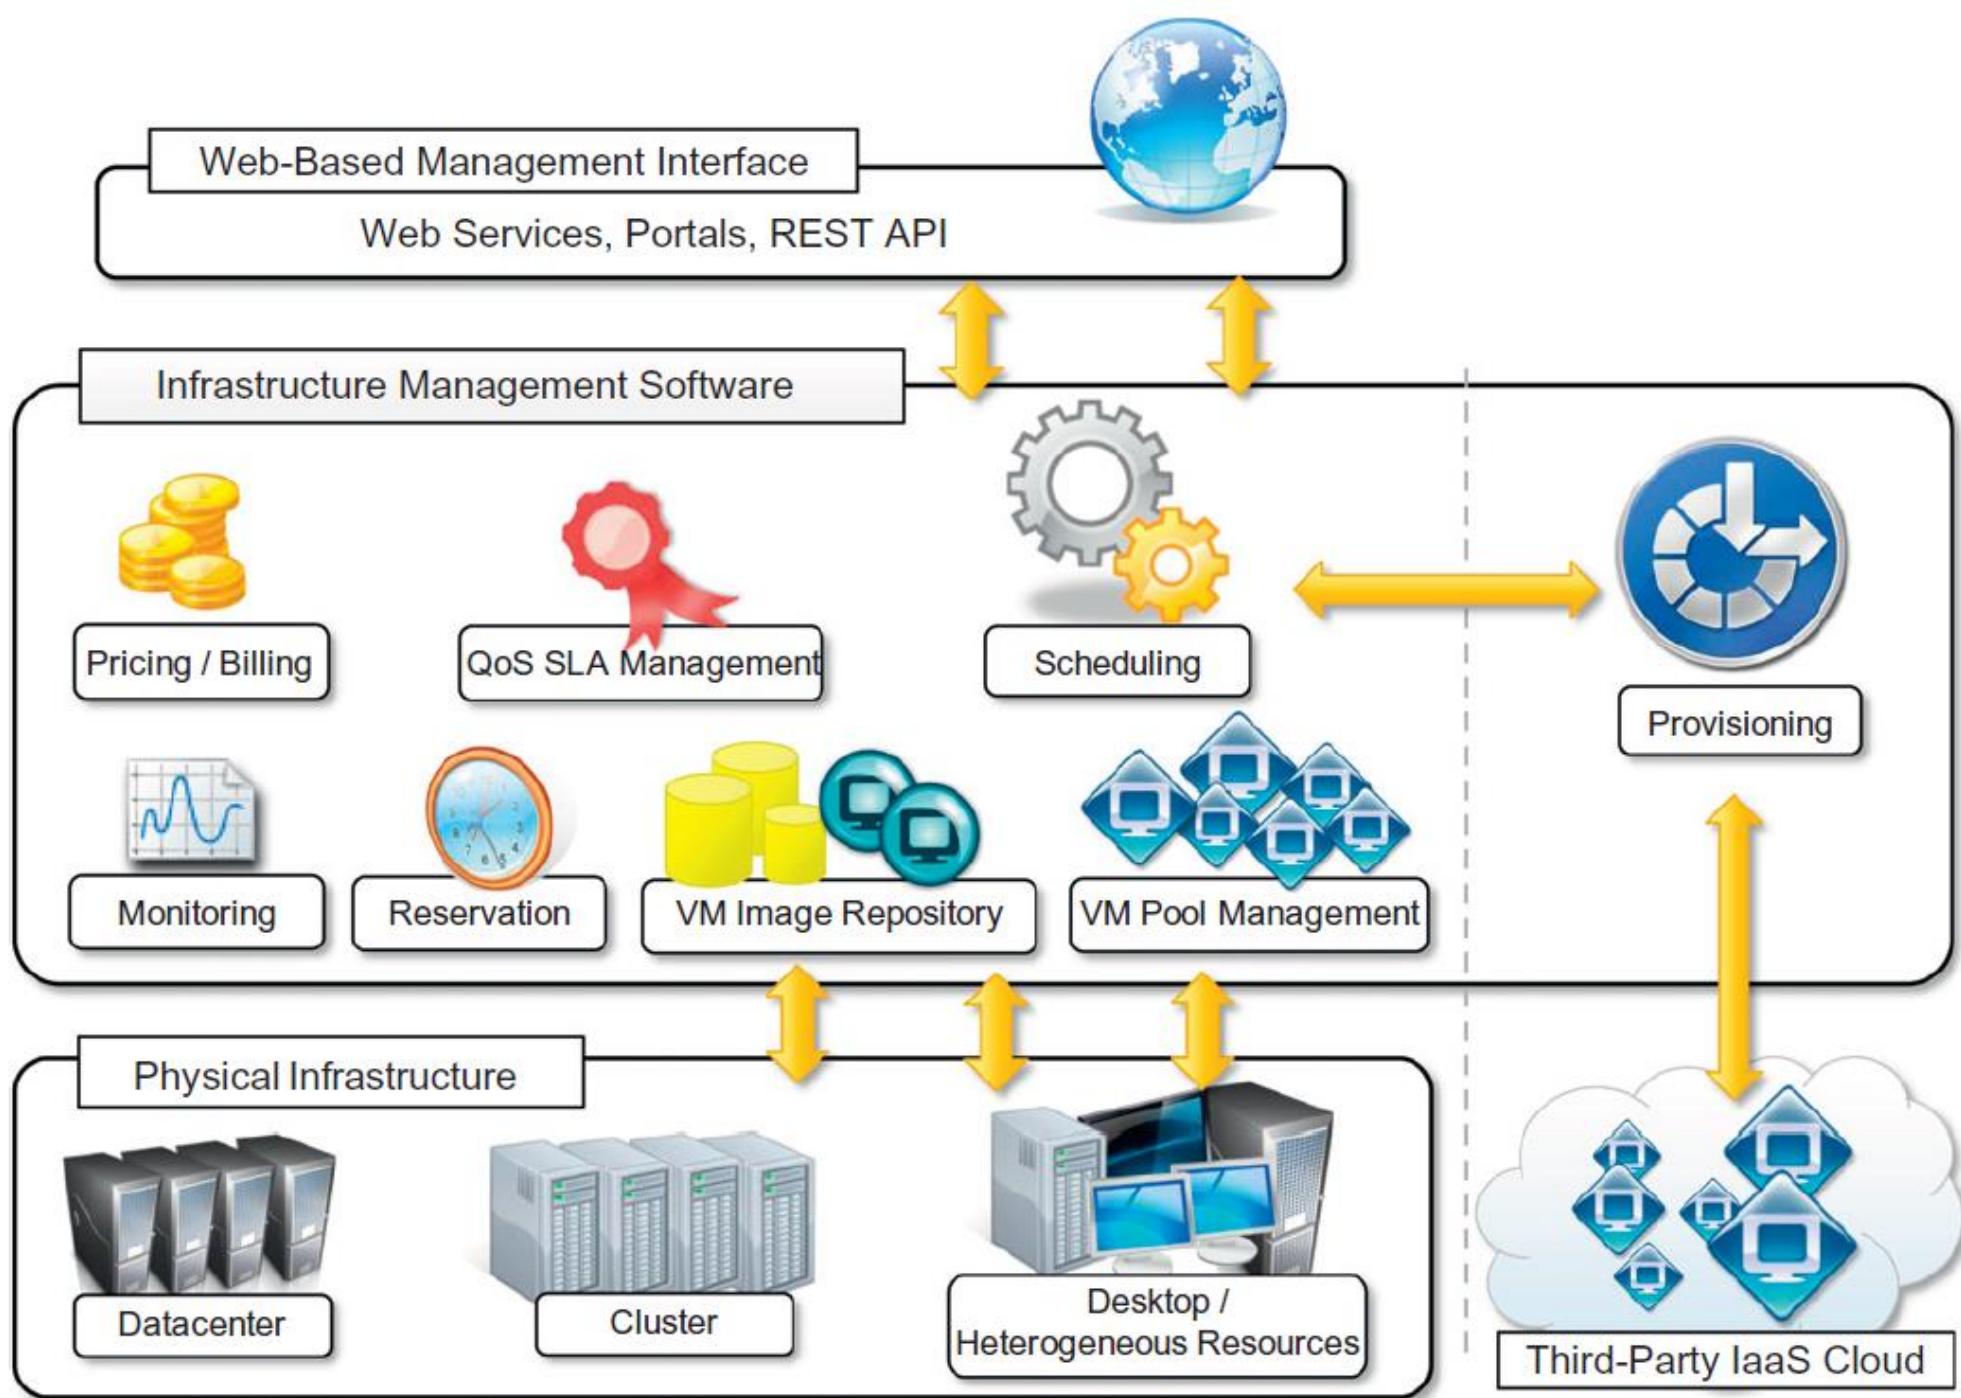
- ▶ Distributed infrastructure
- ▶ Data center(s), cluster(s), composed of desktops, workstations, servers
- ▶ Infrastructure can be owned by provider or rented from third party
- ▶ Typically virtualized for workload isolation and cost optimization
- ▶ Different layers stacked on each other
- ▶ Utility oriented, internet centric, IT services on demand, covering the entire computing stack from hardware infrastructure to development platforms and distributed applications

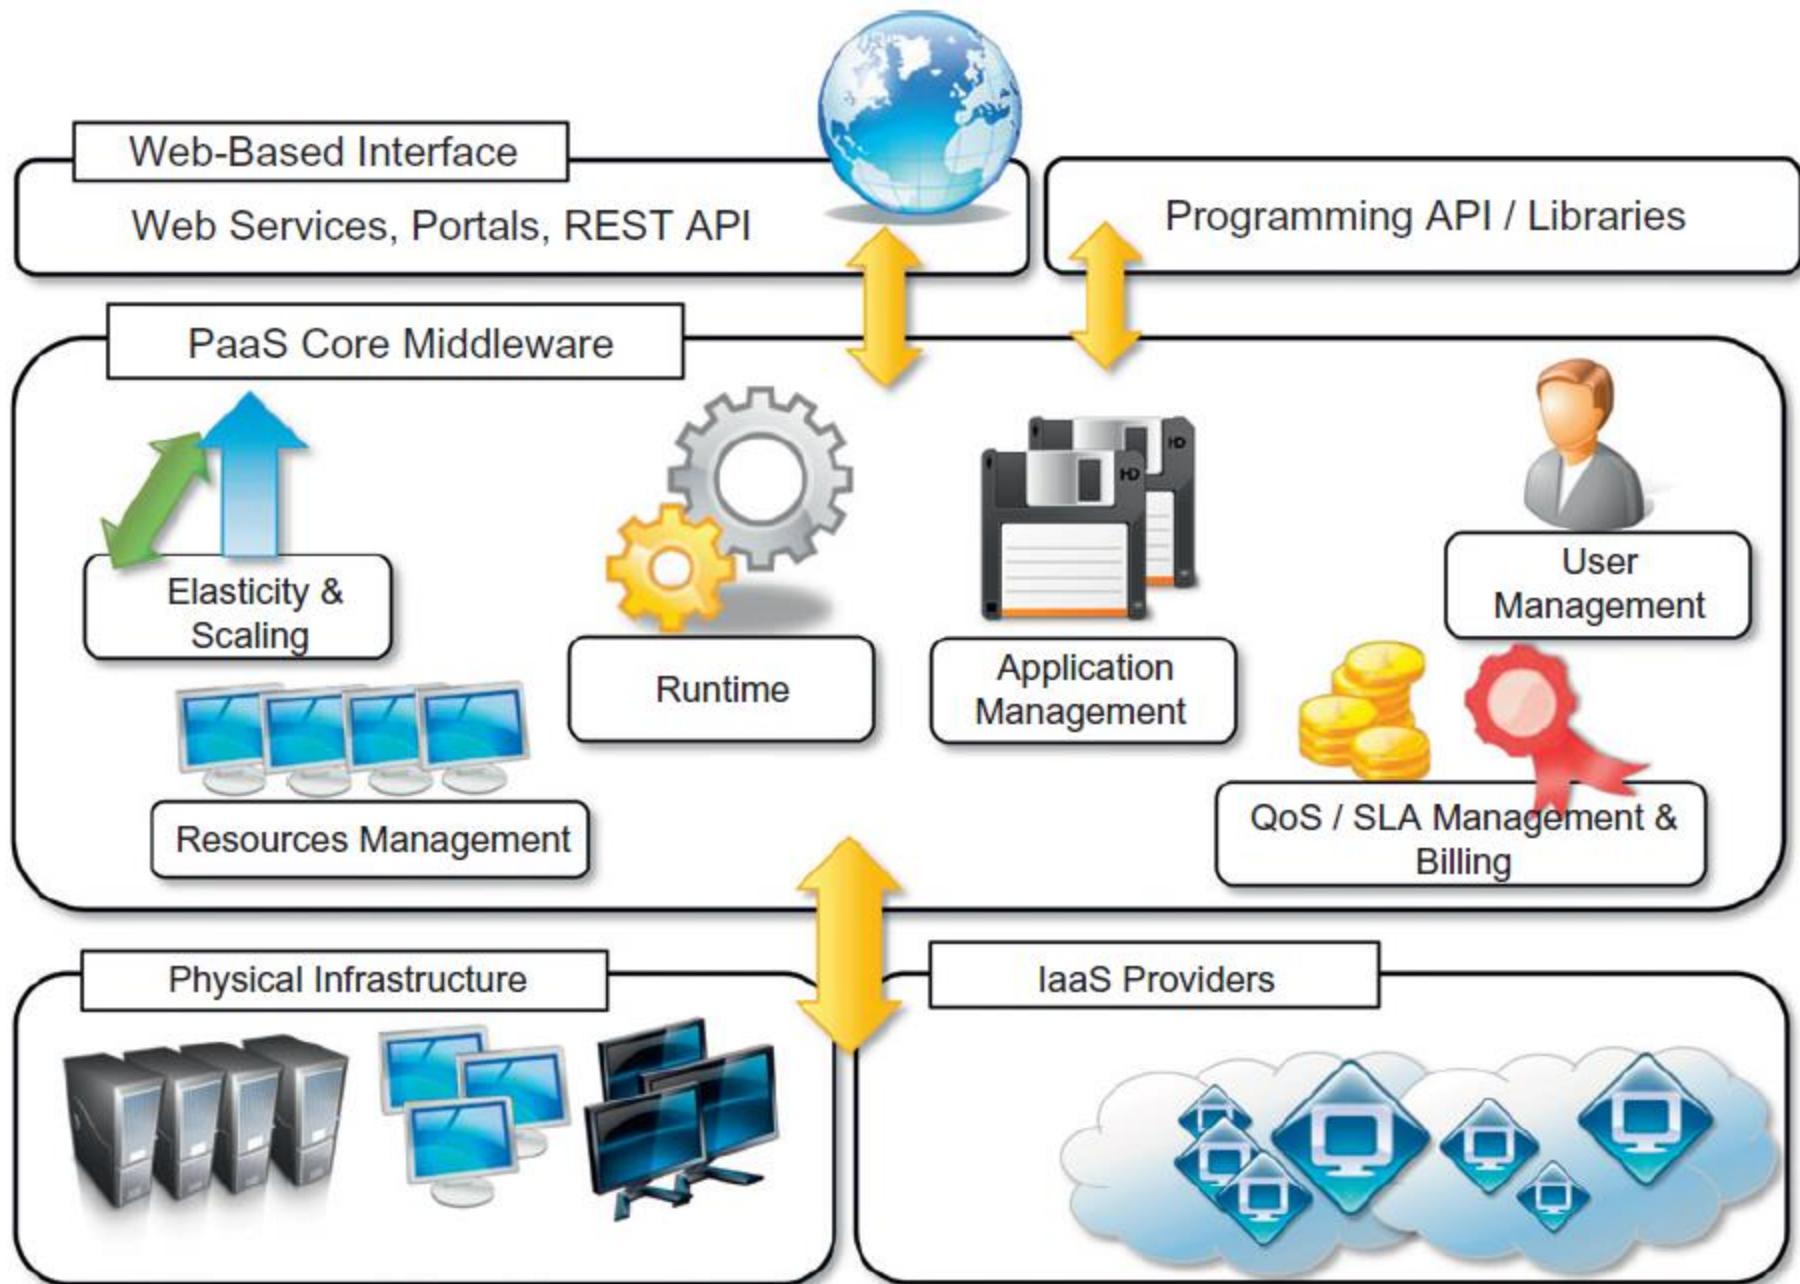
# Reference Model



# IaaS

- ▶ Infrastructure or hardware as a service
- ▶ Most popular and developed market segment of cloud computing
- ▶ from single servers to entire infra- structures, including network devices, load balancers, and database and Web servers
- ▶ Main technology used: hardware virtualization
- ▶ Atomic component of VM priced by memory, #processors, and disk size
- ▶ workload partitioning, application isolation, sandboxing, and hardware tuning
- ▶ LAMP: Linux Apache MySql and PHP





**Table 4.2** Platform-as-a-Service Offering Classification

Category	Description	Product Type	Vendors and Products
PaaS-I	Runtime environment with Web-hosted application development platform. Rapid application prototyping.	Middleware + Infrastructure	<a href="#">Force.com</a>
		Middleware + Infrastructure	Longjump
PaaS-II	Runtime environment for scaling Web applications. The runtime could be enhanced by additional components that provide scaling capabilities.	Middleware + Infrastructure	<a href="#">Google AppEngine</a>
		Middleware	<a href="#">AppScale</a>
		Middleware + Infrastructure	<a href="#">Heroku</a>
		Middleware + Infrastructure	<a href="#">Engine Yard</a>
		Middleware + Infrastructure	<a href="#">Joyent Smart Platform</a>
		Middleware	<a href="#">GigaSpaces XAP</a>
PaaS-III	Middleware and programming model for developing distributed applications in the cloud.	Middleware + Infrastructure	<a href="#">Microsoft Azure</a>
		Middleware	<a href="#">DataSynapse</a>
		Middleware	<a href="#">Cloud IQ</a>
		Middleware	<a href="#">Manjrasof Aneka</a>
		Middleware	<a href="#">Apprenda</a>
		Middleware	<a href="#">SaaSGrid</a>
			<a href="#">GigaSpaces</a>
			<a href="#">DataGrid</a>

# Essential Characteristics of PaaS

- ▶ Runtime framework
- ▶ Higher level of abstraction
- ▶ Automating deployment and scaling
- ▶ Cloud Services and APIs to simplify the creation and delivery of elastic and highly available cloud applications
- ▶ Integration of third party cloud services leveraging SOA
- ▶ Vendor lockin
- ▶ Optimize development, deployment and management of applicaitons

# SaaS

- ▶ Complete applications as a service
- ▶ No hardware, software cost, pay as you go subscriptions
- ▶ One-to-many, hosted, some customization
- ▶ Multitenant
- ▶ SaaS coined in 2001, predates cloud computing, Application Service Providers (ASP)
- ▶ SaaS 2.0 rapid achievement of business objectives, integration, interconnected
- ▶ Salesforce (CRM), Facebook (social networking), linkedin, Google Documents

# Types

Clouds are a type of parallel and distributed system harnessing physical and virtual computers presented as a unified computing resource

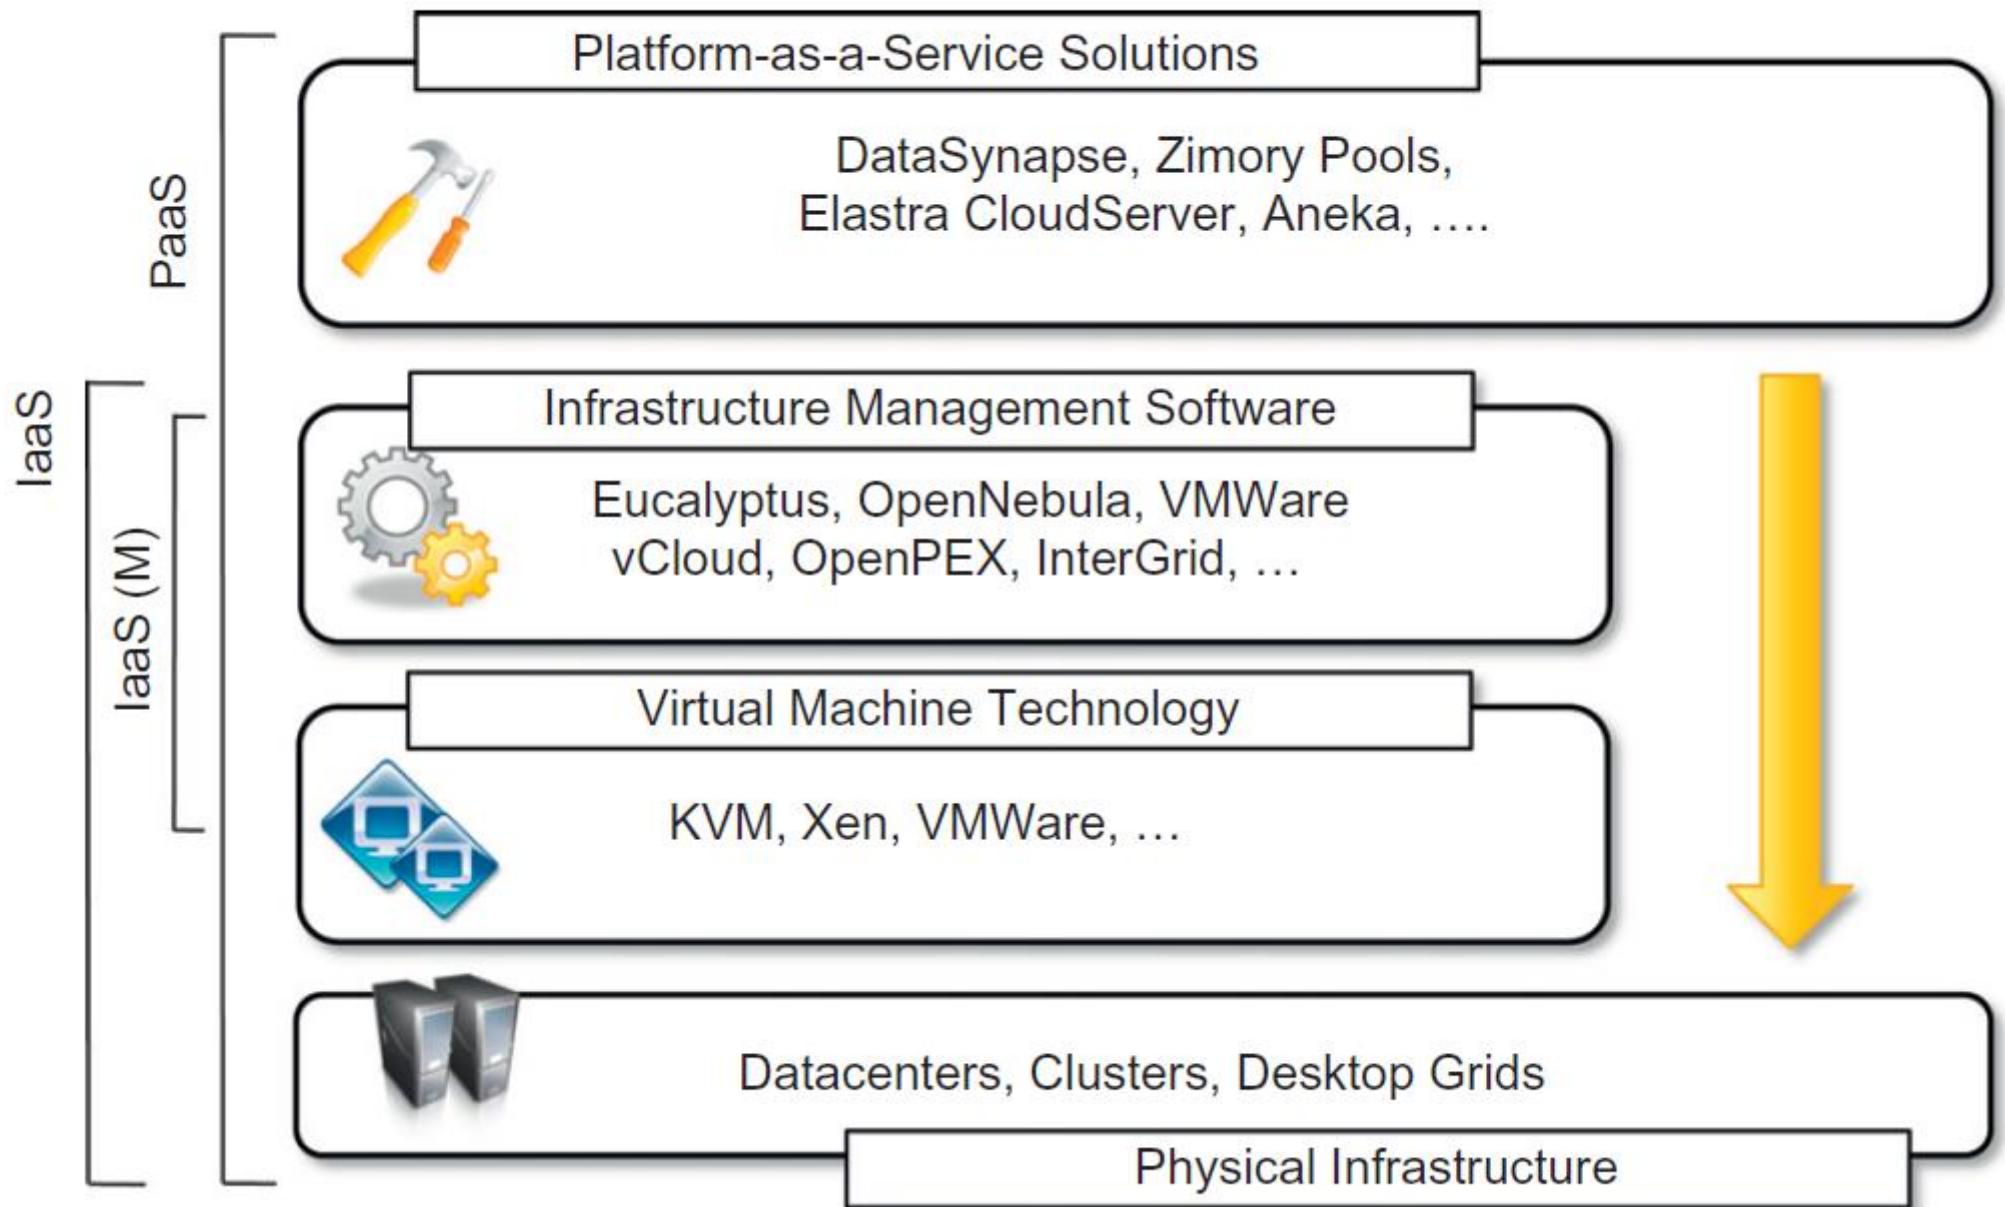
- ▶ Public: open to the wider public
- ▶ Private: limited to an institution, intranet
- ▶ Hybrid or heterogeneous: combination of the above two
- ▶ Community: for a specific industry
- ▶ Multi: combination of on-premises, private, public, and edge

# Public Cloud

- ▶ First expression of cloud computing
- ▶ Anyone, anywhere, any time via Internet
- ▶ Useful to start small enterprises without incurring any capex
- ▶ Grow or shrink according to business needs, sustain peak load
- ▶ Renting infrastructure or subscribing to services
- ▶ Replace or extend
- ▶ Multitenancy, QoS
- ▶ Monitoring: uptime, usage (for billing), activity
- ▶ AWS EC2: IaaS, Google AppEngine: PaaS, salesforce.com: SaaS

# Private Cloud

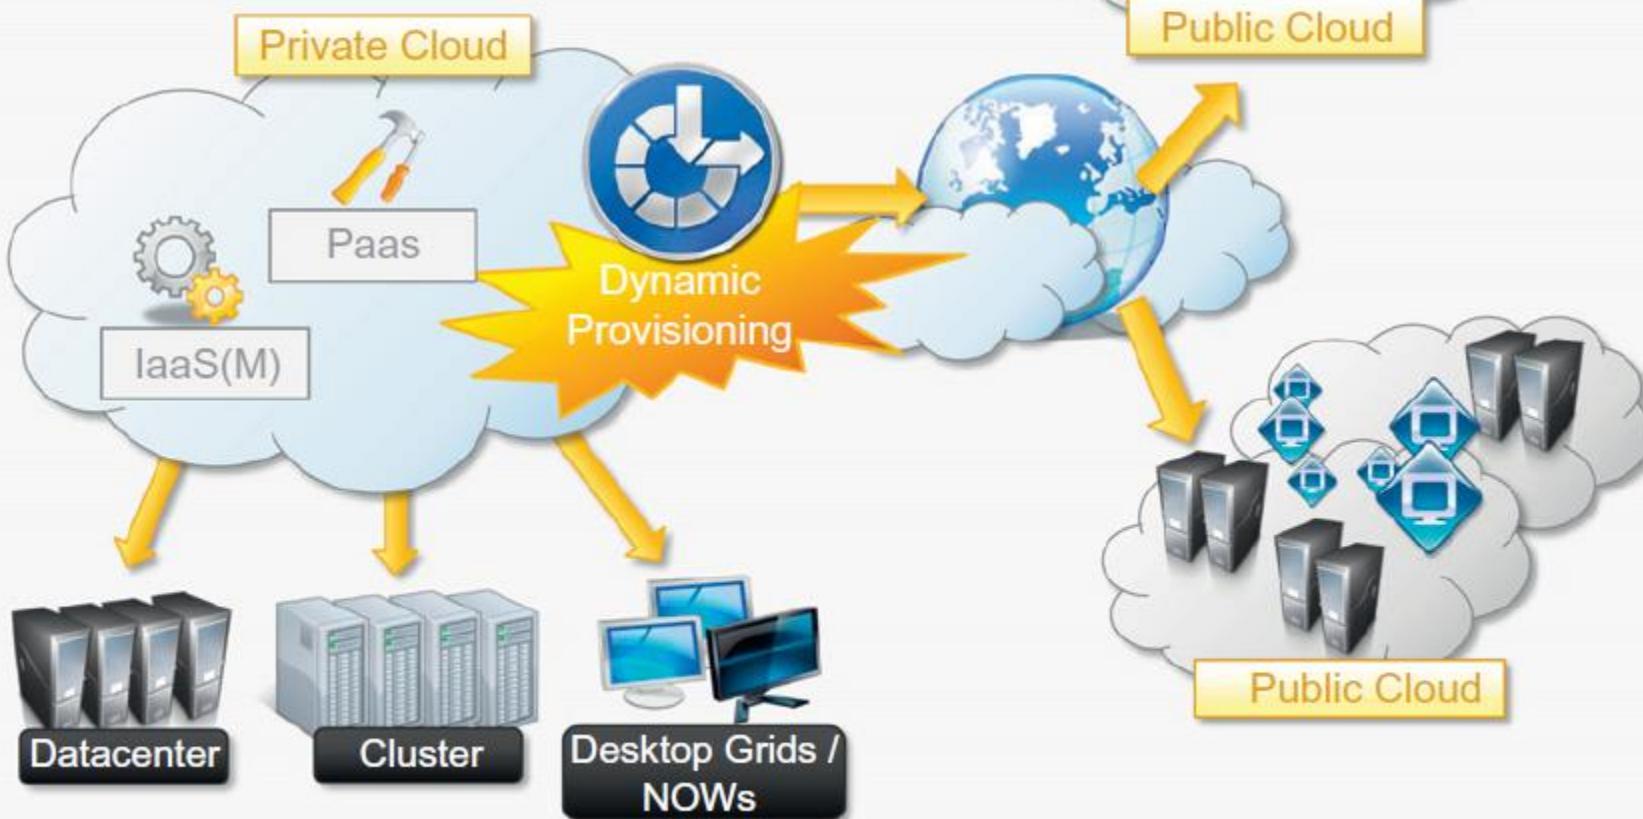
- ▶ Control, privacy, government, military agencies
- ▶ Depending on where data resides: US APATRIOT Act provides its government and other agencies access to information, including that belonging to any company that stores information in the U.S. territory
- ▶ Can leverage existing on prem hardware and software
- ▶ Migration of existing infra
- ▶ IT services on demand
- ▶ Customer information protection, Infra SLAs, Compliance
- ▶ Openstack, VMWare vCloud

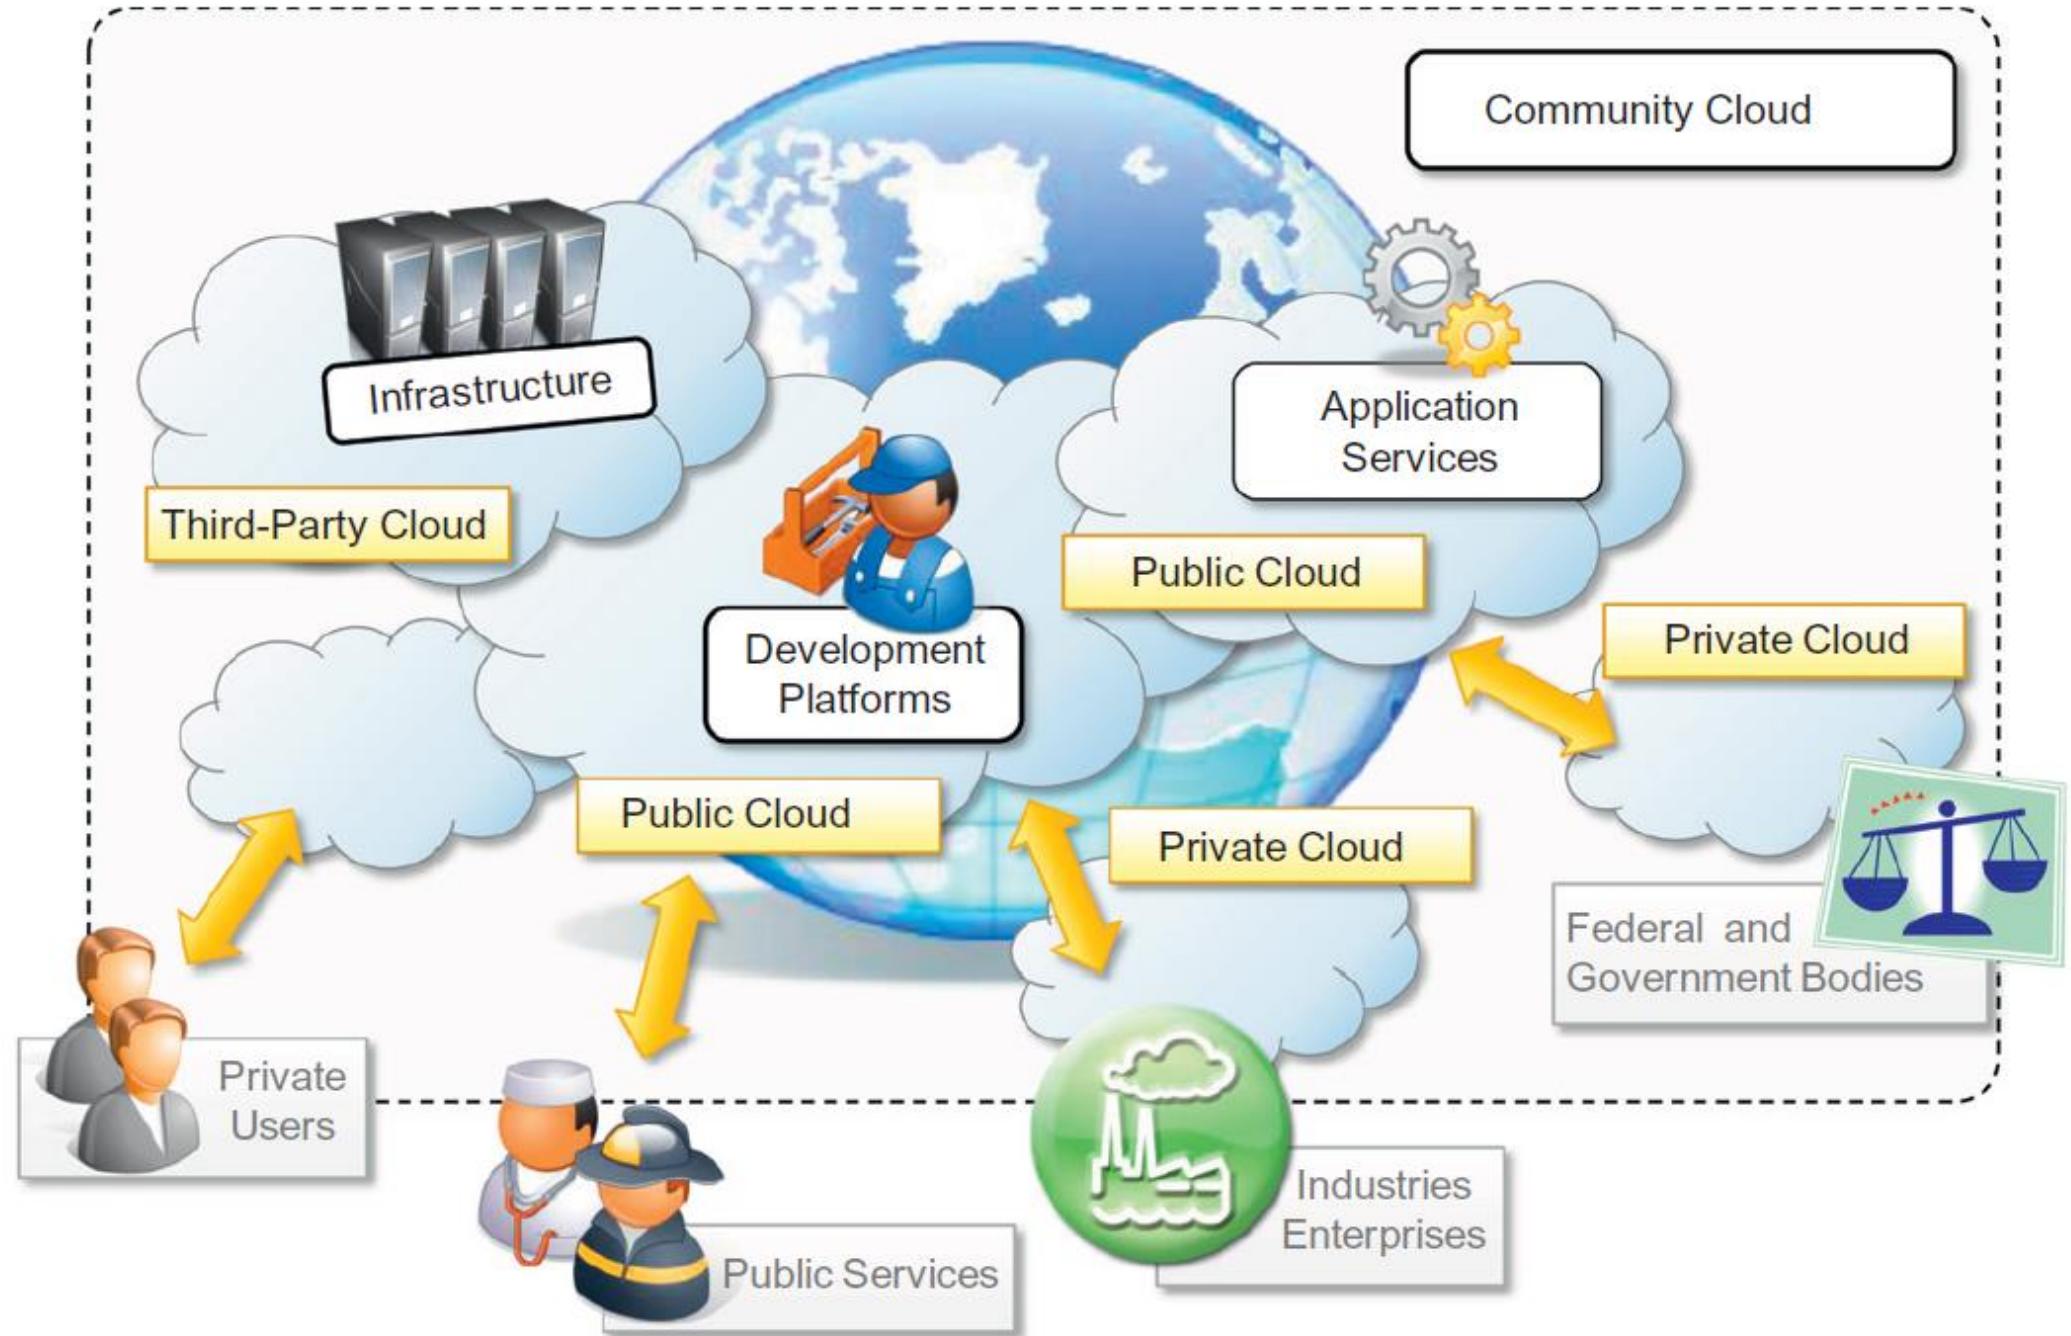


# Hybrid Cloud

- ▶ Keep processing of information within premises
- ▶ Leverage existing hardware and software
- ▶ Scale depending on needs and sustain peak load
- ▶ private cloud that integrates additional services or resources from one or more public clouds
- ▶ Heterogeneous cloud
- ▶ Dynamic provisioning to meet demand – cloudbursting
- ▶ Mostly applies to hardware
- ▶ OpenNebula: cost based advanced Haizea scheduler uses QoS and budget

## Hybrid / Heterogeneous Cloud





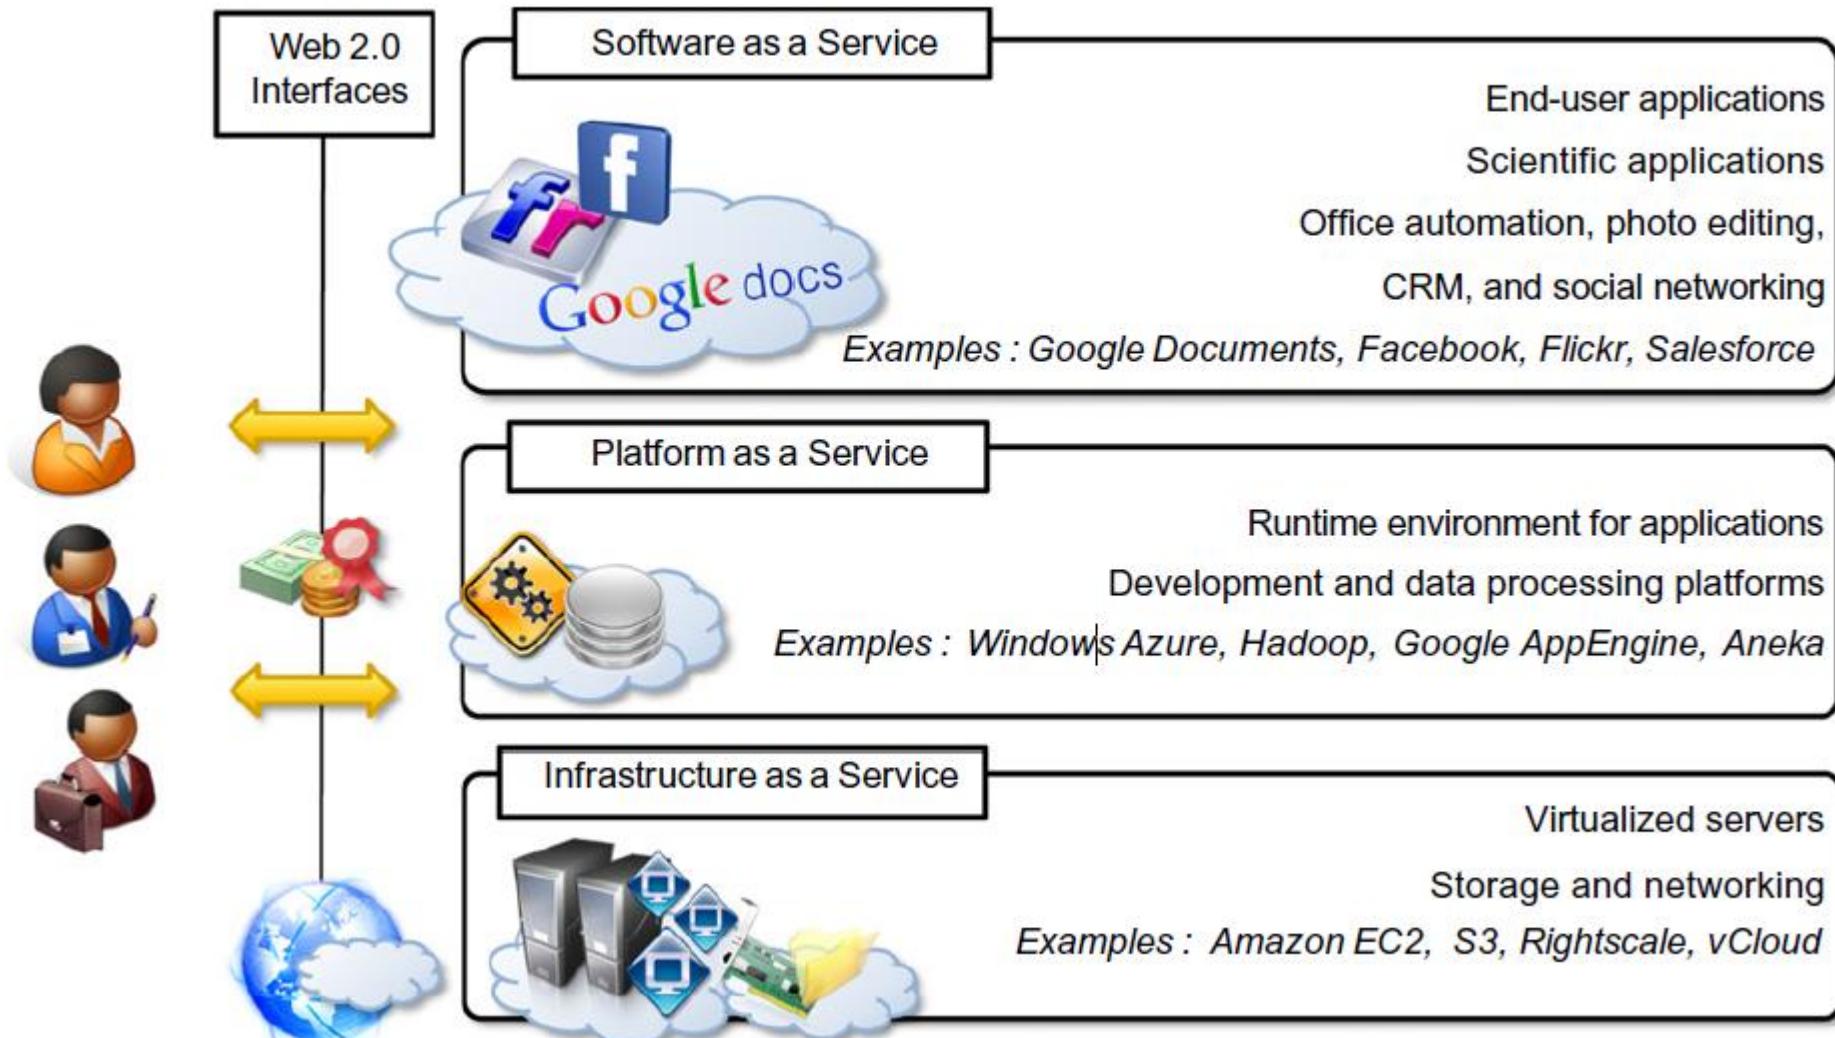
# Pricing

- ▶ Tiered: different types of EC2 instances
- ▶ Per unit: data transfer, per GB storage
- ▶ Subscription based: SaaS, for application(s) / service, per month
- ▶ Combination of above: per GB per month for s3fa

# Assignment

- ▶ Google account
- ▶ Google drive
- ▶ Spreadsheet
- ▶ Note

# Layers / Reference Model



# Cloud Computing Platforms

SAMEER MAHAJAN

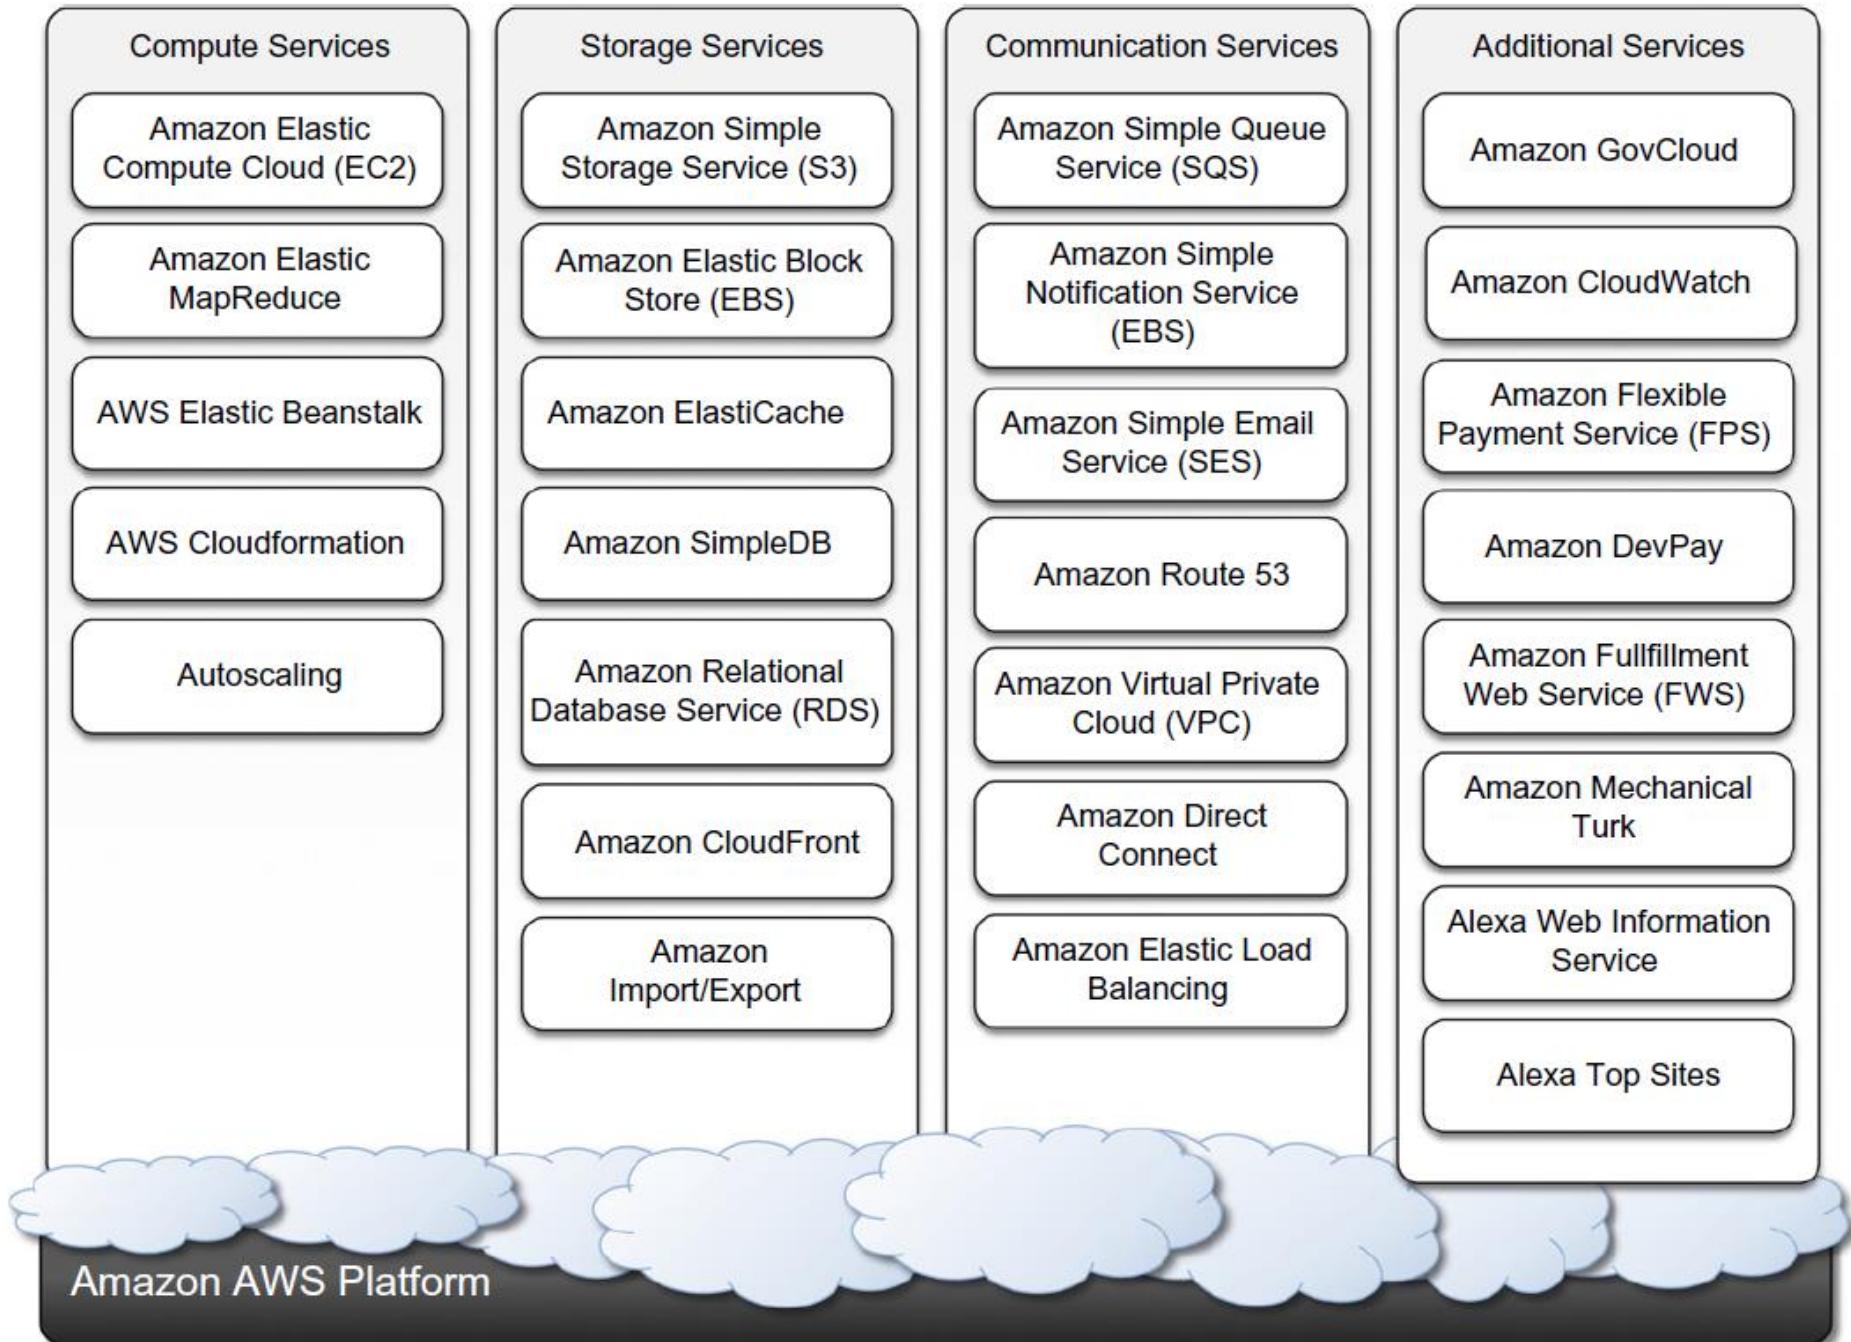
[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

**Table 9.1** Some Example Cloud Computing Offerings

Vendor/Product	Service Type	Description
Amazon Web Services	IaaS, PaaS, SaaS	Amazon Web Services (AWS) is a collection of Web services that provides developers with compute, storage, and more advanced services. AWS is mostly popular for IaaS services and primarily for its elastic compute service EC2.
Google AppEngine	PaaS	Google AppEngine is a distributed and scalable runtime for developing scalable Web applications based on Java and Python runtime environments. These are enriched with access to services that simplify the development of applications in a scalable manner.
Microsoft Azure	PaaS	Microsoft Azure is a cloud operating system that provides services for developing scalable applications based on the proprietary Hyper-V virtualization technology and the .NET framework.
SalesForce.com and Force.com	SaaS, PaaS	<a href="#">SalesForce.com</a> is a Software-as-a-Service solution that allows prototyping of CRM applications. It leverages the <a href="#">Force.com</a> platform, which is made available for developing new components and capabilities for CRM applications.
Heroku	PaaS	Heroku is a scalable runtime environment for building applications based on Ruby.
RightScale	IaaS	Rightscale is a cloud management platform with a single dashboard to manage public and hybrid clouds.

# Amazon Web Services (AWS)

- ▶ SOAP or RESTful web service interface, web based console, CLI, SDK
- ▶ Raw compute: EC2, raw storage: s3
- ▶ EC2 VM with image having predefined software stack, instance configured with memory, processors and storage, remote access
- ▶ Amazon Machine Image (AMI) templates to create a VM, stored in s3 and have unique identifiers, contains Amazon Ramdisk Image (ARI) and Amazon Kernel Image (AKI)
- ▶ vCPU / virtual cores and ECU =1.0 -1.2 GHz 2007 Opteron or Xeon processor
- ▶ Volatile spot instances with user defined price thresholds



Instance Type	ECU	Platform	Memory	Disk Storage	Price (U.S. East) (USD/hour)
Standard instances					
Small	1(1 × 1)	32 bit	1.7 GB	160 GB	\$0.085 Linux \$0.12 Windows
Large	4(2 × 2)	64 bit	7.5 GB	850 GB	\$0.340 Linux \$0.48 Windows
Extra Large	8(4 × 2)	64 bit	15 GB	1,690 GB	\$0.680 Linux \$0.96 Windows
Micro instances					
Micro	< = 2	32/64 bit	613 MB	EBS Only	\$0.020 Linux \$0.03 Windows
High-Memory instances					
Extra Large	6.5(2 × 3.25)	64 bit	17.1 GB	420 GB	\$0.500 Linux \$0.62 Windows
Double Extra Large	13(4 × 3.25)	64 bit	34.2 GB	850 GB	\$1.000 Linux \$1.24 Windows
Quadruple Extra Large	26(8 × 3.25)	64 bit	68.4 GB	1,690 GB	\$2.000 Linux \$2.48 Windows
High-CPU instances					
Medium	5(2 × 2.5)	32 bit	1.7 GB	350 GB	\$0.170 Linux \$0.29 Windows
Extra Large	20(8 × 2.5)	64 bit	7 GB	1,690 GB	\$0.680 Linux \$1.16 Windows
Cluster instances					
Quadruple Extra Large	33.5	64 bit	23 GB	1,690 GB	\$1.600 Linux \$1.98 Windows
Cluster GPU instances					
Quadruple Extra Large	33.5	64 bit	22 GB	1,690 GB	\$2.100 Linux \$2.60 Windows

# EC2 Environment

- ▶ Storage: Ephemeral or EBS
- ▶ Networking: Internal IP (within EC2, client to internet), public IP (remote access over internet), elastic IP (server, failover)
- ▶ Security: Access control
- ▶ Domain name `ec2-xxx-xxx-xxx-xxx.compute-x.amazonaws.com`, `xxx-xxx-xxx-xxx` represents external IP `compute-x` gives availability zone
- ▶ 69 availability zones priced differently, 23 regions (Nov 2019)
- ▶ Key pair, security groups, firewall rules

# Advanced Compute Services

- ▶ CloudFormation: templates, describe resources and services and relations between them required to run an application, integrate EC2 with other services like S3, SimpleDB, SQS, SNS, Route 53, Elastic Beanstalk etc.
- ▶ Elastic Beanstalk: package and deploy applications, provision instances, Web applications with Java / Tomcat, WAR file, flexible, higher level
- ▶ Elastic MapReduce: Hadoop, EC2, S3, Support Pig, Hive etc., dynamic sizing, selects appropriate configurations of EC2 instances, basic web applications to quickly run data-intensive applications without writing code

# Storage Services: Simple Storage Service (S3)

- ▶ distributed object store
- ▶ Buckets: virtual containers to store objects, cannot be further partitioned, top level, no nesting
- ▶ Objects: stored content
- ▶ RESTful interface: HTTP Requests (GET, PUT,DELETE,HEAD, and POST)
- ▶ Renmaing, modifying or relocating not allowed
- ▶ Eventual consistency
- ▶ Requests fail occasionally

# S3 (contd)

- ▶ three different URI ways of addressing
  - ▶ Canonical form: `http://s3.amazonaws.com/bukect_name/[object_name]`
  - ▶ Subdomain form: `http://bucket-name/s3.amazonaws.com/[object_name]`
  - ▶ Virtual hosting form: `http://bucket-name.com/[object_name]` needs DNS
- ▶ S3 is flat data store so unique bucket names across all users
- ▶ Object ACL: `http://s3.amazonaws.com/bukect_name/object_name?acl`
- ▶ Bucket server logging: `http://s3.amazonaws.com/bucket_name?logging`
- ▶ Metadata are manipulated as attributes of in the request
- ▶ Bucket and all its objects belong to a specific region

# S3 (further contd)

- ▶ List content of bucket with GET
- ▶ Name cannot be longer than 1024 UTF-8 encoded bytes
- ▶ Names can have '/', way to achieve folder hierarchy
- ▶ Maximum object size is 5 GB
- ▶ Access Control Policies (ACP): READ, WRITE, READ\_ACP, WRITE\_ACP, FULL\_CONTROL
- ▶ Advanced features
  - ▶ Server access logging
  - ▶ Torrent

# Elastic Block Store (EBS)

- ▶ Persistent storage in the form of volumes
- ▶ Can be mounted at EC2 instance startup
- ▶ Up to 1 TB storage, block device interface (raw, FS etc.)
- ▶ Stored in s3, cloning, snapshot
- ▶ Typically in same AZ, lazily loaded, not shared, multiple volumes
- ▶ Cost by storage size in S3, #I/O requests and throughput
- ▶ SSD / HDD, gp3 / gp2 / io2 / st1 / sc1
- ▶ \$0.08/GB-month, 3,000 IOPS free, \$0.005/provisioned IOPS-month

# ElastiCache

- ▶ Elastic in memory cache based on cluster of EC2 instances
- ▶ Memcached compatible protocol through web services
- ▶ ElastiCache cluster is dynamically (re)sized based on demand
- ▶ Automatic patch management, failure detection, recovery
- ▶ Priced according to EC2 pricing model with some differences due to the use of the caching service installed on such instances

# Structured Storage Solutions

- ▶ Preconfigured EC2 AMIs
  - ▶ IBM DB2, Microsoft SQL Server, MySQL, Oracle, PostgreSQL, Sybase, and Vertica
  - ▶ Priced hourly according to the EC2 cost model
  - ▶ EC2 user has to configure, maintain, and manage the DBMS
- ▶ Relational Data Storage (RDS)
  - ▶ Managed MySQL or Oracle
  - ▶ Advanced features of multi-AZ deployment and read replicas
- ▶ SimpleDB: lightweight, not fully relational, flexible, fully managed
- ▶ Aurora

# Communication Services

- ▶ Among applications and services residing within the AWS infrastructure

## **Virtual Networking**

- ▶ Between compute and storage
- ▶ Virtual Private Cloud (VPC) and Direct Connect (user network to AWS), Route 53 for naming (dynamic DNS), Identity Access Management (IAM)

## **Messaging**

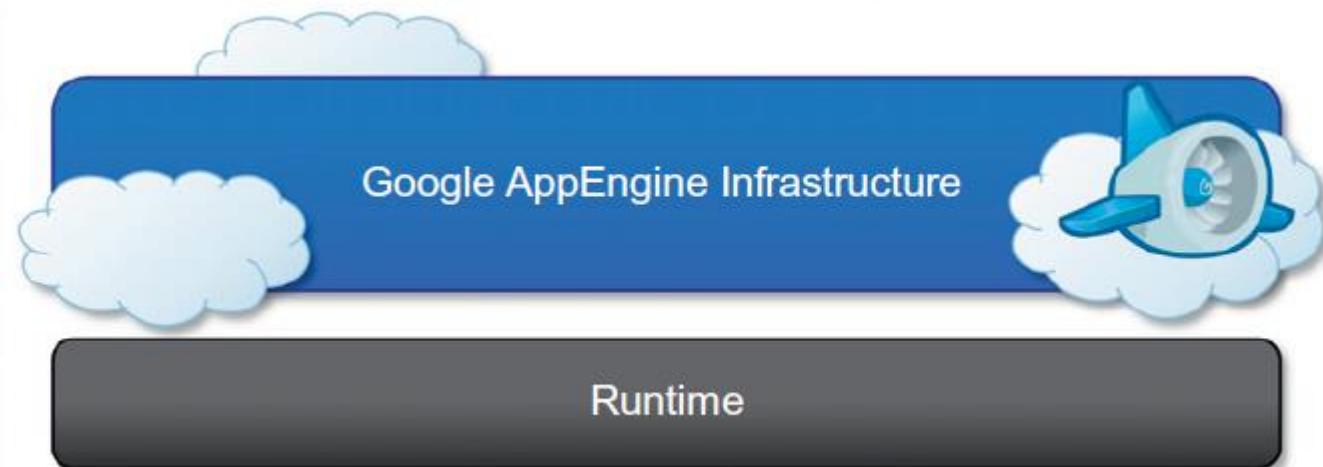
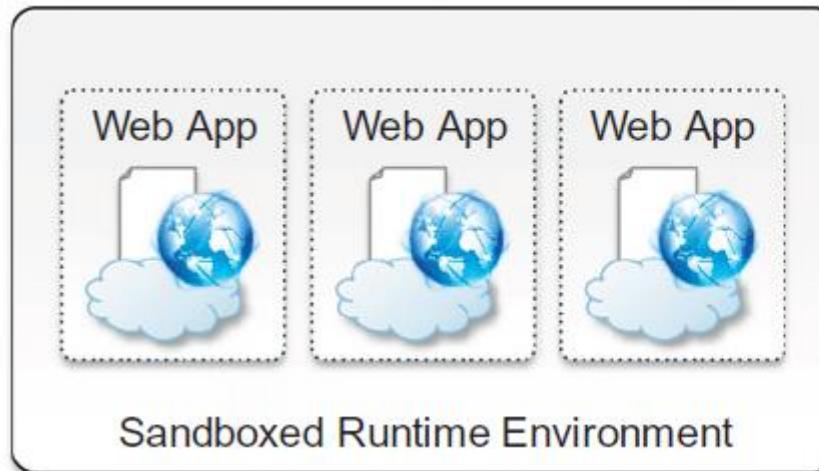
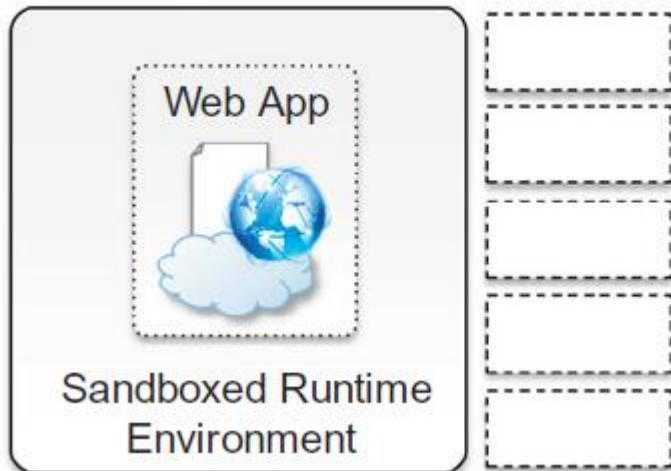
- ▶ Simple Queue Service (SQS), Simple Notification Service(SNS):publish-subscribe, Simple Email Service (SES)
- ▶ Pay as you go

# Additional Services

- ▶ **CloudFront:** content delivery network (CDN) on distributed storage infra, leverages strategically located edge servers, reduced transfer time, Web service APIs, distribution with origin server can be in AWS or outside, static or streaming content, access rules, cheaper than S3
- ▶ **CloudWatch:** monitoring, earlier subscription, now free
- ▶ **Flexible Payment Service (FPS):** sell goods and services to other AWS users
- ▶ **Athena**
- ▶ **Glue**
- ▶ **QuickSight**
- ▶ AWS SDK for Python (**Boto3**)

# Google AppEngine

- ▶ PaaS for developing and hosting scalable web applications
- ▶ Java, Python, Golang
- ▶ Managed or interpreted languages
- ▶ Sandboxing: Cannot write to server's FS, execute more than 30 seconds
- ▶ Java Server Pages (JSP), Python Web Application framework (webapp)



# Storage

- ▶ **Static File Servers:** CSS, html, javascript static content
- ▶ **DataStore:** semi structured data, object database, BigTable based, entities and properties, entity has immutable key, indexes, transactions, query by specifying either the key or conditions on the values of the properties, pre computed indexes (based on queries) – query execution is independent of data size but influenced only by result set size, entity is updated atomically, multiple updates or multiple entities (in the same entity group) in a transaction, optimistic concurrency control, no exception retrieving an entity

# Application Services

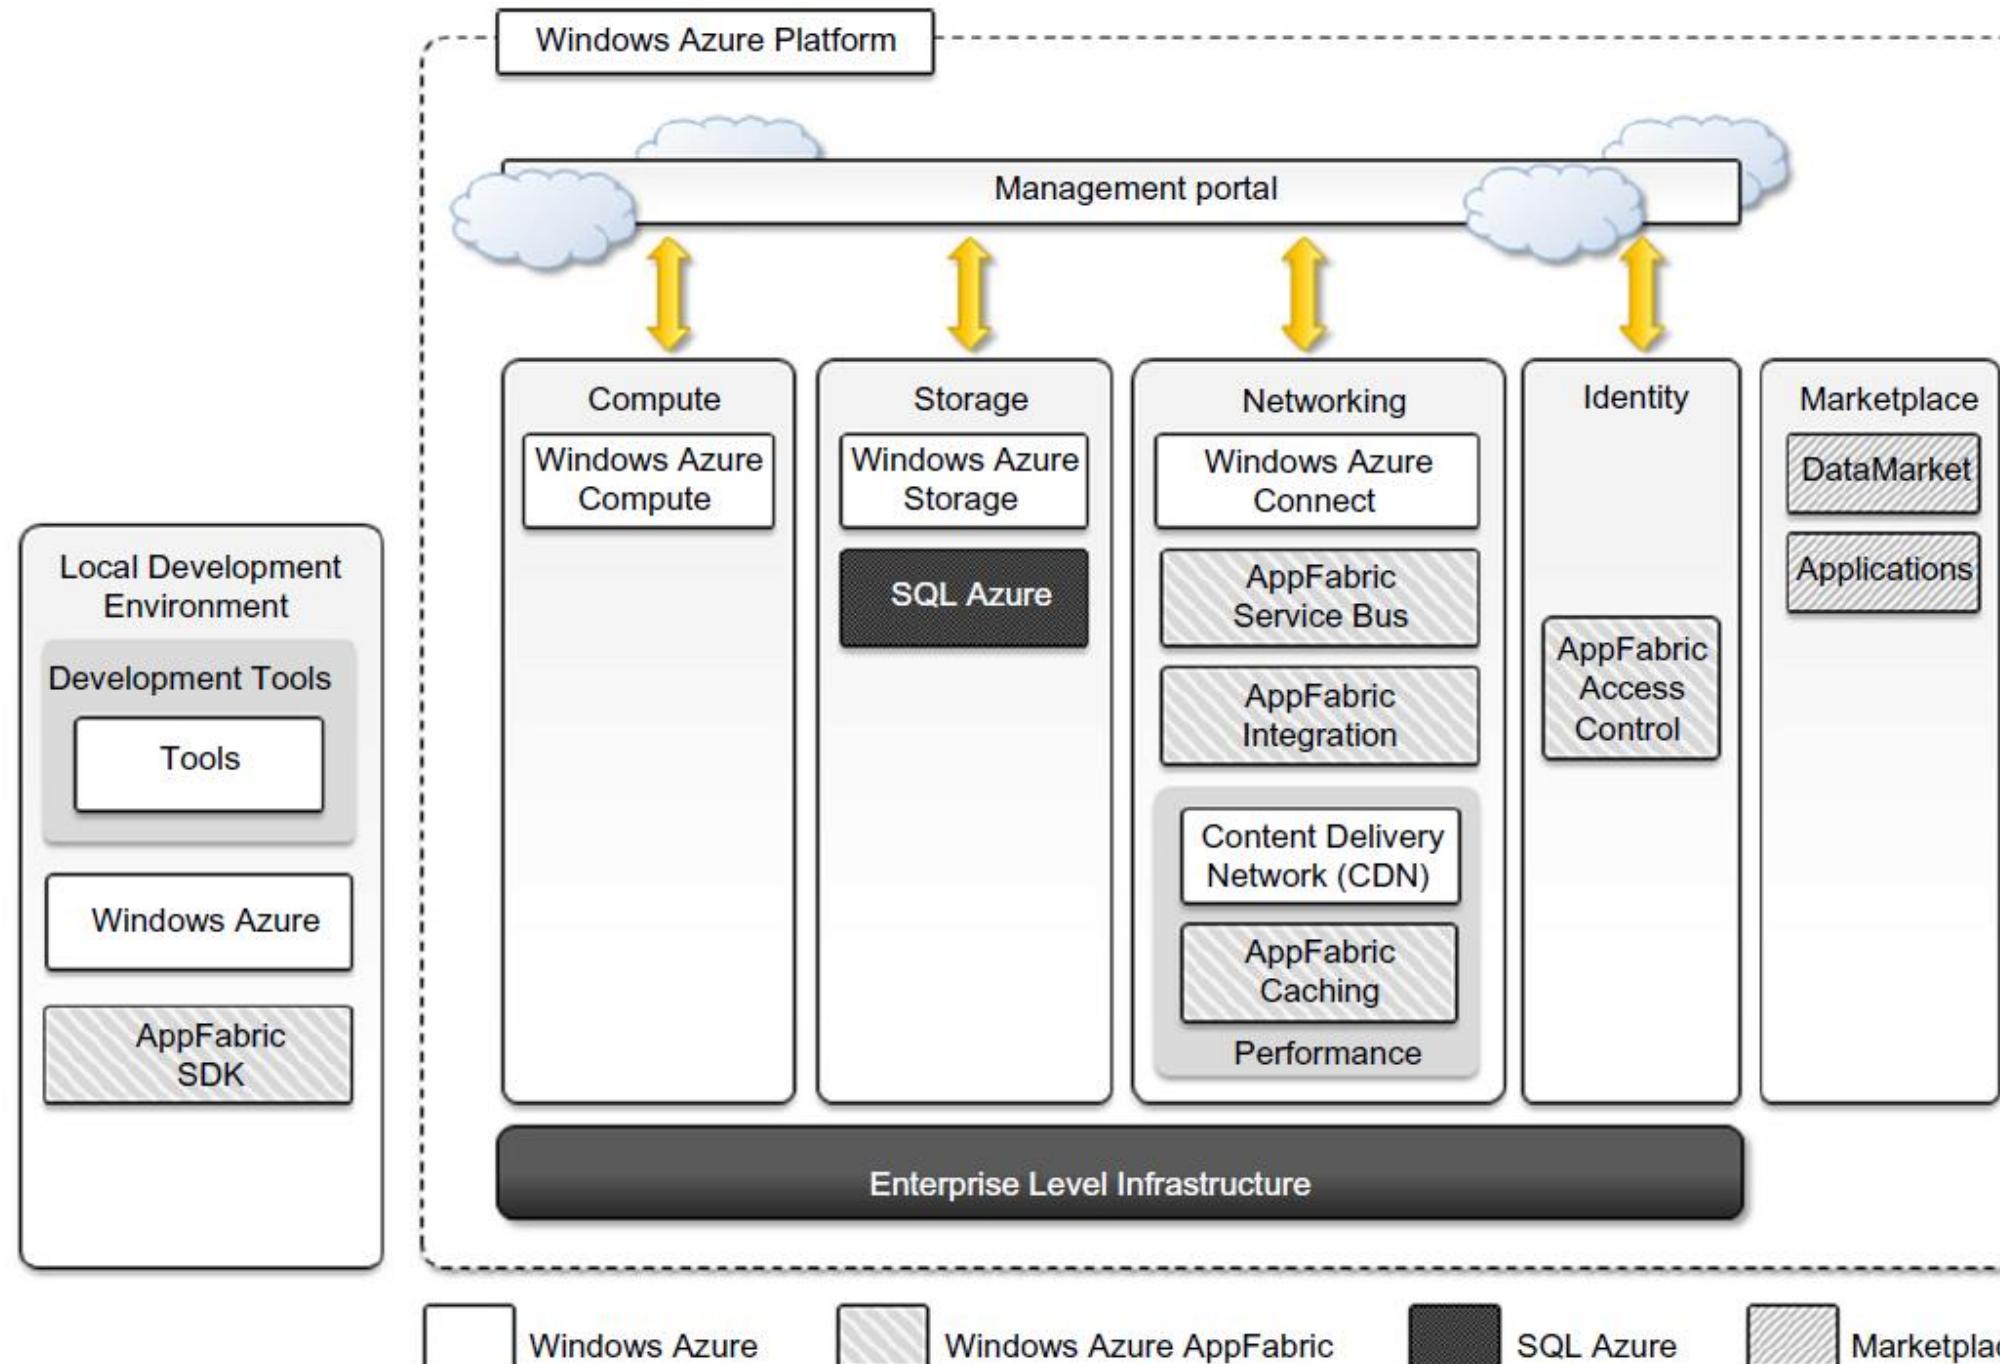
- ▶ **UrlFetch:** retrieve remote resources, request can specify deadline, async, integrate meshes into web pages, leverage remote web service
- ▶ **MemCache:** distributed in memory cache for fast access
- ▶ **Mail and IM:** communication, XMPP for IM
- ▶ **Account Management:** Google Accounts, user profiles, authentication
- ▶ **Image manipulation:** light weight image processing with speed

# Compute Services

- ▶ Real time and synchronous within web request timeframe
- ▶ **Task queues:** submit a task for later execution, long computations, up to 10 queues, queue invokes request handler for task execution, queue re executes failed task until it is successful
- ▶ **Cron jobs:** to be performed at specific time (e.g. nightly), does not re execute in case of failure, maintenance or notifications

# Application Lifecycle

- ▶ Develop on local development environment, simulates AppEngine runtime by providing mock implementations of DataStore, MemCache, services, hosting web applications, monitoring and tuning
- ▶ **Java SDK:** Eclipse with google appengine plugin, servlet, Eclipse web platform, CLI
- ▶ **Python SDK:** standalone GoogleAppEngineLauncher tool, integrated webapp, Django, CLI
- ▶ Create application identifier (.appspot.com), deploy / upload, manage and monitor by administrative console, versioning, billing
- ▶ Free service with limited quotas reset every 24 hours
- ▶ billable, fixed, and per-minute quotas; 403 error once quota reached



# Compute Services

- ▶ **Web role:** web applications, hosted on IIS, auto scaled, native support since .Net 3.5, ASP.NET and WCF applications, PHP runtime via FastCGI, non IIS technologies like JSP on Tomcat can also be hosted
- ▶ **Worker role:** general compute, do not communicate with the external world through HTTP, background processes, continuously running service, any technology that runs on a Windows Server stack
- ▶ **Virtual Machine role:** fully control the computing stack, Hyper-V based, Virtual Hard Disk (VHD) with custom stack

# Storage Services

- ▶ Durable, not local, can be accessed by multiple clients at the same time and from everywhere
- ▶ **Blobs:** binary large objects, metadata, snapshots, CDN
  - ▶ **Block blobs:** sequential access, media streaming, 4 MB block, 200 GB blob
  - ▶ **Page blobs:** random access, 1 TB
- ▶ **Azure drive:** A page blob mounted as part of an NTFS tree
- ▶ **Tables:** semi structured storage, entities as rows, properties as columns, key / index, insert, update, delete, select a subset of rows, no schema, partial result sets, partitions for load balancing across servers, 100 TB, 255 properties, 1 MB row, 1 kb keys
- ▶ **Queues:** message exchange across applications, FIFO, read message -> invisible, processed message -> deleted, peeking
- ▶ Triple Replication

# Core Infrastructure

- ▶ **AppFabric:** middleware for developing, deploying, and managing applications, language agnostic interfaces
- ▶ **Access Control:** rules to access resources, integrates with authentication providers like Active Directory (AD), Windows Live, Google, single coherent identity management framework
- ▶ **Service bus:** messaging and connectivity infra, services as URLs, pub-sub, full duplex, peer to peer, multicast, async, connections per month pricing
- ▶ **Azure cache:** elastic distributed in memory cache, .Net managed objects, any format, easily integrated with applications, priced by allocated size

# Other Services

- ▶ **Virtual Network:** Connect – IP based, VM roles and Traffic Manager – load balancing
- ▶ **CDN:** static or streaming, 24 locations distributed across the world
- ▶ **SQL Azure:** RDBMS, accessible from cloud or any other location having access to cloud, Tabular Data Stream (TDS) protocol, ODBC or ADO.NET, billed by space and edition, Web edition (1 or 5 GB), Business edition (10-50 GB)
- ▶ **Windows Azure Platform Appliance:** windows, SQL, configuration of network, storage, hardware

# Assignment

- ▶ Setting up big data fabric in AWS  
<https://www.youtube.com/watch?v=4VJvjkDDfs>

# Cloud Applications

SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# Cloud Applications

- ▶ Migrate to the cloud
- ▶ Bootstrap (SDK, APIs, programming environments) for the cloud
- ▶ Develop for the cloud
- ▶ Cloud first / Cloud born
- ▶ Rich Internet Applications (RIA) with and without browser
- ▶ Abstraction based at programming language, environment (e.g. browser) level

# Evolution of Abstraction

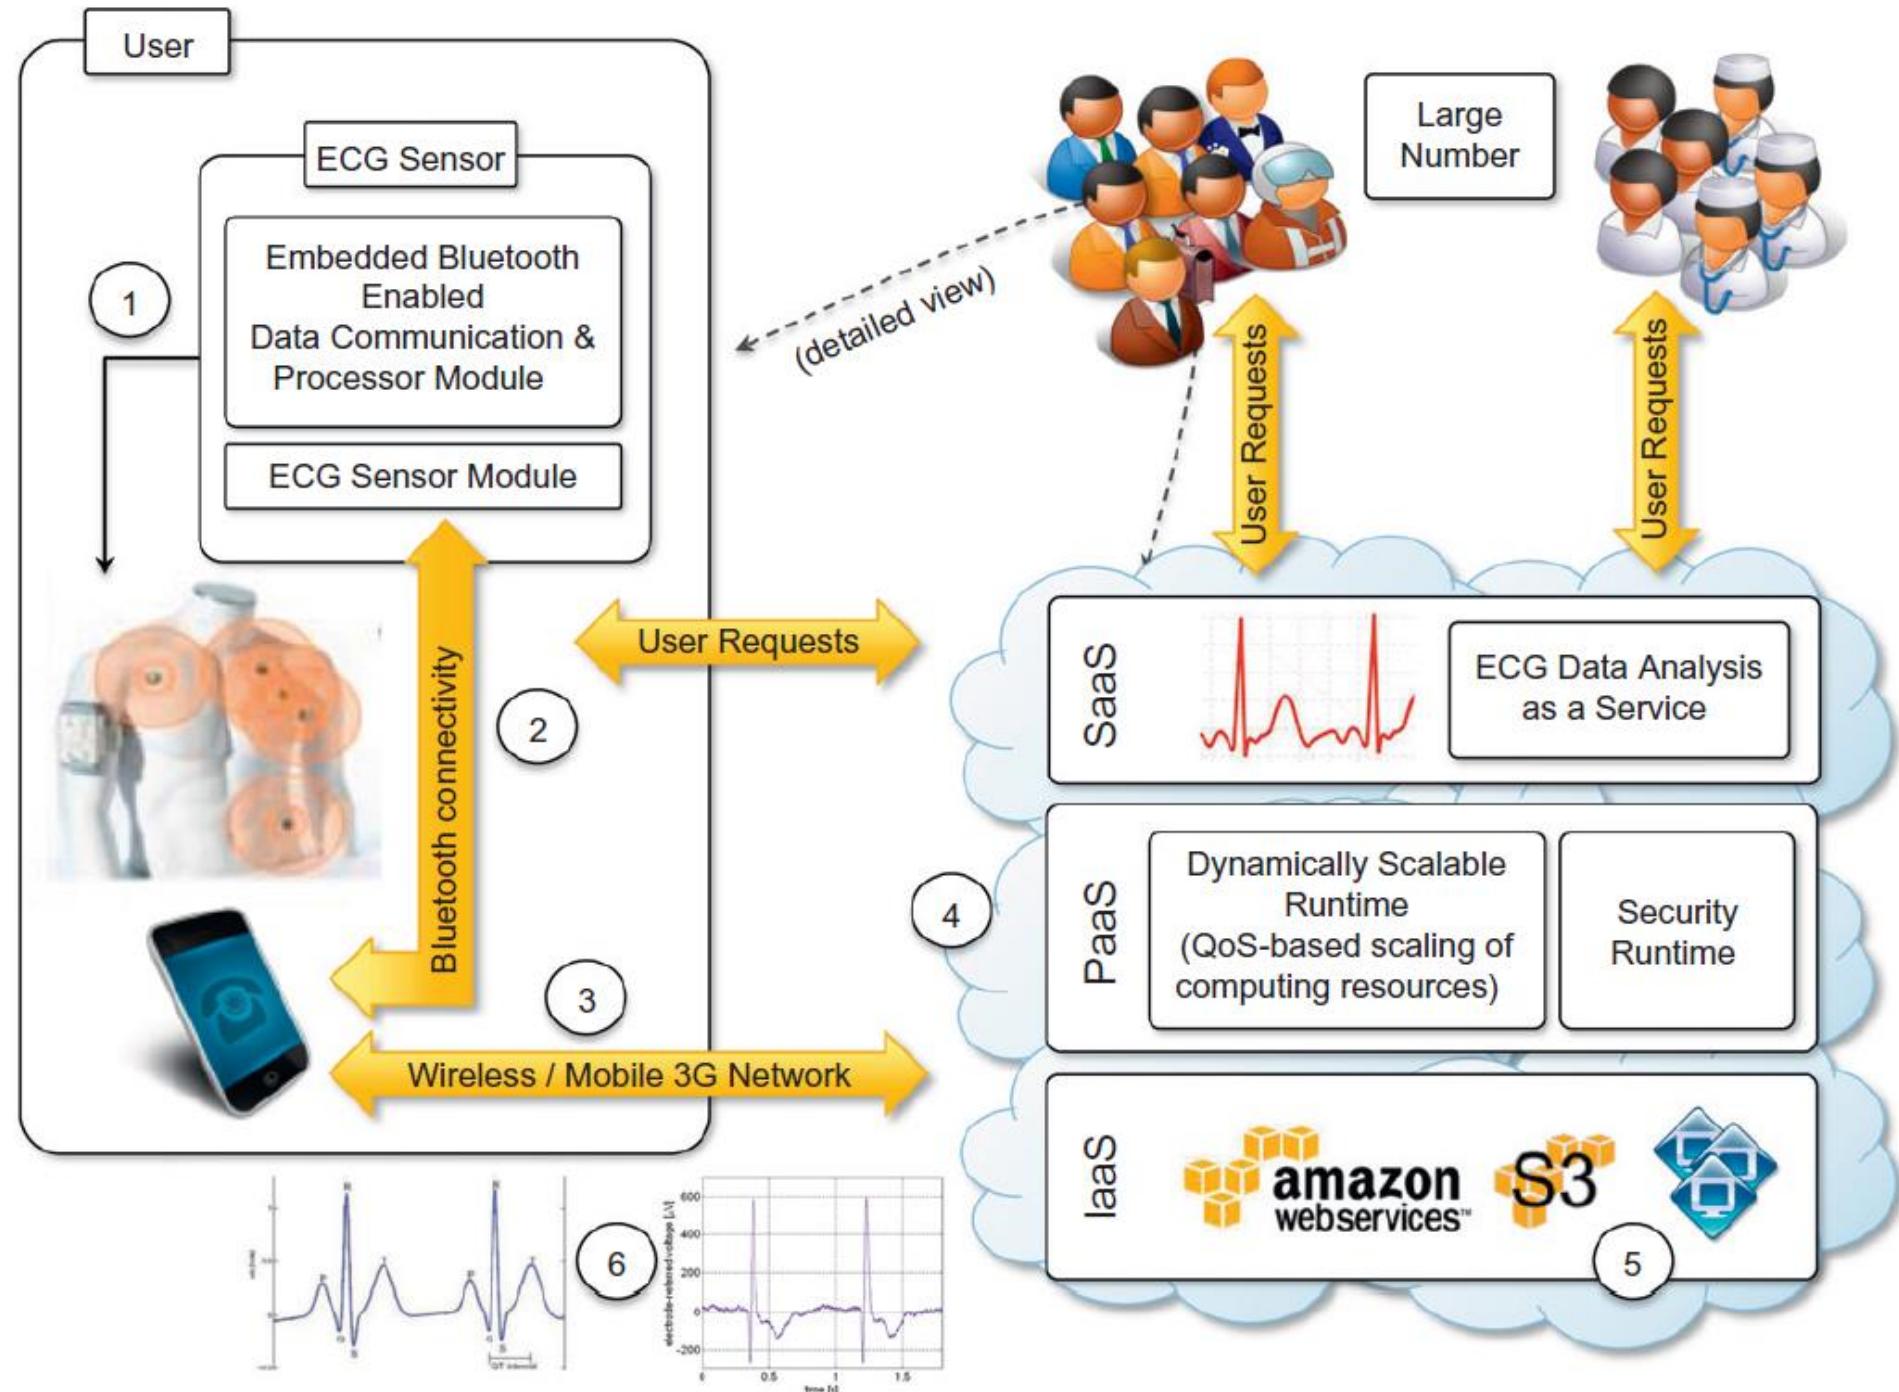
- ▶ Programming Language: assembly -> macro assembler -> cobol -> fortran  
-> C -> Java -> python, Go -> Stats / no code ML / AutoML
- ▶ Browser tightly integrated with OS and hardware, browser plugins instead of applications, simple installation, small footprints
- ▶ Adobe Integrated Runtime (AIR) : Java, Flash, Windows, Linux; Microsoft Silverlight
- ▶ Advantages: ease of development and maintainability
- ▶ Disadvantages
  - ▶ system overhead (speed, final machine code), not tight, concise
  - ▶ security, toolbars have full access to your desktop

# Clouds Use Cases

- ▶ Times: converting 11 million articles into .pdf for online archive. Estimated 100 servers with 4 TB storage taking months. AWS 100 EC2, 4 TB s3 in a day with cost of only \$240
- ▶ Washington Post: 200 EC2 server instances to process 17,481 pages of non-searchable PDF images into a searchable online library. 9 hours, a minute per page. \$144.62
- ▶ Above two are processing and throwaway clouds, use only once
- ▶ Storage Clouds: s3, dropbox, google drive, backup, smart switch
- ▶ Traveling clouds: USB drive with 5 preinstalled windows servers, mounted locally in AWS for a conference demo
- ▶ Occasional use clouds: interopnet VM images / virtual appliance, spun up once for Vegas and once for New York events every year
- ▶ DR: storage side synced up at DR site not incurring compute side cost but still warm to able to be brought up immediately in case of disaster without any delay

# Scientific Applications

- ▶ Researchers and academics
- ▶ Availability of compute and sustainable price of storage compared to complete in-house deployment, minimal changes
- ▶ High performance (HPC), high throughput (HTC), data intensive
- ▶ IaaS, PaaS, MapReduce – data intensive
- ▶ Healthcare – diagnostics, health monitoring, ECG analysis



# Case Study – Data science based technology platform

## About the client

Our client is a technology healthcare startup developing a platform to diagnose mental & physical health of patients based on their voice samples. They have partnered with a premier research lab and licensed speech features pertaining to some health conditions such as depression.

## Challenges

- Create a platform that can score a participant on clinical (depression, Alzheimer's etc.) or wellness (nasality, stress, sleepiness etc.) based on few seconds of voice.
- Develop mobile apps for
  1. Clinical studies with hospitals or pharma partners for discovery of new vocal biomarkers
  2. Organizations to monitor wellness of their employees or end-customers
- Power partner apps through APIs and SDKs
- Accelerate discovery of models for new conditions with help of internal tools

## Solution & outcome

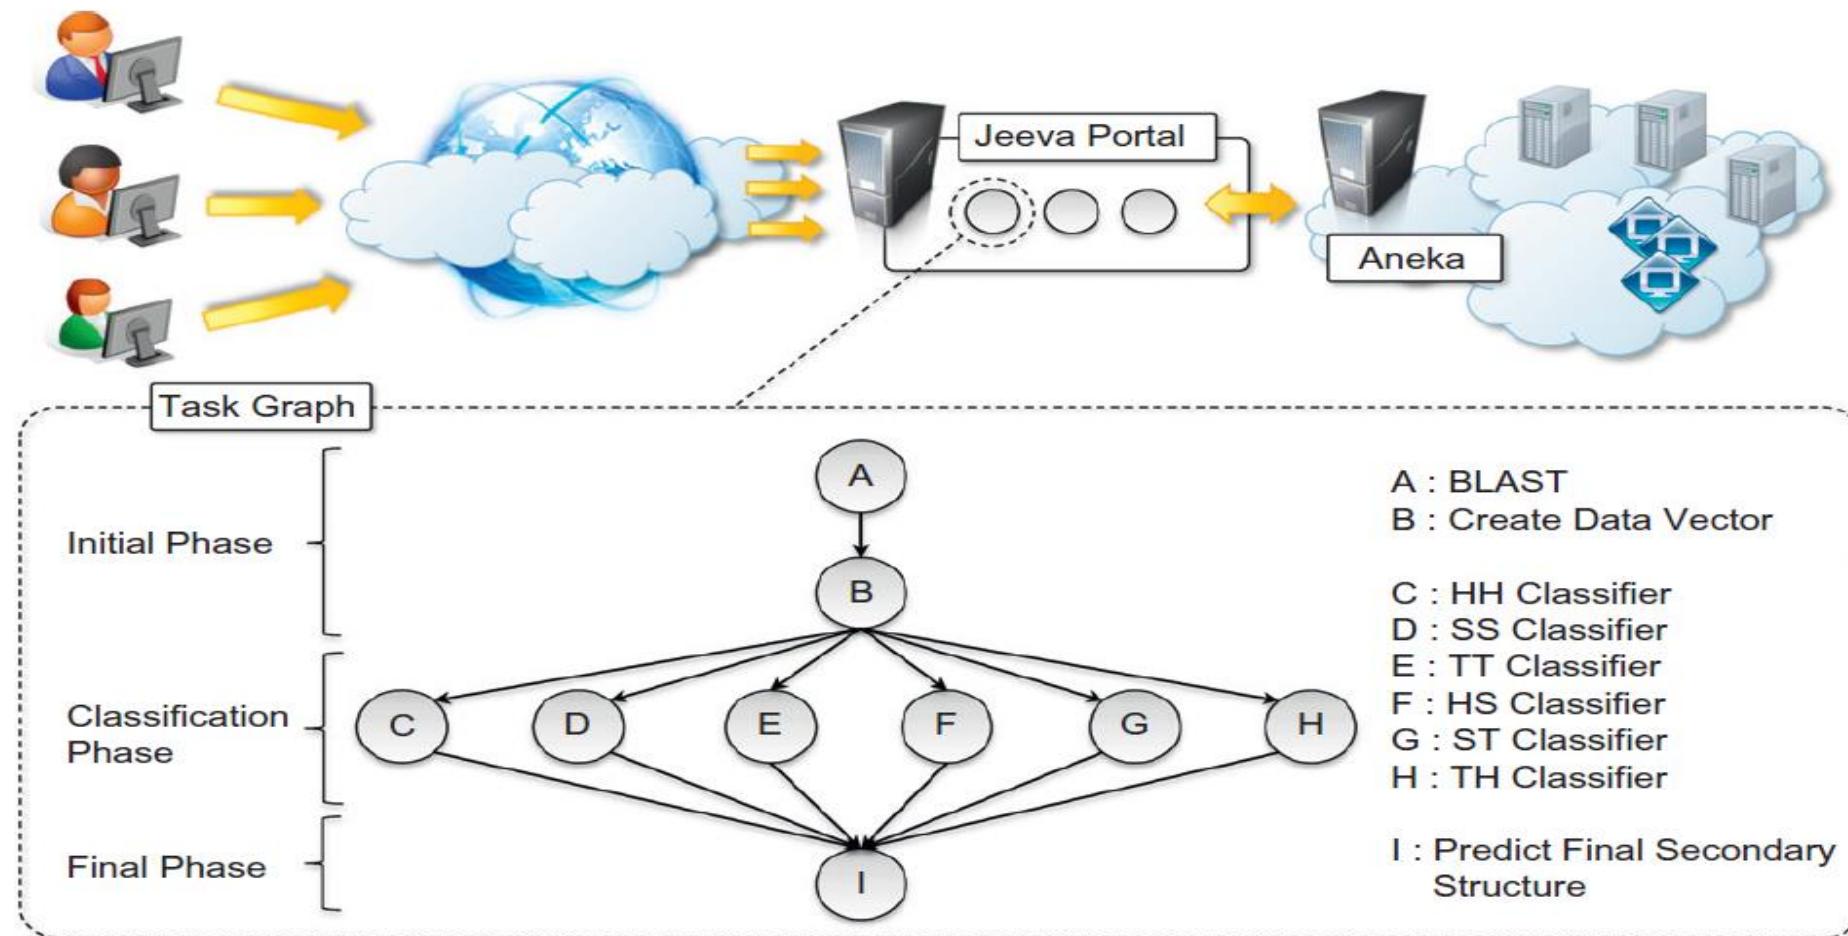


Scalable platform to power mobile apps and partner apps adhering to region-specific legal constraints

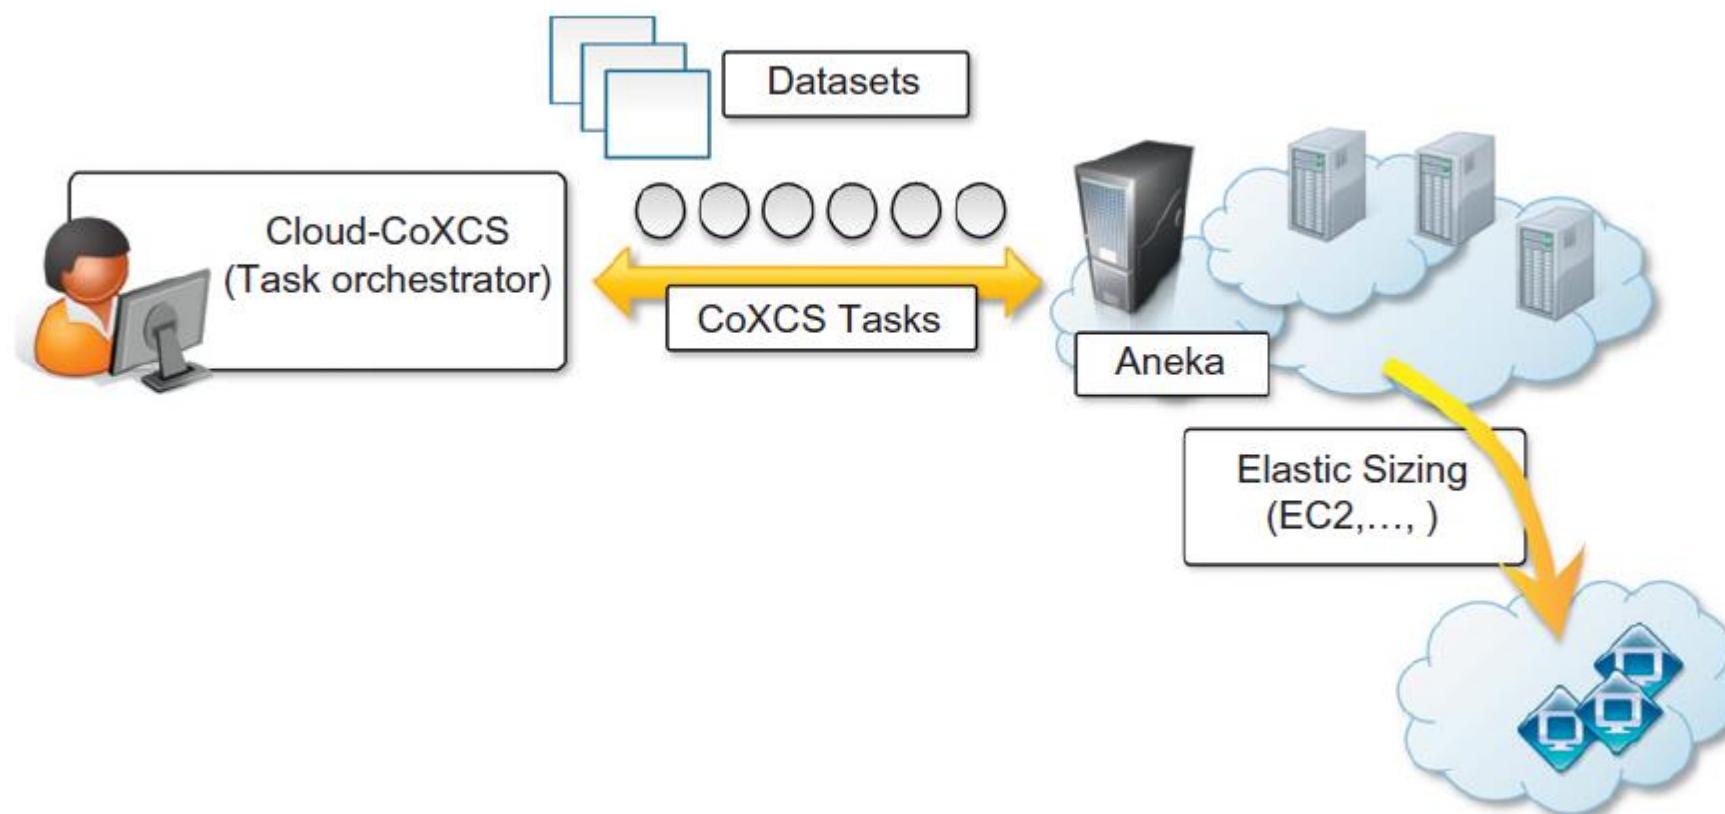


'Discovery' of models for new health conditions reduced from few months to few weeks leveraging rapid experimentation and self-service tools and processes

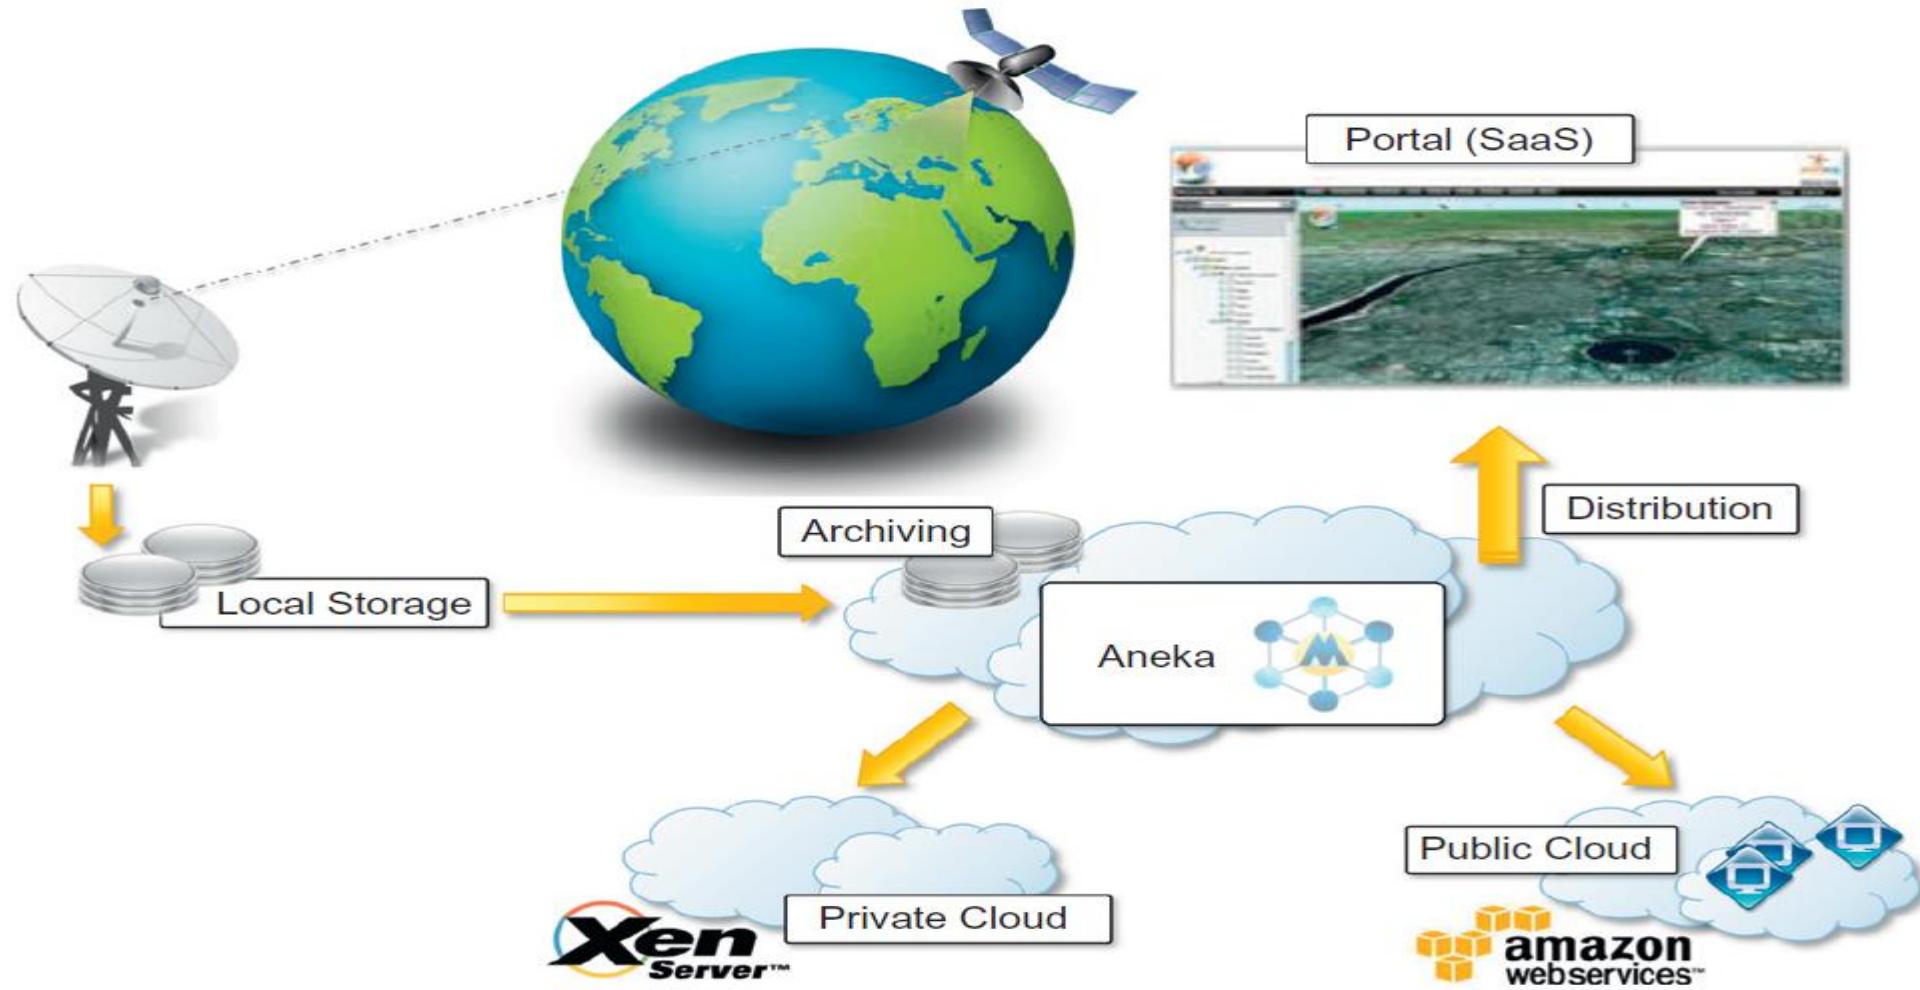
# Biology: protein structure prediction



# Biology: gene expression data analysis for cancer diagnosis



# Geoscience: satellite image processing



# Business and consumer applications

- ▶ Transform capital costs into operational costs
- ▶ Data and services can be accessed from anywhere, anytime, anyone
- ▶ Elastic nature with auto scaling – quickly release products and services
- ▶ Customer Relationship Management (CRM): salesforce.com – 100k customers, Microsoft dynamics CRM – 99.9% SLA, NetSuite
- ▶ Enterprise Resource Planning (ERP): finance, accounting, HR, manufacturing, supply chain, project management, NetSuite

# More applications

- ▶ Productivity
  - ▶ Dropbox: synchronize any file across any platform, any device, through browser or client
  - ▶ Google docs
  - ▶ Cloud web desktop: EyeOS, Xcerion XML Internet OS/3 (XIOS/3)
- ▶ Social networking: facebook, twitter, Instagram, linkedin
- ▶ Media applications: Netflix, youtube
- ▶ Multiplayer online gaming

# Provider Support

- ▶ AWS: supports all languages, environments
- ▶ GCP: web and google centric
- ▶ Azure: .Net and windows centric
- ▶ Blurring lines between mobile and desktop, 1000+ contacts on mobile phone backed up in cloud

# Cloud Security

SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# The Spectrum

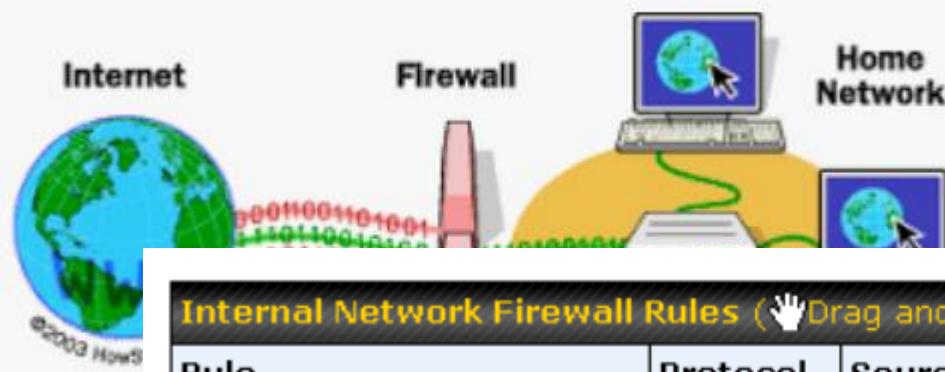
- ▶ Two categories
  - ▶ Assets (hardware, software, network infrastructure) – traditionally
  - ▶ Data – more important today
    - ▶ At rest
    - ▶ In motion
- ▶ User Authentication: only authorized users can access resources
- ▶ Filtering: only authorized data flows into or out of network, unauthorized data should be unusable
  - ▶ Firewall: rule based (devices, appliances)
  - ▶ Deep packet and pattern inspection: signatures
- ▶ Monitoring: preventing any breach

# Authentication

- ▶ Logically centralized but physically distributed authentication directory
- ▶ HQ Master – cloud / sites slaves
- ▶ Read only copies at sites
- ▶ Identity federation, Single Sign On – single authentication to identify and authorize across multiple systems
- ▶ Network Access Control (NAC) – endpoints are secure, especially mobile
- ▶ Agent based, communicates status and configuration to the central server, receives and performs any necessary updates to keep the endpoint latest and secure, polling from server based
- ▶ No access at all if not authenticated in case of NAC / cloud
- ▶ Microsoft's version is called NAP

# Cloud Filtering

- ▶ Firewall operates at low network level, ports, protocols, IP addresses, filter incoming traffic
- ▶ IDS, Intrusion Prevention System (IPS), Unified Threat Management (UTM) look for patterns, filter incoming as well as outgoing traffic
- ▶ Web filters operate at higher level and look at data stream content
- ▶ Data leakage prevention systems filter outgoing traffic
- ▶ IDS is like burglar alarm, IPS is real time
- ▶ Email filtering, application blocking
- ▶ Takes significant computing power
- ▶ Only tap at source to redirect and process elsewhere, process at source leveraging auto scale and reduce traffic



### Internal Network Firewall Rules (Drag and drop rows by the left to change rule order)

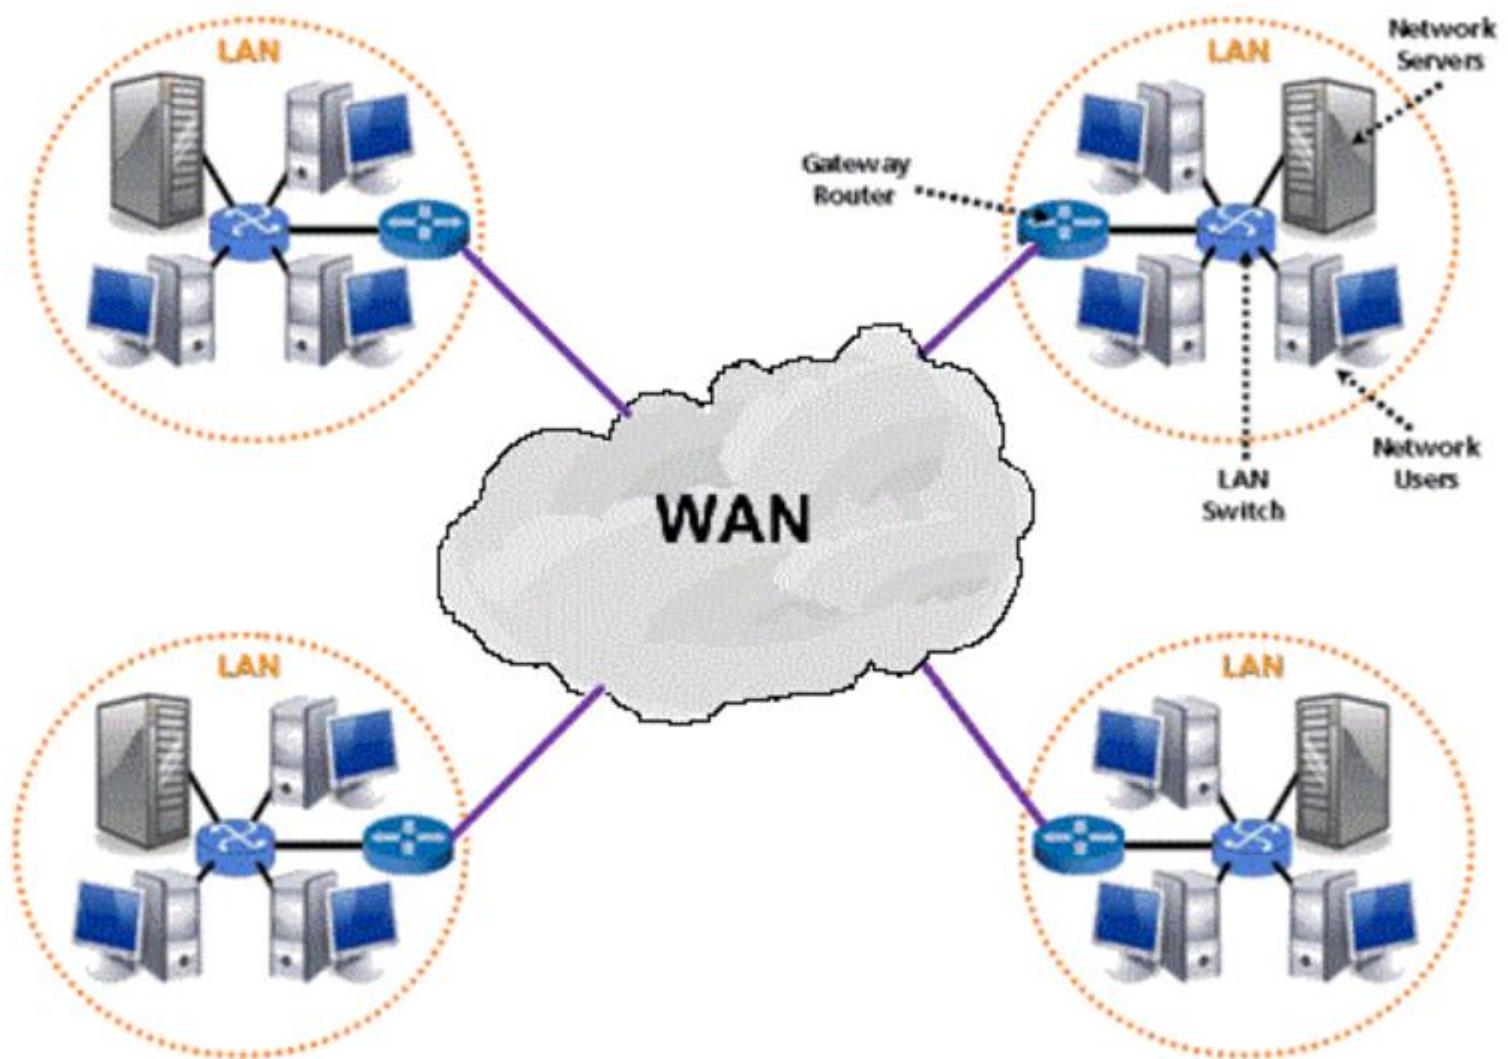


Rule	Protocol	Source	Destination	Action	
AlexJumpBox	Any	192.168.10.72	192.168.10.0/24	🚫	X
AdminAllowAll	Any	192.168.10.104/29	Any	✅	X
CLJumpBox	Any	192.168.10.104/29	10.88.11.4	✅	X
R-CELive	Any	192.168.10.157	10.88.11.4	✅	X
N-CEDemo	Any	192.168.10.158	10.12.0.12	✅	X
LocalAllow	Any	192.168.10.128/26	192.168.10.192/26	✅	X
Default	Any	Any	Any	🚫	

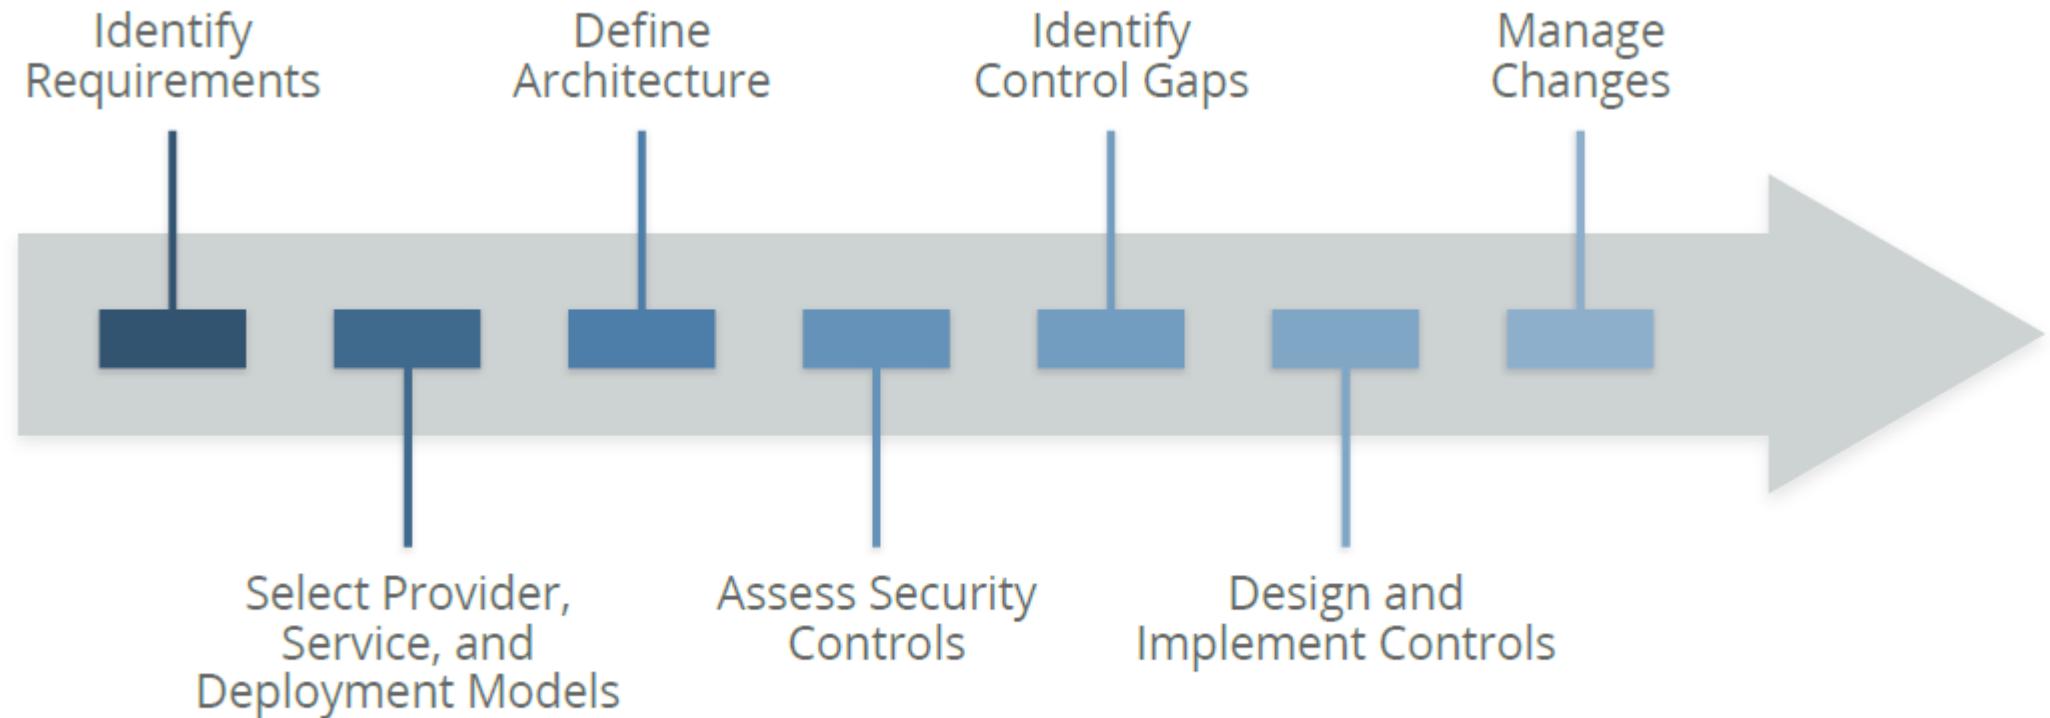
**Add Rule**

# Firewalls

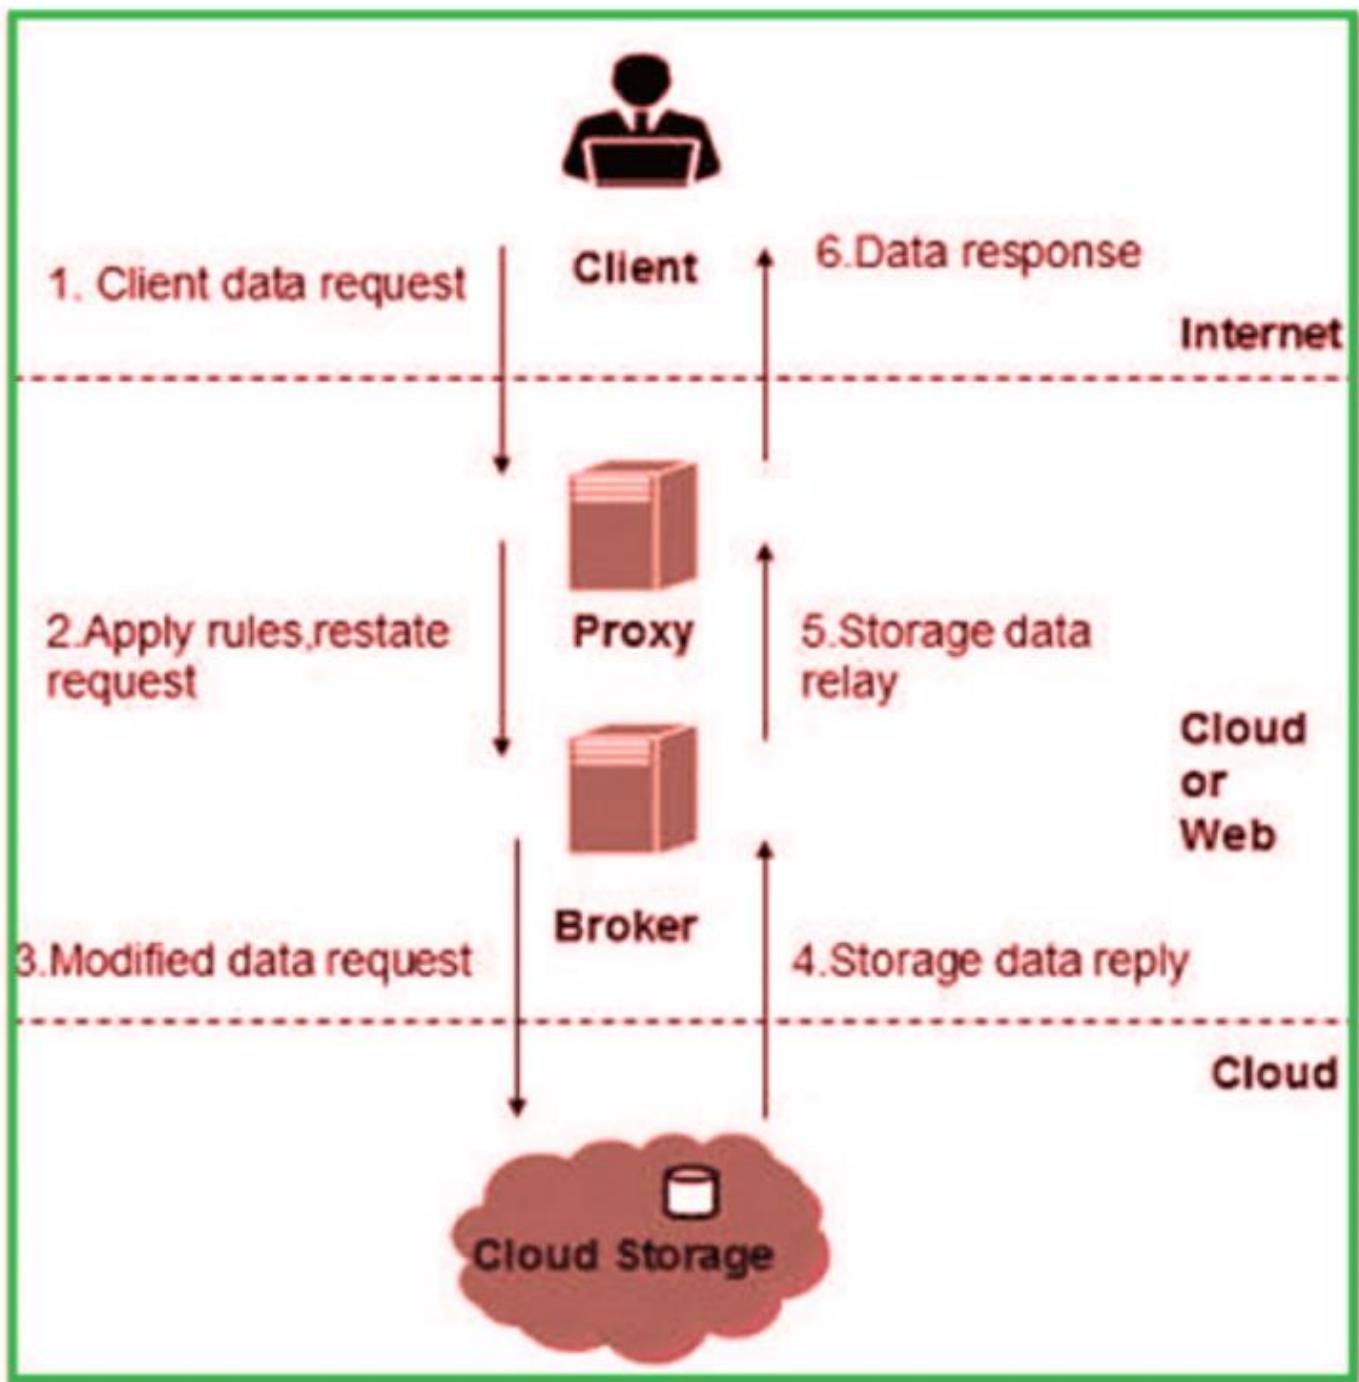
- ▶ Cisco
- ▶ SonicWall NSA 16 core caveum
- ▶ Astaro VM
- ▶ Autoscale firewall VMs from load balancers
- ▶ Autoscale servers and firewalls in front of them
- ▶ VM in cloud can also be monitored for its performance using Simple Network Monitoring Protocol (SNMP) or Windows Management Instrumentation (WMI)
- ▶ Virtual Private Network (VPN) to connect to protected virtual networks



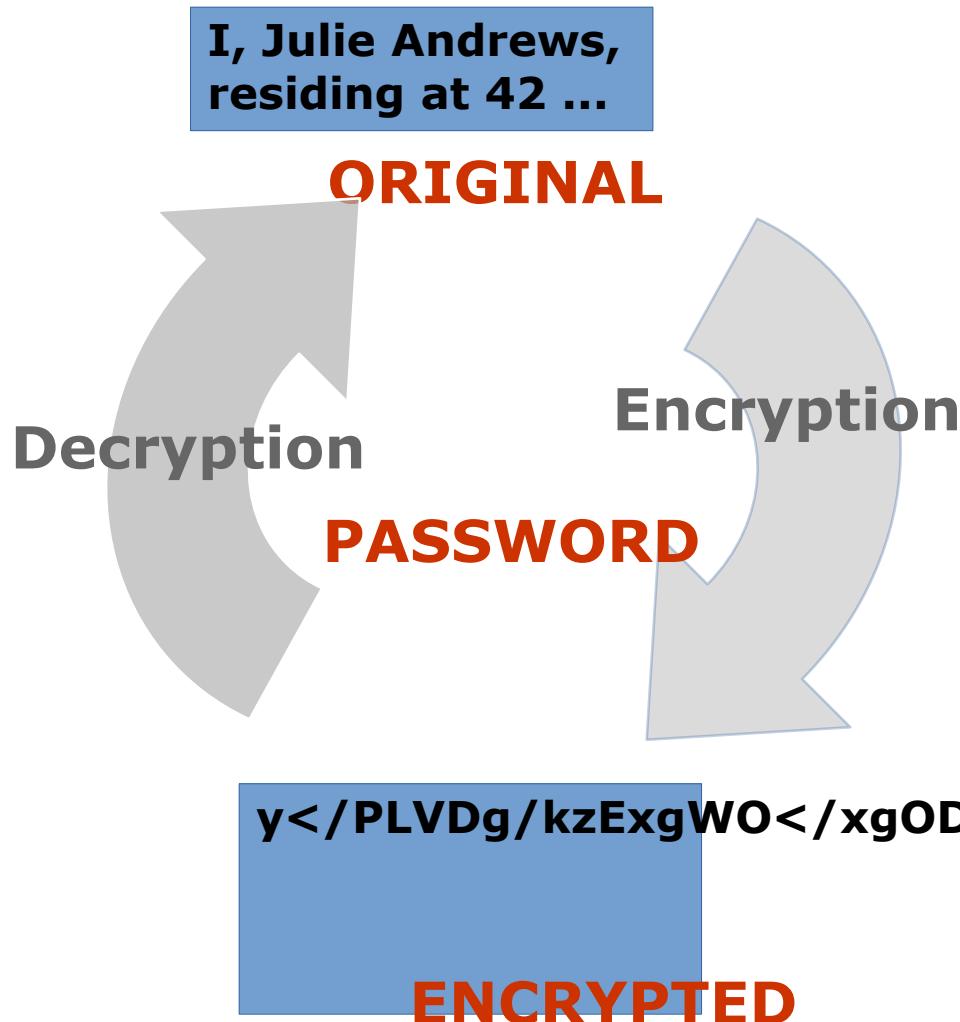
# Cloud Security Alliance



Broker  
Cloud  
Storage



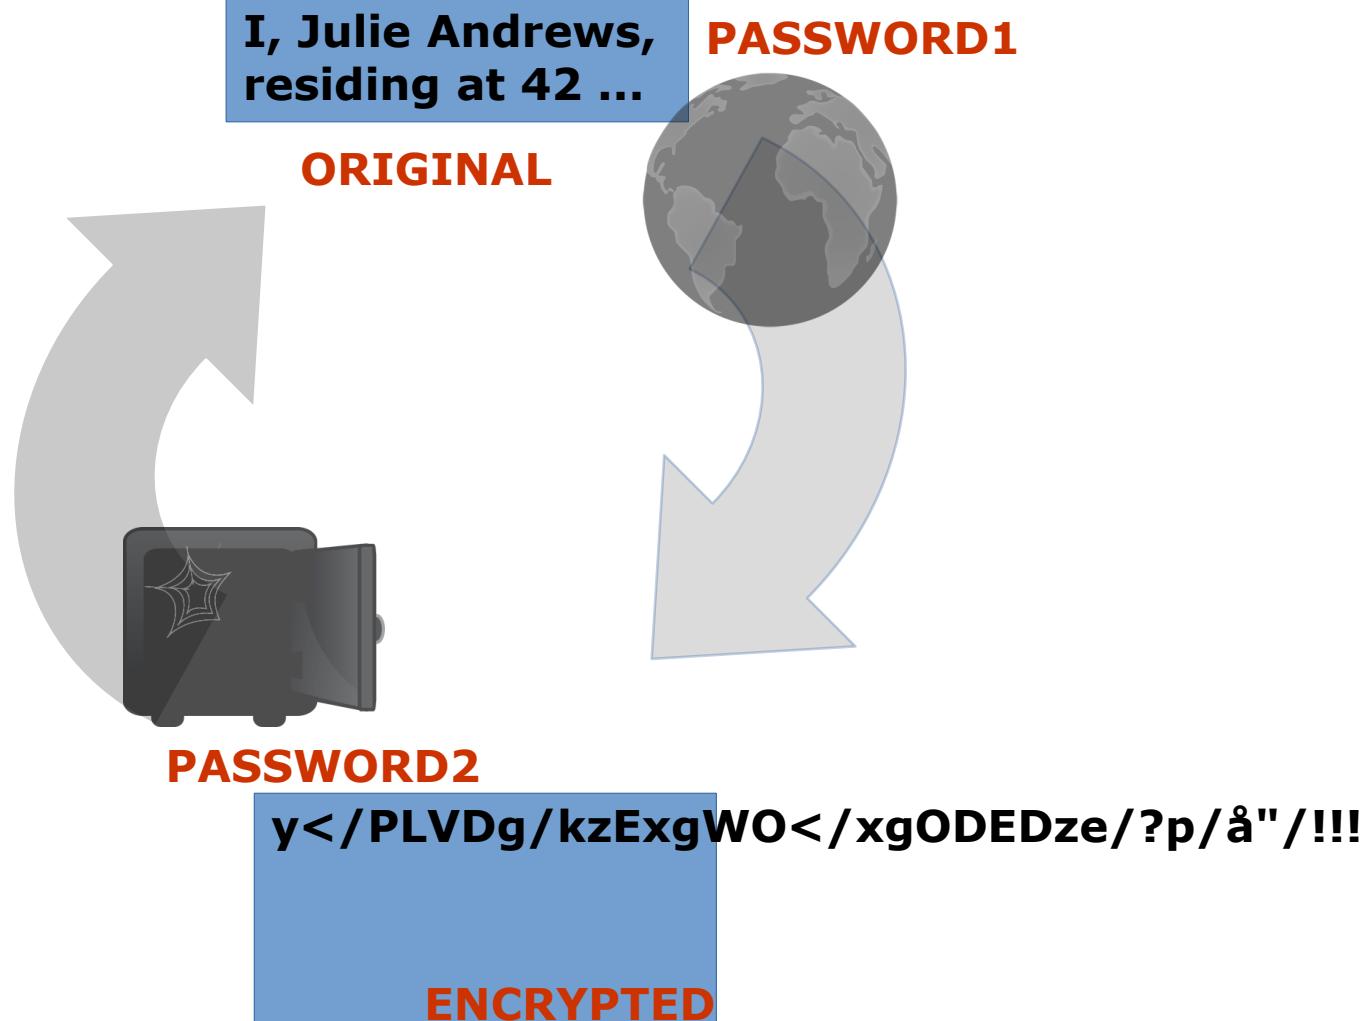
# Encryption Using Normal Passwords – data at rest



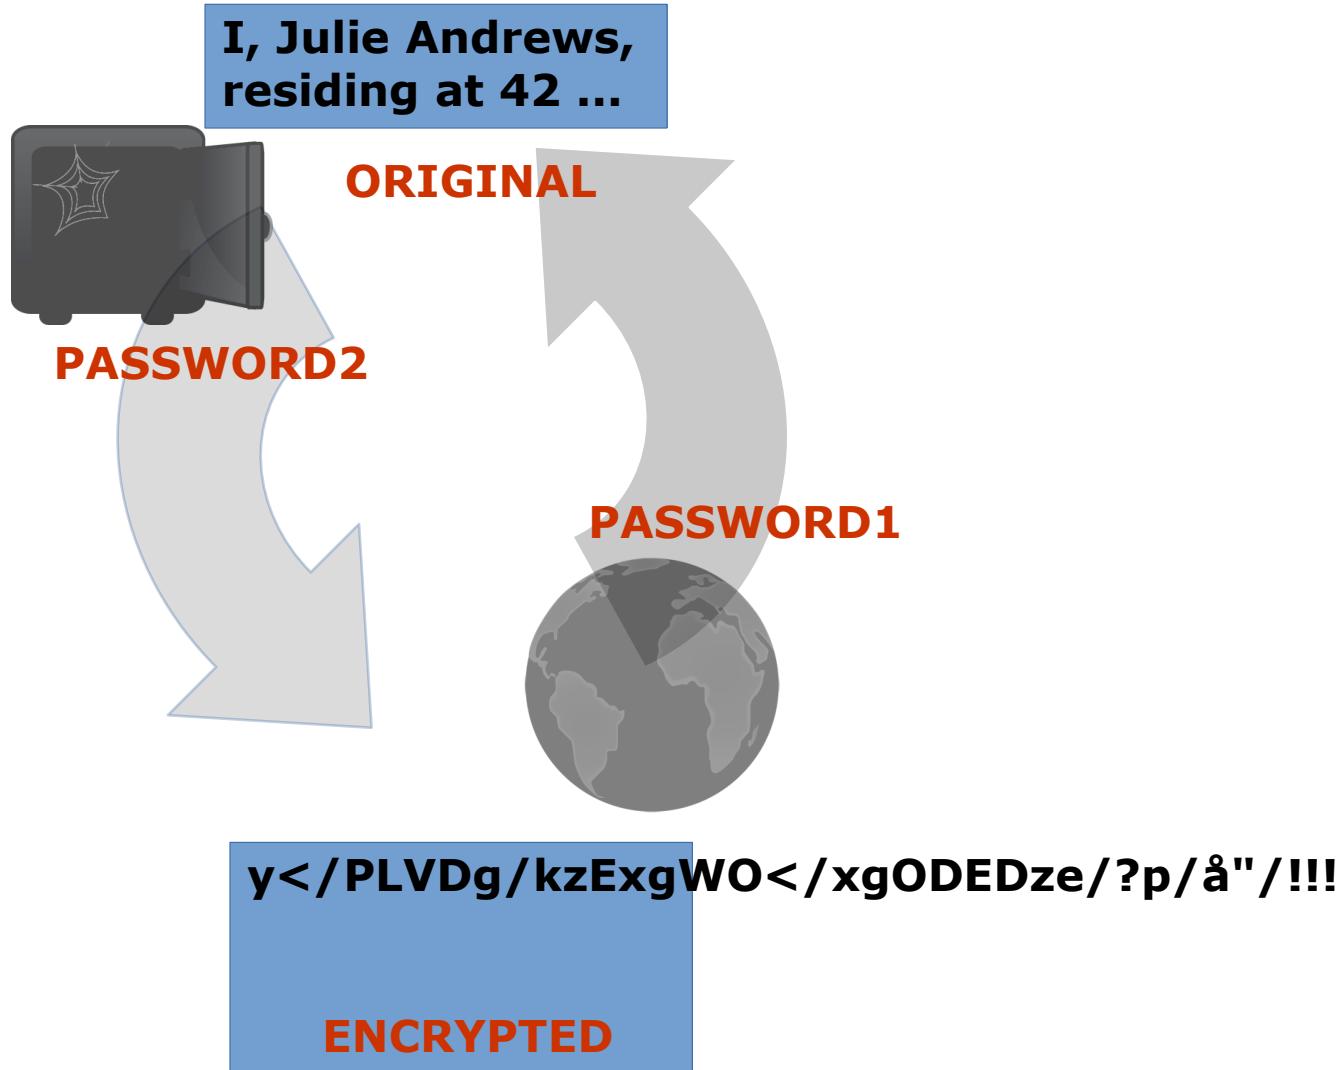
## How it works:

- At encryption
  - » Software asks for password
  - » You can pick any password
- At Decryption
  - » Software asks for password
  - » Must provide same password

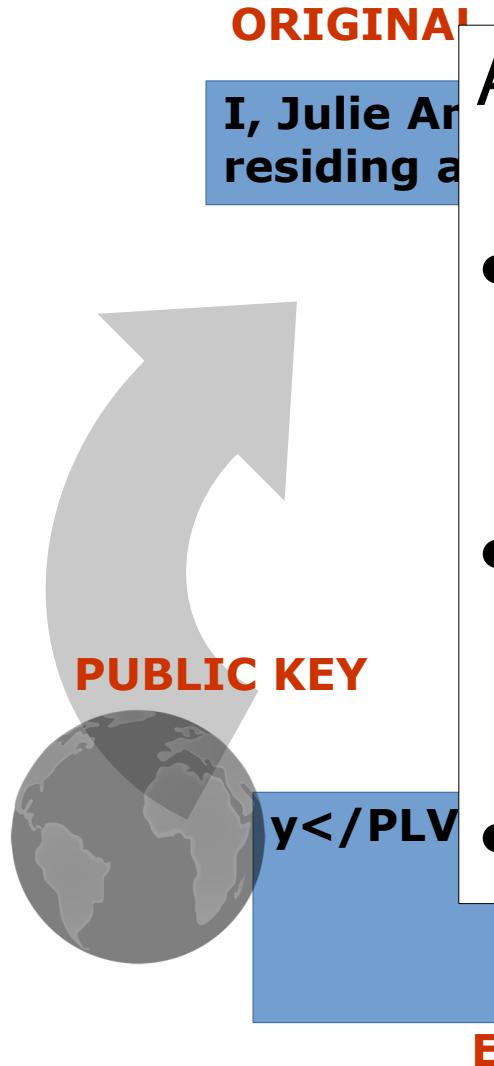
# Encryption Using 2 Passwords – data in motion



# Works in Reverse Too



# Authentication



## Authentication Usage:

- Initial setup for authentication:
  - Public key + Private key = Keypair
  - Everyone creates keypairs
  - Publicly publish all public keys
  - Private keys kept hidden
- None else can do this

# Advantages

- ▶ Bad traffic can be intercepted and stopped before it ever reaches the enterprise network to avoid any liability, threats
- ▶ Mobile devices are vulnerable to malware, not many antimalware for mobile and have overhead, cleaner data for lighter clients
- ▶ Filter at source, spam email filters, easy to catch in transit malware
- ▶ Security systems as building blocks, can be changed easily to adapt
- ▶ Spin up another IDS when malware detected requiring more resources
- ▶ Eliminates single point of failure
- ▶ Abilities on demand – download karate like Neo in The Matrix

# Limits

- ▶ Speed or bandwidth related
- ▶ Willingness of management to trust security in cloud
- ▶ Regulators' reluctance, consider abstraction layers as obfuscation
- ▶ Existence of on prem firewalls preventing cloud based firewalls adoption
- ▶ Customization of authentication systems (granularity of identifying groups of users, user defined fields)
- ▶ Regulators and cloud providers working together towards cloud security

# Future

- ▶ Improving compute and bandwidth is helping
- ▶ Fuzziness of network boundary due to mobile devices, telecommuting or work from home (WFH) employees
- ▶ Use case: Card swipe for physical entry and VPN login for the same user from another country at the same time
- ▶ Single security infrastructure for various OS, platforms, from anywhere
- ▶ More security
- ▶ Ease of management with single pane

# Cloud Challenges

SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# Stability

- ▶ How many internet feeds you have? Do they come into your data centers through different cable paths? Consider road digging and fiber cut!
- ▶ Kind of backup power. Kind of fuel for generators and how much they have. Are they tested and serviced regularly.
- ▶ Where is DNS hosted? Its availability on site as well as in Internet.
- ▶ NAT, network demarcation point, out of band network
- ▶ Remote access to consoles
- ▶ Change management system to track power down
- ▶ Physical security like access cards etc.
- ▶ Administrators, on site spare parts

# Partner Quality

- ▶ Provider's relationship with its partners
- ▶ (Not so normal) circumstances like slammer
- ▶ How fast provider can work with partners setting up ACLs to block attacks
- ▶ Inadvertent changes in internet routing e.g. BGP peering table
- ▶ .edu to .com within city of Honolulu went all the way through Chicago!

# Longevity

- ▶ How long your provider can support you
- ▶ Cloud provider shuts down: PictureBug closed in October 2007
- ▶ Cloud failure can bring your business to standstill
- ▶ Imagine google (along with its gmail and google accounts) close down!
- ▶ Can I get my data, apps, images back?
- ▶ Advance notice, contracts, enforcements
- ▶ Look out by provider's financial reports
- ▶ Cash rich like Amazon or backed by experience and trustworthy
- ▶ Provider's staffing, automation, available templates, documentation

# Service Level Agreement (SLA)

- ▶ Few risks and many benefits as possible
- ▶ Troubleshooting capabilities
- ▶ Should / should not be negotiable
- ▶ Cover everything that can possibly go wrong
- ▶ Cover damages in case of business impact
- ▶ As quantitative and objective as possible
- ▶ Exercise regularly and not only when things go wrong
- ▶ Keep data ownership with you
- ▶ Have plan B ready
- ▶ Ensure physical jurisdiction, tech and legal hand-in-hand

# Reaching agreement

- ▶ Communication between VMs on the same host, tapping traffic for analysis
- ▶ Promiscuous mode, peeking into vSwitch traffic
- ▶ Quality of Service: understandable voice over VoIP, not much interruptions due to buffering while watching a video, performance expectations, user experience, profiling, measurable metrics
- ▶ Don't rely completely on the provider for security but have your own layer
- ▶ Control over failover and load balancing in cloud
- ▶ Regulatory : putting in cloud considered as publishing giving up patenting rights, auditing

# Business factors

- ▶ Who has control over data
- ▶ Training and personnel cost, hidden costs
- ▶ Regulations, where is my data (load balancing)
- ▶ Abstraction layers make it difficult to enforce policy (proxied internet access)
- ▶ Does cloud support your applications or is migration required?
- ▶ Bandwidth support (cost, apps, firewalls) to and from external cloud
- ▶ Speedtest.net, iperf, ixia chariot application
- ▶ Load balancers: traditional, dedicated servers for peek load, vMotion
- ▶ On prem – cloud, salesforce - SugarCRM combos
- ▶ Websites, liability, SMBs partnering with VARs
- ▶ Collaboration through cloud, shared s3 storage without compromising privacy
- ▶ Identity federation (LDAP, eDirectory, Active Directory)

# Planning and Future for Cloud

SAMEER MAHAJAN

[HTTPS://WWW.LINKEDIN.COM/IN/SAMEERSMAHAJAN/](https://www.linkedin.com/in/sameersmahajan/)

# Licensing

- ▶ Per CPU, VM has single CPU but physical machine may have 8
- ▶ Multiple copies on VMs on the same physical machine
- ▶ License dongles
- ▶ License servers
- ▶ Some versions may not support virtualization
- ▶ Overhead of abstraction, tuning for optimal throughput

# First Steps

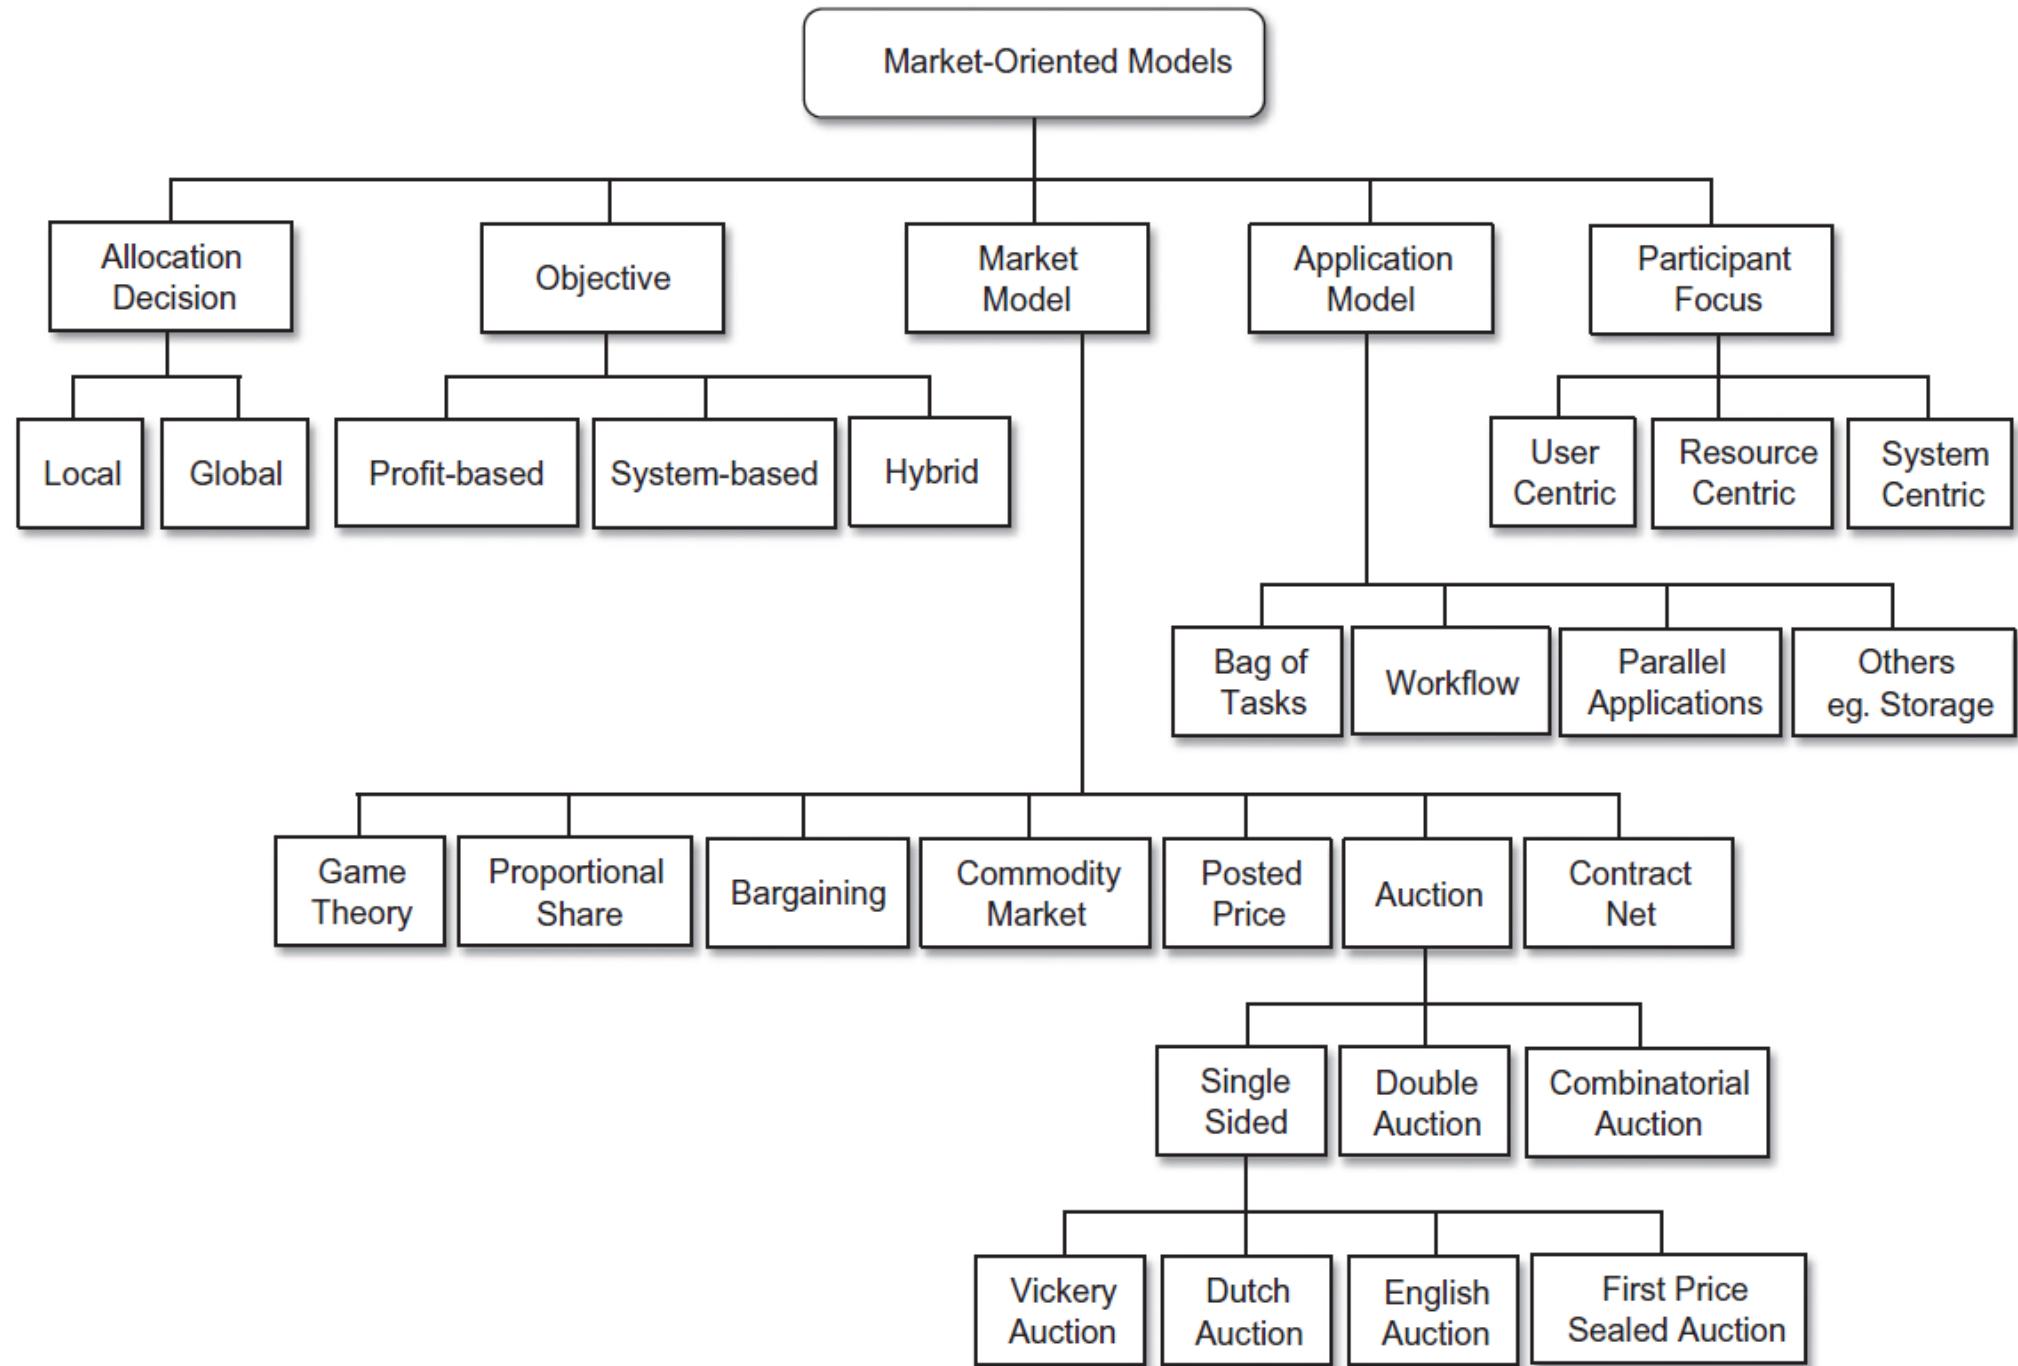
- ▶ Short term and long term, revisit at every technology refresh cycle
- ▶ Keep asking: “does it make sense for my organization?”
- ▶ Might have to adopt some cloud due to partners / vendors etc.
- ▶ Test drive / trial
- ▶ If no 24/7 data centers in house, you can shutdown at nights, weekends
- ▶ Collocation as first step towards moving to cloud completely
- ▶ Remote access, IP KVM, console access, VMWare, Hyper-V support

# Further steps

- ▶ Peaks and Valleys – auto scale, blade servers can run full load all the time
- ▶ Energy efficiency: power supplies are most efficient with in the middle of their load range, green devices, green incentives and initiatives
- ▶ Starter kit, migration utilities, eat learning curve / training / getting used to
- ▶ Virtualization as a stepping stone towards cloud, free trial software, old / retired servers
- ▶ Planning for success, cloud just an aid, well understood application
- ▶ Websites for contests, globally accessed CRM

# Advancements

- ▶ Energy efficiency: single datacenter consumes as much energy as 25000 households, doubles every 5 years, 42% of Amazon's budget
- ▶ The Green Grid: increase energy efficiency, reduce carbon footprint, \$23 (2010) -> \$16 (2020) billion, 28% less emission
- ▶ Green scheduling
- ▶ Market oriented, broker, bid / negotiate
  - ▶ EC2 spot instances
  - ▶ SpotCloud
  - ▶ AppSpot

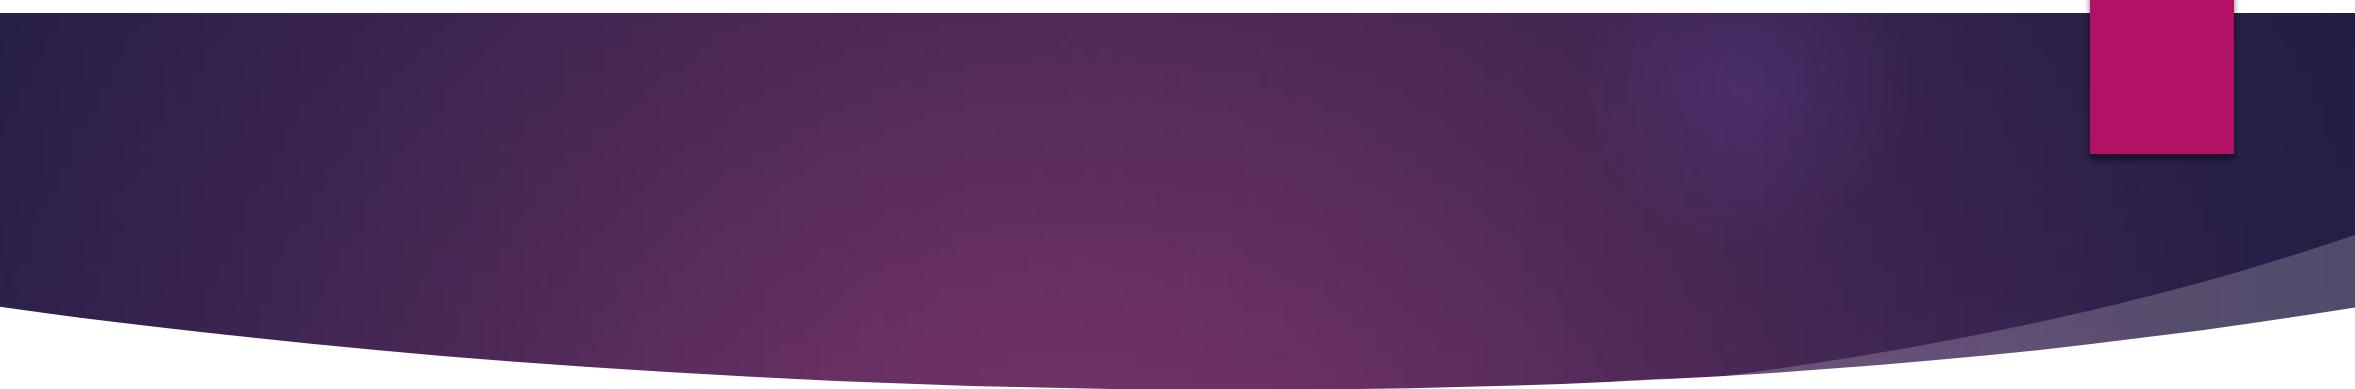


# Advancements (contd)

- ▶ Federated clouds
  - ▶ ad hoc
  - ▶ agreement between cloud providers to use each other's services
  - ▶ Proprietary interfaces
- ▶ InterCloud
  - ▶ Standards
  - ▶ Interoperability
- ▶ Third party cloud services: MetaCDN (over storage clouds), SpotCloud

# Future

- ▶ Standardization
- ▶ BOINC, Blockchain, leverage distributed available compute
- ▶ Specialized clouds
- ▶ Google docs, Microsoft office live, quora
- ▶ Mobile cloud: heavy lifting by back end cloud, maps
- ▶ Multi programming -> multiple VMs
- ▶ Cloud extensions: AutoCAD -> Building Information Modeling (BIM)
- ▶ Cloud Descriptor Language (CDL): capturing resource requirements
- ▶ Cost, control, bottom line – profit
- ▶ Full cycle; mainframe – dumb terminals, cloud – thin clients
- ▶ Ubiquitous computing, abstraction -> focus on business rather than tools



Thank You  
&  
Good Luck